Meridian 2 guide

The new interface: every screen, what its numbers mean and where they come from, what each role can do, the keyboard, and what to do when something is wrong.

Printed 2026-10-30 from docs/guide.

Contents

The screens

- Meridian in three minutes
- Landscape
- An object and its tabs
- The rule canvas and the rules library
- The Work ledger
- Approvals and exceptions
- People and inbox
- Runs
- Run and Compare
- Systems
- Golden records
- Glossary
- Lineage
- Processes and the designer
- Reports and schedules
- Waves and gates

Administration and settings

- Administration
- Preferences
- Keyboard and the command line

Walkthroughs by role

- Walkthrough: programme lead
- Walkthrough: data steward
- Walkthrough: data owner
- Walkthrough: SAP Basis
- Walkthrough: process analyst
- Walkthrough: auditor and viewer
- Walkthrough: administrator

The screens

Meridian in three minutes

Meridian checks your SAP data against fixed rules, scores every object, lists every record that fails and routes each fix through a second person before it goes back to SAP. It puts all of that on three surfaces: the **map** of systems and objects on the left, the **canvas** in the middle and the **inspector** on the right. There is no menu of pages. Every screen is a place on the map, every piece of work is a row in the ledger, and anything you select opens in the same inspector.

The three surfaces

- **The map** lists **Landscape**, **Work**, **Runs** and **People** at the top, then each connected system with its objects. Every object row carries its score and, where there is work, the number of records to fix. On a phone the map is a sheet you open from the top bar.
- **The canvas** shows the place you chose: the landscape, one object with its tabs, one rule, the Work ledger, a run or a comparison of two runs, a system, People, Reports or Administration.
- **The inspector** opens beside the canvas when you select a row: a record, a rule, a run, a job, a batch, a person, a system. It always has the same order: title, status line, facts, actions, history. `Esc` closes it and returns you to where you were.

The top bar holds the scope (which systems), the run you are reading, the command line (`⌘K` or `/`), the inbox count, help (`?`) and your account menu with Preferences and, for administrators, Administration.

The numbers on every screen

Every score in Meridian is drawn on one **ruler**: six segments whose widths are the dimension weights (completeness 25, accuracy 25, consistency 20, timeliness 10, uniqueness 10, validity 10), each filled to its own score, with a tick at the **measured** score. The number beside the ruler is the score **shown**. When a critical rule is over tolerance the shown score is capped and the word **capped** appears beside it.

The cap rule is fixed: one critical rule over tolerance caps a score at 85, two or more at 70. A critical rule counts as over tolerance when it fails at least the share of records or the number of records set under Administration, Settings, Cap rule and tolerance. The receipt under every object score states it in one line, for example: "Measured 71.8, the tick on the ruler. Capped at 70 by MM014 and MM031, two critical rules over tolerance."

Verdicts come from the score bands: **Go**, **At risk** and **No-go**. Value at risk comes from the cost model named in the footer, such as "cost model ZAR v3". Nothing is estimated by a language model: every figure is a rule result over stored data, and each one has an **i** beside it that shows the arithmetic.

Where each role starts

Everyone lands on the **Landscape**, shaped for their role by one **next-step button**:

- Programme lead: "Clear the 3 no-go objects in Wave 1: finance foundations".
- Data steward: "Accept the suggested values".

- Data owner: "Approve batch B-241".
- SAP Basis: "Extract SuccessFactors" (the system not yet extracted).
- Process analyst: "Open PTP step 3".
- Auditor and viewer: the same landscape, read-only.

Under **Preferences**, **Start screen**, you can open Meridian as another role's start screen. Your access does not change.

Keyboard

- `⌘K` or `/` opens the command line. Type a material number, a rule id such as `MM014`, an object, a run, or a word such as approve or export.
- `g` then `l`, `w`, `r` or `p` goes to Landscape, Work, Runs or People.
- `?` shows every shortcut. `⌘Z` undoes the last change. `Esc` closes the top layer.

When something is wrong

- **A grey dash** means the number is unknown, never zero. Hover it for the reason.
- **"Not extracted yet. Extract"** on a system means nothing has been read from it; **"Extracted, not analysed. Run analysis"** means no scores exist yet.
- **The red line at the bottom of the screen** names a failed job, such as "Extraction of Employee central failed: OData refused credentials (HTTP 401)". Open it for the stages and a retry; SAP Basis owns connection errors.
- **"... needs the ... permission. Ask an administrator to grant it."** means your role does not hold that action. Your administrator grants it under Administration, Users.
- For anything else, open help (`?`), copy the diagnostics block under **Contact support** and send it to your administrator. It holds no record values and no secrets.

Landscape

The landscape is the start screen for every role. It answers three questions in order: how ready is the data across every system and object, what holds it back, and what is the one thing you should do next. It opens at `/`, from **Landscape** at the top of the map, or with `g` then `l`.

What is on it

- **The title and scope**: "Landscape" and the systems in scope, such as "All systems". Change the scope and the run from the top bar.
- **The sentence**: one line that says what holds the landscape back, for example "11 of 15 objects are held under 85 by 70 critical rules over tolerance; fixing MM014 and MM031 would release 2 of them."
- **The provenance line**: records checked, rules measured and the run, for example "400,662 records checked, 5,178 rules measured, run 16 Oct 2026, 05:17".
- **The next-step button**: one per role (see below), with the wave and its due date under it for the programme lead.
- **The coverage matrix**: objects as rows, systems as columns. Each cell is a small ruler with the shown score, the word **capped** where a cap applies, the verdict (Go, At risk, No-go), records checked and rules measured. A cell flagged **low coverage** has fewer than 70% of its rules producing a result. A system never extracted shows "Not extracted yet. Extract"; one extracted but not scored shows "Extracted, not analysed. Run analysis".
- **Coverage, All objects and Waves**: the view switch. **Waves** groups the objects under each migration wave with its due date, its go, at-risk and no-go counts and the gate state (see the Waves and gates chapter).
- **Your next work**: the top five rows of the Work ledger for you, ranked by value at risk, with **View all**. SAP Basis sees **Your next work** as system tasks instead: extract again, run analysis, extract a system never read.

The numbers and where they come from

- **The score in each cell** is the object's shown score in the run in the top bar: the weighted dimension scores, capped at 85 by one critical rule over tolerance and at 70 by two or more. Open the **i** beside the object name for the receipt.
- **The verdict** comes from the score bands under Administration, Settings: at or above target is go, at or above warning is at risk, below that is no-go. A critical rule over tolerance keeps an object out of go.
- **The sentence's counts**: "held under 85" counts the objects whose shown score is capped; "critical rules over tolerance" counts the rules behind those caps; "would release" names the rules whose fix lifts the most objects off their cap.
- **Value at risk** in Your next work is the cost model named under the table: "Rows ranked by value at risk, cost model ZAR v3."

- **The footer** repeats the run, the records checked, the rules measured, the scoring time, the cost model and the licence validity, and shows a failed job in red when one needs attention.

What you can do, by role

- **Programme lead**: follow the next step ("Clear the 3 no-go objects in Wave 1: finance foundations"), switch to **Waves**, open any cell, share the page from **Share and export**, export the grid to Excel.
- **Data steward**: "Accept the suggested values" opens Work with your rows; Your next work lists your highest-value records.
- **Data owner**: the line under the sentence says how many batches wait for you and your SLA, for example "2 batches wait for your approval, SLA met on 92% of your open work"; the button is "Approve batch B-241".
- **SAP Basis**: the button extracts or analyses the system that needs it first; the task list below shows every system's state.
- **Process analyst**: "Open PTP step 3" opens the first blocked process step.
- **Auditor and viewer**: read everything; actions are not offered.

Selecting a cell opens the object canvas. Selecting a system heading opens the system in the inspector.

Keyboard

- `g` then `l` returns here from anywhere.
- `⌘K` opens the command line; type an object name and system, such as "material master ECC".
- In the matrix and in Your next work, `j` and `k` move, `Enter` opens.

When something is wrong

- **Every object reads 70.0**: most objects have two or more critical rules over tolerance, so the cap holds them at 70. The measured score differs; read it on the ruler tick or in the receipt.
- **A system row says "Not extracted yet"**: ask SAP Basis to extract it, or extract it yourself if you hold trigger sync.
- **A cell shows a dash**: no run has scored that object on that system. "Not analysed" and "Not extracted" say which step is missing.
- **Low coverage**: some rules could not run because a table, field or value list was missing. Treat the score as indicative and ask SAP Basis to check the extraction scope.
- **The footer shows a failed job**: open it from the footer link; the job inspector says why it stopped and offers a retry.

An object and its tabs

An object is one area of SAP data scored as a whole on one system, such as Material master on ECC production. The object canvas is where you read its score with the receipt that proves it, and where every view of the object lives as a tab: Rules, Records, Fixes, Runs, Readiness, Processes, Lineage, Glossary and Golden records. Open it from a cell on the landscape, from an object row in the map, or by typing its name in the command line. Its address is `/o/{system}/{object}`, for example `/o/ecc-production/material-master`.

The header

- **Watch**: "You hear about runs, score drops, new critical breaches and comments on Material master." Watching is per person; what reaches you is set under Preferences, What you hear about.
- **Data contracts**: the contracts that name this object.
- **Related objects**: each related object with its relation, **upstream** or **downstream**. Fix the upstream object first: records there fail before this object can pass.
- **The facts line**: the run, records checked, rules measured, records to fix, value at risk, S/4 features blocked and the wave verdict, for example "Run 16 Oct 2026, 05:17, 61,816 records checked, 913 rules measured, 25,712 to fix, ZAR 4.2 m at risk, 3 features blocked, No-go for wave 1".

The ruler and the receipt

The large ruler has the 0 to 100 axis with the 60 critical, 75 at risk and 90 go marks, and two cap lines, **cap 1 critical** (85) and **cap 2 or more** (70). Each of the six segments is filled to its dimension score, which is printed under it. The tick is the measured score; the number beside it is the score shown, with **capped** when a cap applies.

Under it the receipt says it in words, exactly as the product states it: "Measured 71.8, the tick on the ruler. Capped at 70 by MM014 and MM031, two critical rules over tolerance."

The receipt table then adds the score up: one row per dimension with its weight, score, points, rules measured, records checked and coverage; then **Weighted**, **Measured**, **Cap, 2 critical rules over tolerance** and **Score shown**. The formula line under it is the full rule: the six weights, "severity weights critical 4, high 3, medium 2, low 1; tolerance 10% of rows or 25 rows, as set on each rule; cap rule: 85 at one critical rule over tolerance, 70 at two or more." Every row has an **i** that explains the figure from the run itself.

Critical rules over tolerance lists the rules behind the cap with their failing records and share, for example "MM014 Base unit of measure not valid, 3,640 failing, 5.89%".

The tabs

- **Rules**: every rule of the object, ranked by value at risk, with field, dimension, severity, tolerance, over tolerance, failing rows, failing %, coverage and value at risk. Each rule says what goes wrong in SAP while it fails, for example "Goods movements and MRP cannot post for a material whose base unit is not an ISO unit." Where suggested values exist, **Accept all** accepts them for

that rule. Pivot the list **By rule**, **By field**, **By dimension**, **By severity** or **By S/4 feature**; **Learned** and **Drafts** list rules waiting for review. The line under it reads "Ranked by value at risk; total ZAR 4.2 m across 913 rules measured. Failing % is of 61,816 records checked."

- **Records**: the Work ledger scoped to this object. See the Work ledger chapter.
- **Fixes**: the fix batches of this object. "No fix batches for Material master yet. A batch starts when a steward accepts suggested values or imports a filled fix sheet."
- **Runs**: the measured score over the last eight runs with the cap line and the baseline, a table of measured and shown scores per run ("Scores are as stored with each run; the measured score is before the cap, the shown score after it."), **Compare with** to open two runs side by side, and **Why the score moved** against the baseline: rules newly failing, fixed, worse and better.
- **Readiness**: the migration waves this object belongs to with the verdict, blocking gaps, score, readiness and records blocked, and "What the failing rules do in S/4HANA": the features blocked and degraded, each with the rules behind it.
- **Processes**, **Lineage**, **Glossary**, **Golden records**: each has its own chapter.

What you can do, by role

- **Everyone**: watch the object, read every tab, export any grid with **Share and export**, comment on a rule in its inspector.
- **Data steward**: **Accept all** suggested values on a rule (needs apply), open the records to fix.
- **Data owner**: apply the steward template to objects you own, from the object's steward template action.
- **Programme lead**: share the page as a link, compare runs.
- **Process analyst**: the Processes tab and its what-if simulation.

Keyboard

- In every tab table, `j` and `k` move, `Enter` opens the row in the inspector, `x` selects.
- `[` and `]` step to the previous or next row while the inspector is open.
- `Esc` closes the inspector.

When something is wrong

- **"No run yet"** and an empty ruler: this object has never been analysed on this system. Ask SAP Basis to extract and analyse it.
- **The measured and shown scores are the same, but the receipt names a critical rule**: "A critical rule is over tolerance, but the score is already below its cap."
- **The Records tab says the ledger could not be loaded**: open Work from the link under the message; the ledger reads the latest run of each object.
- **Low confidence**: fewer than 70% of the rules attempted produced a result. The receipt row **Coverage** shows which dimension is short. SAP Basis can widen the extraction.

- **A figure looks wrong**: open its **i**. If the arithmetic there does not match, copy the diagnostics from help and send them to your administrator.

The rule canvas and the rules library

Every number in Meridian comes from a rule: a fixed, deterministic check with an id such as `MM014`, a field, a dimension and a severity. The **rule canvas** shows one rule across every object and system it runs on. The **rules library** lists every rule of the tenant, with drafts and suppressions. Open a rule by typing its id in the command line, from any rule row, or at `/r/{id}`. Open the library from the command line ("Rules") or at `/rules`.

The rule canvas

The title is the rule id and name, for example "MM014 Base unit of measure not valid". The facts line under it reads like "Critical, Validity, 4,212 failing across 2 objects, 4.99% of records checked, ZAR 1.18 m at risk, over tolerance, blocks 3 S/4 features".

Four sections follow:

- **Definition**: the check class (such as `domain_value_check`), the field (`MARA.MEINS`), the expression (`MARA-MEINS in T006-MSEHI`), its parameters, the source ("Shipped with Meridian" or a custom rule), the authority ("SAP hard constraint") and what goes wrong in SAP. **Versions** lists every version with who changed it and why; the shipped definition names its file and content hash.
- **Failing records**: the scope, for example "Scope: 2 objects, 1 system", with each object, its run and its failing count. **Open failing records** opens the Work ledger filtered to this rule.
- **Impact**: the S/4 features this rule blocks or degrades and the lineage from field to rule to feature.
- **History**: each run's failing count next to the version changes, so a jump in failures can be read against a change to the rule.

Contract shows the data contracts that name this rule.

The rule inspector

Selecting a rule anywhere opens the inspector with the same facts: severity, dimension, field, **Tolerance** (for example "25 rows"; only critical rules carry a tolerance), **Over tolerance**, failing records by object, **Failing %** ("4.99% of 84,456 records checked") and **Value at risk**. Below come the definition, scope, the features it blocks, lineage, versions, contract, history, **Watch** and **Comments**.

The rules library

The library title says what is there, for example "38 rules across 19 objects, 38 enabled". Three views: **All rules**, **Drafts** and **Suppressed**. Filter by object, dimension, severity, check class and state; the table shows rule, object, severity, dimension, class, failing records, pass rate, state and source. **New rule draft** starts a rule in authoring; a draft runs nowhere until it is reviewed and enabled.

Rules can also be imported and exported in bulk under Administration, Rules import and export: imported rules are validated against the data dictionary and land as drafts for review, never live.

The numbers and where they come from

- **Failing** is distinct records failing the rule in the latest run of each object in scope.
- **Failing %** divides by the records checked in those runs.
- **Over tolerance** applies to critical rules only: a critical rule is over tolerance when it fails at least the share or the count of records set as its tolerance. "Only critical rules count: when one is over tolerance it caps the score (85 for one rule, 70 for two or more); a critical rule within tolerance still lowers its dimension but never caps."
- **Value at risk** is from the cost model of the run, in the tenant currency.
- **Pass rate** is the share of checked records that pass.

What you can do, by role

- **Everyone**: read, watch, comment, export the library.
- **Data steward and rule maintainers** (manage rules): **New rule draft**, **Edit in authoring**, **Suppress** with a reason, **Hand to a steward**.
- **Process analyst** (handoff rule): hand a blocking rule to a steward from the step inspector.
- **Platform admin**: enable, disable and change severity; bulk import and export.
- **Auditor**: read the versions and the history; the rule's change trail across versions and runs is read side by side, there is no single joined trail yet.

Keyboard

- ⌘K, then a rule id, opens the rule.
- In the library table, `j` and `k` move, `Enter` opens the rule in the inspector.
- `Esc` closes the inspector.

When something is wrong

- **A rule fails far more records than last run**: compare **History** with **Versions**. A version change on the same day usually explains it; otherwise the data changed.
- **Coverage below 100%**: the rule could not read some records because a table, field or value list was missing from the extraction.
- **A rule is wrong for your business**: suppress it with a reason, or mark the affected records not an issue in Work. Both are audited. Ask the rule's maintainer or your administrator to change its severity or tolerance.
- **Your edit buttons are missing**: authoring needs the manage rules permission.

The Work ledger

Work is the one list of everything that needs a person: every failing record and field, with its current value, the values the rule accepts, a suggested value, a New value cell and the decision, on one row. The row is the fix. Open it from **Work** at the top of the map, with `g` then `w`, or at `/work`. On an object, the **Records** tab is the same ledger scoped to that object.

What is on it

- **The title and scope**: "Work" and the objects in scope, such as "All objects". Click an object in the map to scope the ledger to it.
- **The sentence**: what the current filter holds, for example "25,712 records to fix, ranked by value at risk." and, under it, "Rows are read from the latest run of each object. Value at risk is from cost model ZAR v3."
- **The state chips**: **To fix**, **Fix proposed**, **Waiting for approval** and **Exceptions**, each with its count. The chips are the filter.
- **The controls**: **Flat** or **By rule** grouping; **Sort** by value at risk, due, age or owner; **Compact** or **Standard** rows; **Views** (saved and shared views); the object filter; **Export fix sheet** (`e`) and **Import filled sheet**.
- **The rows**: the record key (for example Material number and Plant), name, rule, severity, field description, SAP field (`MARA.MEINS`), problem, current value, possible values, suggested value, New value, decision, first seen, days open, owner and value at risk.
- **The fix-order hint**: where an upstream object still fails, a line says so first, for example "Fix Material master first: 3,640 upstream failures. Records there fail before this object can pass."

The ledger loads 100 rows at a time ("Showing 100 of 25,712. Sorted by value at risk within the rows loaded.") with **Show 100 more**.

The numbers and where they come from

- **The chip counts** are distinct records in each state in the latest run of each object in scope, whatever the rows loaded.
- **Value at risk** is the cost-model estimate of what the failing value costs, in the tenant currency.
- **Possible values** are the values the rule accepts; open the row for the full list with meanings.
- **Suggested value** is the value the rules propose. "Blank when no single value is right."
- **Days open** counts from the day the oldest failing issue of the row was first seen.

Editing a row

Type in **New value** and press `Enter`; until you do the cell reads "Not saved yet: press Enter". Or choose a **Decision**: **fix** (accept the suggested or entered value), **not an issue** or **accept risk**. A record in the inspector shows every failing field with the same cells, plus its facts, feature impact, provenance and history. On a phone, editing happens in the inspector.

Every change shows a notice at the bottom of the screen, such as "Accepted 1 suggested value", with **Undo**. ⌘Z takes the last change back while focus is not in a field: "Undone: accepted 1 suggested value". The chip counts move at once.

Working in bulk

- **Select rows** with the checkboxes, x, or Shift+J and Shift+K to extend. The bulk bar offers **Value for all selected rows**, accept the suggested values, not an issue, accept risk and **Assign...**.
- **Fill down**: select the rows, then press ⌘D in the New value you want to copy: "Filled EA down 12 rows".
- **Paste a column** of values from Excel into a New value cell to fill the selected rows.
- **Accept all suggested** in a rule group, or the next-step button "Accept the 112 suggested values".
- **By rule** groups the rows to fix under their rule, so you can decide a whole rule at once.

Accepted and entered values go into your draft fix batch. They wait there for a second person: see the Approvals and exceptions chapter.

Export and import

Export fix sheet (e) saves the rows on screen as an Excel sheet with the same columns, in business words. Fill in New value and the decision offline, then use **Import filled sheet**: "Drop the filled fix sheet here, or choose a file (.xlsx). Meridian checks keys, SAP fields and values before anything changes." The result lists what was accepted and what was refused, row by row.

Share and export, above every grid, also offers **Export to Excel** (the visible columns, as a branded workbook), **Save these columns as a view** (shared with everyone if you choose) and **Share this page as a link** (read-only, for people of your tenant, until it expires).

What you can do, by role

- **Data steward** (apply): edit values, accept suggestions, bulk actions, import, assign to yourself.
- **Data owner** (approve): sign records off as not an issue or accept risk, approve batches.
- **Programme lead** (assign): assign rows, balance work from People.
- **Auditor, viewer, analyst**: read and export; the ledger says "Read-only for your role".

Keyboard

- j or ↓, k or ↑ move; Shift+J and Shift+K extend the selection; x selects; Enter opens the row.
- a accepts the suggested value on the selected rows; n marks them not an issue; r accepts the risk.
- e exports the fix sheet of what is on screen; f moves to the filter chips.
- ⌘D fills down; ⌘Z undoes; [and] step through rows in the inspector.

When something is wrong

- **"The work list of the latest run is not built yet, so no rows can be listed. It is built after each run finishes."** Wait for the analysis to finish, or ask whoever started it.
- **"No rows in this view."** The filters exclude everything. **Reset to the default view** from Views.
- **"Select rows first: x selects the highlighted row, Shift+J extends the selection."** A bulk key was pressed with nothing selected.
- **"The file could not be read as a fix sheet."** Export a fresh sheet and keep its columns and key fields unchanged.
- **A value you typed is refused**: it still fails the rule. Pick one of the possible values.
- **A record you fixed is back in To fix**: the next run still found it failing after it was sent to SAP. Check the load reached the right client, then decide again.

Approvals and exceptions

Every value decided in Work waits in a fix batch until a second person approves it. That is the four-eyes rule, and it has no bypass: whoever drafted a batch cannot approve it. Exceptions are records signed off without a change, as **not an issue** or **accept risk**, with a reason. Both live in the Work ledger, under their own chips: **Waiting for approval** and **Exceptions**. Open approvals from the command line by typing "approve", or at `/work?filter=approvals`.

Waiting for approval

The sentence reads, for example, "612 records wait for approval: the values are in a draft fix batch that a second person approves." The rows are grouped by batch. Each batch heading names the batch (for example **B-241**, "Unit of measure and material type"), who built it, the records, the value at risk, the due date and the state, with **Approve** and **Send back** on the heading.

Open a batch to read it in the batch inspector:

- **Facts**: batch, name, status, owner, steward, created, due, records and value at risk.
- **Records**: each change as the SAP field and the pair of values, for example "`MARA.MEINS` STK to EA". The first 50 show in the inspector; **Open all 328 in the ledger** opens the rest.
- **Approval trail**: who drafted it and who proposed the values, with times.
- **Actions**: **Approve**, **Send back** (with a reason the steward reads), **Download sheet** and **Reassign...**. "Approving acts on the whole batch. You can take it back from the notice that follows."
- **Watch** and **Comments**.

After approval

Approval does not touch SAP. It is Meridian's record that the values are right. The approved batch is exported as a load file for your SAP process, and records then read **Sent to SAP**. Where live write-back is switched on for the deployment, an approved change can go to SAP through that route instead, and each request needs someone other than the requester. Either way the work is not finished until a run says so: the next finished run reads the records again, and those that pass become **Fixed**. A record that still fails returns to **To fix**.

The record statuses, in order: **To fix**, **Fix proposed**, **Waiting for approval**, **Approved**, **Sent to SAP**, **Fixed**. Exceptions are **Not an issue** and **Accept risk**.

Exceptions

The sentence reads "48 records are signed off as not an issue or an accepted risk." The columns change to **Resolution**, **Reason**, **Signed off by** and **When**, for example "Accept risk: Replaced by the new plant in wave 2; the risk is accepted until then". **Reopen** "Puts the record back to rows to fix. Needs the apply permission."

The numbers and where they come from

- **The chip counts** are distinct records in each state in the latest run of each object in scope: "612 records wait for approval" counts records, not batches.
- **Records** and **Value at risk** on a batch are the records in it and their cost-model estimate, in the tenant currency.
- **Due** comes from the SLA days for the batch's worst severity, set under Triage and SLA.
- **Fixed** is set only by a run: a value waiting for approval, approved or sent to SAP never counts as fixed.

Out of office and escalation

- When an approver is away, their approvals go to the deputy they named under Preferences, Out of office. The batch inspector says so, for example "You cover for T. Mokoena until 30 Oct 2026: their approvals and assigned work come to you as well." Four eyes still holds: nobody approves their own batch.
- A batch past its due date is escalated by the escalation chain under **Triage and SLA** (see People and inbox).

What you can do, by role

- **Data owner** (approve): approve or send back batches you did not draft; sign records off as exceptions.
- **Data steward**: draft batches, read their state, reopen exceptions; a steward who also holds approve can approve other stewards' batches, never their own.
- **Programme lead**: approve, reassign batches.
- **Auditor**: read batches, trails and exceptions as four-eyes evidence; export them.

Keyboard

- ⌘K, then "approve", opens the approvals.
- `j` and `k` move, `Enter` opens the batch or the row; `n` and `r` sign the selected rows off.
- ⌘Z undoes an approval or a sign-off from the notice.

When something is wrong

- **"No value is waiting for approval. Accepted and entered values land here until a second person approves them."** Nothing is pending.
- **Approve is missing on a batch**: you drafted it, or your role does not hold approve. Ask another approver.
- **"Sent to SAP" records come back as To fix**: the load did not take. Check the file was loaded completely, into the right client, and that the run you read finished after the load.
- **A sign-off is missing its reason**: the reason is required when you sign off; older records may show "The API does not send the reason yet".

People and inbox

People lists the owners and stewards with their open work, approvals waiting, overdue records and SLA, and is where work is balanced, contracts are read and out of office is set. The **inbox** holds what Meridian tells you: approvals, SLA warnings, breaches, finished runs, mentions and digests. Open People from the map, with `g` then `p`, or at `/people`. Open the inbox from **Inbox** and its count in the top bar.

People

The sentence sums the team up, for example "4 of 6 people hold open work: 270 records, 34 overdue. 2 batches wait for your approval." and "Counts are read from the work assigned in the latest state."

The table: person and role, objects (owned or stewarded), open work, approvals waiting, overdue, **SLA met**, breaches this month and last action. Your own row reads "you".

Above it: **Triage and SLA**, **Balance work**, **Data contracts** and **Out of office**.

The person inspector

Open a person for their facts (role, objects, open work, approvals waiting, overdue, SLA met, the SLA, breaches this month, last action), **Their objects**, **Open work by object**, what waits for their approval, **Open their work**, **Reassign…** and **Activity**: their own audit trail. Your own activity is always visible to you; another person's needs view audit or manage users.

Balance work

"An assignment plan, applied with undo." Choose how to balance:

- **Even counts**: "Every person ends within one record of the others."
- **Whole objects**: "A person keeps all of one object's records together."
- **By the assignment template**: "Each record goes to the steward named for its object and plant, company code or organisation."

Propose a plan shows the moves before and after. Applying moves the records; the notice that follows has **Undo**. Balancing needs the assign permission.

Data contracts

"What each consumer is promised." Each contract names a producer object, a consumer and a promise, such as "Score at 85 or more", with its compliance (**Breached** or met), when it was last checked, its **Versions** and **Breaches**. Owners and rule maintainers create drafts with **New contract**; a data owner approves a contract, never their own; **Revise as a new version** starts the next version as a draft.

Out of office

"Your approvals, SLA breaches, mentions and assigned work are copied to them between these days. Four eyes still holds: nobody approves their own batch." Choose who covers, the first and last day and an optional reason, then **Set out of office**. **You cover for** lists whom you deputise for, with **End it**.

Triage and SLA

SLA days to resolve by severity (by default critical 1 day, high 3, medium 7, low 14), when to escalate, and the **Escalation chain**: "After a breach the assignee and the team lead hear first. Each level below notifies once more when an item stays past due that long." Changing SLA days and the chain needs the manage triage permission; anyone can choose their own digest here.

The numbers and where they come from

- **Open work** counts the records assigned to the person in the latest state of the work list.
- **Overdue** counts their open records past the due date their SLA gives them.
- **SLA met** is the share of their open work still within its SLA days; a dash means they hold no open work.
- **Breaches this month** counts the records that passed their SLA this calendar month.
- **Approvals waiting** counts the batches waiting for that person to approve.

The inbox

The inbox panel says how many items are unread and how many mentions, groups items under **Today** and **Earlier**, and gives each one its kind, time, sentence and one action: **Approval** ("Batch B-241 waits for your approval. 328 records of Material master, ZAR 412,600 at risk."), **SLA at risk**, **SLA breached**, **Notice** (a finished run), **Data quality** (a rule that rose), **Exception**, **Digest**, **Report**, **Monitor**, **Cleaning**. **Mark read** and **Mark all read** clear them. **Notification settings** opens your preferences.

Comments with an `@` mention of you count as unread mentions until you open them.

What you can do, by role

- **Everyone**: read the inbox, set out of office, choose a digest, read People and contracts, see their own activity.
- **Programme lead and steward with assign**: balance work, reassign a person's work.
- **Data owner**: approve contracts, receive escalations, set a deputy for approvals.
- **Administrator and triage managers**: SLA days and the escalation chain.
- **Auditor**: a person's activity trail.

Keyboard

- `g` then `p` opens People.
- In the inbox and the People table, `j` and `k` move and `Enter` opens.
- `Esc` closes the panel.

When something is wrong

- **SLA met shows a dash**: the person has no open work to measure.
- **"Reassigning needs the assign permission, which your role does not hold."** Ask your lead.

- **A deputy cannot approve**: the deputy needs the approve permission themselves, and cannot approve a batch they built.
- **You hear about too much or too little**: Preferences, What you hear about, and per object or rule the **Watch** button.

Runs

A run is one measurement of your data at a point in time: an extraction from a system scored against the rules as they stood. Every score, failing record and report belongs to a run. **Runs** lists them on a timeline with their jobs, so you can see what ran, what failed and which run is the baseline. Open it from the map, with `g` then `r`, or at `/runs`.

What is on it

- **The sentence**: "The last 8 runs, the latest on 16 Oct 2026. The baseline is 2 Oct 2026."
- **The run timeline**: one mark per run, the baseline labelled. Select a mark to open that run.
- **Runs and Jobs**: two views. **Start a run** opens the start panel.
- **Runs** table: date, system, label, trigger (Scheduled or Started by hand), records, rules, score range across objects (for example "70.0 to 93.4"), status (Complete, Failed, Queued or running) and how long it took. Filter by system and status; **Show archived runs** brings back archived ones.
- **Jobs** table: every background job (extraction, analysis, configuration load, configuration sync, upload, simulation, process discovery) with its object, system, state, progress, start and duration.

The job line and the job inspector

A running job shows as a thin progress line under the top bar, never a tray. Open it, or any job row, for the job inspector: its stages and tables with row counts, the percent, the error in plain words for a failed job (for example "OData refused credentials (HTTP 401). Check the client id and secret in the system settings.") and, where the job recorded how it started, **Retry**.

The numbers and where they come from

- **Records** and **Rules** are the records checked and the rules measured in that run, across every object.
- **Score** is the lowest and highest shown score of the run's objects.
- **Took** is from the job registry: extraction and analysis timings are recorded by the workers.
- The baseline is the run you pinned; every comparison and "since the baseline" figure measures against it.

What you can do, by role

- **SAP Basis and analyst** (trigger sync, analyse): start an extraction or an analysis, retry a failed job.
- **Anyone with analyse**: start a run ("Starting a run needs the analyse permission"), pin a baseline, archive older runs.
- **Anyone with export**: download a run's report.
- **Everyone**: read runs and jobs, open a run, compare two runs.

Starting runs, schedules and retries for one system at a time also live under Administration, **Runs and jobs**: "Start an extraction, an analysis or a configuration load, schedule them, and retry what failed."

Keyboard

- `g` then `r` opens Runs.
- `j` and `k` move through runs or jobs, `Enter` opens one.
- `Esc` closes the job inspector.

When something is wrong

- **A run reads Failed with no records**: open its job. Connection errors (HTTP 401, RFC logon failures) belong to SAP Basis; the job says which credential or table failed.
- **A run is complete but scores are missing for an object**: the extraction read no records for it, or the analysis skipped its rules. Open the object; its receipt says how many rules produced a result.
- **Retry is missing**: the job did not record how it started; start it again from the system or from Administration, Runs and jobs.
- **The job line shows nothing but a job is running**: the job registry updates every few seconds; reload the page if it does not appear within a minute, then ask your administrator to check Jobs and system update.

Run and Compare

The **run canvas** shows one run: what it scored, object by object, and the jobs that made it. **Compare** puts two runs side by side and explains why each score moved. Open a run from the timeline or the Runs table (`/run/{id}`); open a comparison with **Compare with** on a run or an object's Runs tab (`/run/{earlier}/vs/{later}`).

One run

- **The title**: "Run 16 Oct 2026, 05:17" and the system.
- **The facts line**: the system, records checked, rules measured, scoring time and trigger, for example "ECC production, 400,662 records checked, 5,178 rules measured, scored in 11 m 42 s, trigger scheduled".
- **The sentence**: "15 objects scored: 2 go, 2 at risk, 11 no-go."
- **Actions**: **Show scores from this run** (sets the run in the top bar so every canvas reads it), **Pin as baseline**, **Archive the 8 older runs**, **Download report**, **Compare with**.
- **By object**: "Measured score, then the score shown after the cap". Each object with its measured score, shown score (with **capped**), verdict, records checked, rules and open work.
- **Jobs**: "The extraction and analysis that made this run", with state, start and duration.

Compare

The title names both runs, for example "Run 16 Oct 2026, 05:17 vs baseline 2 Oct 2026". The sentence sums it up: "12 of 15 objects scored higher and 3 lower. 14 rules started failing and 37 stopped. 12 objects are still capped by a critical rule, so the shown score hides part of it."

The method is stated under it, exactly as the product states it: "Rules are compared on the common set: the 5167 rules that ran cleanly in both runs. 11 rules added since the baseline of 2 Oct 2026 are excluded. A rule counts as changed when it flips between failing no records and some, or fails at least 10% and 10 records more or fewer. Coverage is the share of records a rule could evaluate."

- **By object**: measured before and after, the change in points, shown before and after, records checked, coverage, and the counts of rules new, fixed, worse and better. The total row adds them.
- **What changed**: "Rule counts match the table above." Four lists: **Newly failing rules**, **Fixed rules**, **Worse** and **Better**, each rule with its object and its failing records before and after, for example "MM201 Base unit of measure not in the check table, Material master, 0 to 316 records failing".

On an object's Runs tab the same comparison appears as **Why the score moved**, and ends with what the cap did, for example "The shown score moved only 69.8 to 70.0, because the cap holds it at 70.0."

The numbers and where they come from

- Measured and shown scores are as stored with each run; nothing is recomputed when rules change later.
- Only the common set of rules is compared, so a rule added later never shows as "newly failing".
- The change threshold (10% and 10 records) stops noise from counting as a change.

What you can do, by role

- **Everyone**: open runs, compare any two, export the tables.
- **Anyone with analyse** (programme lead, analyst, SAP Basis, admin): pin the baseline, archive older runs.
- **Anyone with export**: download the run's report.

Keyboard

- `j` and `k` move through the tables, `Enter` opens an object or rule.
- `g` then `r` returns to Runs.

When something is wrong

- **The shown score did not move although many records were fixed**: the cap holds it. Read the measured score, and fix the critical rules over tolerance to release it.
- **A comparison says few rules were compared**: one of the runs ran fewer rules cleanly (a failed extraction or a narrower scope). Compare runs of the same scope.
- **Pin as baseline or Archive is missing**: changing a run needs the analyse permission; archived runs stay readable from **Show archived runs** on Runs.
- **"This run could not be read. It may belong to another tenant or have been removed."** Check the link, or open the run from the Runs table.

Systems

A system is one SAP or cloud source Meridian reads from: ECC, S/4HANA on premise or Cloud, Extended Warehouse Management, SuccessFactors, Concur, Ariba or BTP. Meridian reads from SAP and never writes to it. The system canvas is where SAP Basis connects a system, tests it, extracts with a scope, schedules extractions, loads the configuration and reviews custom objects. Open a system from its heading on the landscape or the map (`/s/{id}`); connect a new one from the command line ("Connect a system") or at `/s/new`.

Connect a system

"Meridian reads from SAP and never writes to it. Enter the connection, test it, and connect. Then extract the objects you want checked."

1. Name the system, choose the **System type** and the **Environment** (Production, Quality, Development).
2. Fill in the connection. For ECC and S/4HANA on premise: **Host**, **Client** (3 digits), **Instance number** (2 digits) and an optional **User** ("Leave empty to use the deployment's default RFC user"), then the password under **Sign-in secrets**. Cloud systems ask for their base URL and OAuth details instead.
3. **Test connection**. A passing test reads like "RFC connection succeeded (client 100, 41 ms)".
4. **Connect system**. "You can connect now and test later. Nothing is read from SAP until you extract."

The system canvas

The header gives the health ("Connected"), the type and protocol ("ECC, RFC"), the last extraction and its records, and the last connection test. **Test, load or extract** opens the system inspector; **Connection settings** opens the connection editor. Five views:

- **Overview**: "What the latest run scored", object by object, and **Recent jobs** from the job registry.
- **Extractions**: the stored extractions of this system and whether each was analysed.
- **Configuration**: the configuration Meridian read from the system's customising tables, by area (enterprise structure, financial accounting, controlling, materials management and so on), each with "Loaded, 3 of 3 objects" and what it found, such as "4 company codes, 11 plants". **Load again**, **Discover now**, the **Drift log** and **Compare with target** (the target system chosen in the connection settings, or the SAP standard baseline). A table that could not be read says why, for example "T500P: No authorisation to read this table. The user needs read access to it."
- **Schedules**: "Each schedule extracts the object again and re-runs its checks, so trends and alerts see every run. Times are UTC." Add a schedule with an object, a rhythm (daily, weekly, monthly, hourly or a custom cron) and a mode: **Full extraction** or **Changes only**. "Changes only reads the records SAP changed since the last download, with a full download at least once a week." Pause or delete a schedule from its row.

- **Custom objects**: "Custom (Z) objects found across every connected system in this workspace, from the latest analysis." Facts (custom objects, active anomalies, ghost and dormant, mapped to standard), **Detect now**, and the registry, anomalies, ghost, dormant, drift, rules and mapping views. Open a Z object to map it to its standard equivalent or add a Z rule.

Extract with a scope

In the system inspector, **Extract** asks for the objects and the scope: "Choose objects to set the scope: company codes, plants, organisations and a date window." With no filter the scope reads "Scope: everything, no organisational filter." The same extraction form under Administration, Runs and jobs adds company codes, plants, a document date window and **Analyse straight after extracting**, which scores the new extraction when it lands.

Dry run shows "What an extraction of 3 objects would read, without reading it": the rows per table last time and an estimated duration, in the form "rows in tables, about a duration (this system's recent speed, rows a second)". Where the system has no recent extraction the speed reads "an assumed speed". From there **Extract changes only** starts a delta extraction that reads the change documents since the last download, where the system supports it; otherwise the reason is shown.

Run analysis on ECC production scores a stored extraction against the rules as they stand now. "Earlier scores are kept as they were."

The numbers and where they come from

- **Records** is what the last extraction stored, across every object.
- **Configuration** counts areas loaded out of those the connector knows, for example "11 of 12 areas".
- The dry run counts rows from the last extraction of each table; a table never extracted is not counted and says so.

What you can do, by role

- **SAP Basis** (manage systems, trigger sync): connect, test, rotate credentials, extract, dry run, delta, schedule, load configuration, remove a system.
- **Analyst and programme lead** (analyse, trigger sync where granted): run analysis, extract.
- **Rule maintainers** (manage rules): edit a custom object's registry entry, add a Z rule.
- **Everyone else**: read.

Keyboard

- ⌘K, then "connect", opens Connect a system.
- In every table, `j` and `k` move and `Enter` opens.
- `Esc` closes the inspector.

When something is wrong

- **"Sign-in refused"** or **HTTP 401**: the user or secret is wrong or expired. Open **Connection settings**, type the new values ("Stored secrets are never shown. Type new values to replace them; an empty field keeps the stored one.") and test again.
- **"No authorisation to read this table. The user needs read access to it."** Grant the RFC user read access to that table in SAP (S_TABU_DIS or S_TABU_NAM), then **Load again**.
- **"The table does not exist in this release. Nothing to fix."** Expected on some releases.
- **"This system is extracted but no rule has run on it. Run an analysis."** Use Run analysis.
- **"Testing, loading and extracting need the trigger sync permission on this system."** Ask your administrator for a grant on that system.
- **A schedule did not run**: due runs start within 5 minutes of their time; check Runs, Jobs for a failed extraction.

Golden records

A golden record is the one trusted version of a thing your systems hold more than once, such as a material kept twice or a supplier in two systems. The **Golden records** tab of an object lists the match candidates waiting for review, the golden records already built and the survivorship rules that decide, field by field, which value survives. Open it from the object's tabs, for example `/o/ecc-production/material-master/golden`.

What is on it

- **Match candidates**: "148 pairs to review", best score first, each pair with its match score and both keys, for example "0.94, MATNR=5000 and MATNR=6000", and **Review**. "Showing the first 5 of 148, best score first." **Show 20 more** loads more.
- **Golden records**: "1,284 records", with golden key, name, sources, survivorship, confidence, last merged and open conflicts. The line under the table explains them: "Confidence is the match engine's score for the surviving record. Last merged is when the record last changed. Open conflicts are fields where the sources disagree and no rule has settled it."
- **Survivorship rules**: the rule per field, for example "MARA-MEINS Trusted source, ECC production" or "MARA-MATKL Most recent", with **Edit the rules**.

The golden record inspector

Open a golden record for **Golden fields and how each value was chosen** (the value chosen, the rule that chose it, or "No rule decided this field"), its **Source records**, its open conflicts and its history ("No merge has changed this record yet."). **Choose records to split out** separates a record merged in by mistake.

Review a pair

Review opens the pair: Record A and Record B **Side by side**, "The two records side by side, pick a value per field", with the **Fields that differ** and **Keep the key of** to choose the surviving key. "Each value you pick on a field that differs is kept in the merged record; the others follow the survivor." Merge the pair, or mark it **Not a duplicate** with the reason ("Why these are different materials"): "This pair will not be offered again." Every decision is in the audit log.

Survivorship rules

"A rule decides which record's value survives in the golden record, field by field. A field with no rule keeps the survivor's value. Rules are listed by field." The rule types and what each says:

- **Most recent**: "The value from the record changed last survives."
- **Most complete**: "The longest non-empty value survives."
- **Most frequent**: the value most sources hold survives.
- **Trusted source**: "The value from the first trusted system that has one survives." List the trusted systems in priority order.
- **Chain of systems** and **Manual override**.

Each rule has a weight from 0 to 100 and can be switched off. A rule the miner suggested is marked "suggested by the miner" and takes effect only when a person saves it. **Save rules** stores them; they apply at the next match run.

The numbers and where they come from

- **Match score** and **Confidence** come from the deterministic match engine, from 0 to 1 and as a percentage.
- **Sources** counts the records merged into the golden record.
- **Open conflicts** counts fields where the sources disagree and no rule settled it.

What you can do, by role

- **Data owner** (approve): merge a pair or mark it not a duplicate. "Merging, or marking a pair as not a duplicate, needs the approve permission. A data owner can decide this pair."
- **Data steward** (mdm create and edit): read pairs, edit survivorship rules.
- **Everyone else**: read, export.

Keyboard

- `j` and `k` move through candidates and golden records, `Enter` opens.
- `Esc` closes the inspector.

When something is wrong

- **No match candidates**: matching has not run for this object, or found nothing above its threshold.
- **Many open conflicts**: a field without a survivorship rule. Add one for the fields downstream processes depend on.
- **A golden value differs from SAP**: golden records live in Meridian. SAP changes only through an approved batch or an approved write-back.

Glossary

The glossary is the shared language of an object: one business term per SAP field, with its definition, its owner, its status and the rules that cite it. The **Glossary** tab of an object lists its terms and the fields that still have none. Open it from the object's tabs, for example `/o/ecc-production/material-master/glossary`.

What is on it

- **Search terms** and **Add terms from the rule library**.
- **Business terms**: "6 terms", each with term, definition, SAP field (`MARA-MEINS`), owner, status (**Active** or **Under review**) and rules. For example: "Base unit of measure: The unit in which stock of the material is managed. Every other unit is converted to it."
- The line under the table: "Owner is the owner of the material master object. Rules counts the rules that cite the term."
- **Fields without a term**: the fields read by failing rules that have no term yet, each with the rule that reads it, for example `MBEW.BKLAS` and `MM091`. "The fields are those read by the failing rules of this object in the latest run. Terms are drafted from the rule library with Add terms from the rule library, and a data owner approves each."

The term inspector

Open a term for its definition, **Fields this term governs**, **Rules that cite this term**, its review cycle and when it was last reviewed, and its history ("The change is in the term's history."). **Mark as reviewed** keeps an active term current. "Editing and approving a term needs the approve permission. A data owner can do both." Approving moves a term from **Under review** to **Active**.

The numbers and where they come from

- **Rules** counts the rules that cite the term.
- **Fields without a term** comes from the failing rules of the latest run, so it changes as rules start or stop failing.

What you can do, by role

- **Data owner** (approve): edit, approve and mark terms as reviewed.
- **Data steward**: draft terms with Add terms from the rule library, comment on terms.
- **Everyone**: read, search, comment, export.

Keyboard

- `j` and `k` move, `Enter` opens a term.
- `Esc` closes the inspector.

When something is wrong

- **A term stays Under review**: it waits for the object's owner. People shows who owns the object.
- **No owner shown**: the object has no owner yet. Administration, Stewardship lists the gaps.
- **A field is missing from Fields without a term**: only fields read by failing rules are listed.

Lineage

Lineage draws why a failing value matters: from the source table, to the field, to the rule that fails on it, to the S/4 feature it blocks or degrades. The **Lineage** tab of an object draws it for the failing rules with the highest value at risk. Open it from the object's tabs, for example `/o/ecc-production/material-master/lineage`.

What is on it

- **Four columns**: **Source tables** (`MARA`, `MARC`, `MBEW`), **Fields** (`MARA.MEINS`), **Rules** (`MM014`) and **S/4 features** (Migration cockpit: material load, Material ledger actual costing, Embedded EWM product master and so on).
- **Node facts**: "Hover or focus a node to see its facts here. Click a rule or a feature to open it."
- **Export as SVG**: the drawing as a file for a slide or a document.
- The line under it: "Drawn from the 8 failing rules of Material master with the highest value at risk. Features come from the impact model of the current run."

Related objects

The object header lists the related objects with their relation, **upstream** or **downstream**. An upstream object's failures pass into this one, so the Work ledger and the record inspector say which to fix first, for example "Fix Material master first: 3,640 upstream failures. Records there fail before this object can pass."

The numbers and where they come from

- The rules drawn are the failing rules of this object in the current run, highest value at risk first.
- The features and whether each is blocked or degraded come from the feature impact rules of the run: the same mapping the Readiness tab shows.

What you can do, by role

- **Everyone**: read, open a rule or a feature, export as SVG.

Keyboard

- `Tab` moves through the nodes; the node facts follow the focus. `Enter` opens a rule or a feature.

When something is wrong

- **The drawing is empty**: no rule of this object fails in the current run, or the object has no run.
- **A feature you expect is missing**: no impact rule maps the failing rules to it. Impact rules are shipped with Meridian; ask your administrator.

Processes and the designer

The **Processes** tab of an object shows which business process steps its data blocks: Procure to Pay, Plan to Produce, Order to Cash and the others, down to the step, the transaction and the field. The **process designer** is where the process model is edited. Open the tab from the object's tabs (`/o/{system}/{object}/processes`) and the designer with **Open the designer** (`.../processes/design`).

What is on the Processes tab

- **The sentence**: "2 of 4 process steps are blocked by Material master: PTP step 3, PP step 1."
- **Each process**: its name and code ("Procure to Pay PTP"), its state ("Blocked, 1 of 3 steps blocked"), its description and its steps. Each step names the transaction (`ME51N`, `ME21N`, `MIGO`), its state (**Ready** or **Blocked**) and, when blocked, the activities and fields behind it with status, pass rate and failing records, for example "Enter Quantity and Unit: `MARA.MEINS` Base unit of measure, Blocked, 94.1%, 3,640".
- The rule under it, as the product states it: "A step is blocked when a field it writes fails a critical or high rule, or passes under 70%. Failing counts are records, and a record can fail several rules."
- **What if these rules were fixed** (simulation): "Tick the rules to fix. The steps of Material master are read again without their failures; nothing is changed and no data is read." Each rule shows "3,640 records at least, in 1 step". The result shows which steps would be ready.
- **Process designer** with **Open the designer**.

The step inspector

Open a step for its transaction, its activities and fields, the rules failing on them and **Hand off**, which hands a blocking rule to a steward with a note.

The process designer

"Process designer, Material master", with the model's state, for example "Draft, version 1, updated 12 Oct 2026".

- **Model**: choose a model, **Create a model** (start from the reference model or the group model), **Publish model**, **Rename model**, **Delete model**.
- **The tree**: levels L1 to L4 (for example Procure to pay, Sourcing, Purchase order, Create purchase order with its activities), each with **Up**, **Down** and **Delete**, and **Add a process at the top level**.
- **The model panel**: "Choose an item in the tree to rename it, reorder it, add what sits under it, or map fields to its activities."
- **What changed**, **Save as a new version**, **Discard changes**, and **Versions**.
- The rule at the foot: "A step is blocked in process readiness when a field it reads fails a critical or high rule. Map the fields each activity reads here, and the Processes tab shows them on the next run."

The numbers and where they come from

- Step states come from the fields mapped to each activity and the rules that failed on those fields in the run.
- The what-if simulation is computed in the browser over the process document and the receipt; it reads no data and changes nothing.

What you can do, by role

- **Process analyst**: read readiness, run what-if, hand off a blocking rule (handoff rule), export the as-is process (export process), edit and publish the model.
- **Data steward**: receive handed-off rules in Work.
- **Everyone**: read.

Keyboard

- `j` and `k` move through steps, `Enter` opens a step.
- In the designer tree, `Tab` moves between items.

When something is wrong

- **"Choose at least one rule."** The simulation needs a ticked rule.
- **Every step is Ready but the object is No-go**: the failing rules are on fields no activity reads. Map them in the designer.
- **Your edits do not show on the Processes tab**: they show on the next run, and only once the model is published.

Reports and schedules

Reports are files built from a stored run: the board pack, the analyst workbook, the analysis report, the trend report and the rule catalogue. Schedules build a report from the latest run and send it on a rhythm. Open Reports from the command line (type "reports") or at `/reports`.

What is on it

- **The sentence**: "6 reports made, 2 of 3 schedules are on. Reports are built from stored runs, so a number in a report matches the run it names."
- **Create a report** and **Schedule a report**.
- **Reports**: report, run ("Run of 16 Oct 2026"), format (PDF or XLSX), created, by, size and **Download**. The rule catalogue is "Not tied to a run".
- **Schedules**: "Each builds a report from the latest run and sends it." Report and format, cadence ("Every Monday at 07:00 (Africa/Johannesburg)"), recipients, next run and last result (**Delivered**, **Partly delivered**, **Failed** or "Not run yet"), and **Paused** for a paused schedule.

Create a report

Choose the report, the run it is built from ("Built from the run you choose") and, optionally, one object. The report is built by a worker and appears in the list when it is ready.

The schedule inspector

Open a schedule for its facts (next run, last run and its result), then: name, report (Board pack (PDF), Analyst workbook (XLSX), Analysis report (PDF), Trend report (PDF), Rule catalogue (PDF)), **Which run** ("The latest finished run", "The latest file upload" or the latest run of one system), an optional object, the cadence (every day, week or month, the day and the time in the tenant's time zone), the recipients (people of the tenant and the alert channels that take reports) and **Schedule is on**. Actions: **Save the schedule**, **Pause**, **Run it now** ("It is built on the worker and delivered to the recipients; the next scheduled run is unchanged.") and **Delete** ("No more reports are built or sent from this schedule. Files it already made stay in the list.").

The numbers and where they come from

Every figure in a report is read from the run it names, so the board pack of 16 Oct says what the landscape said for that run. Money is in the tenant currency from the run's cost model. Alerts and channel messages carry counts, rule ids and links, never record values.

What you can do, by role

- **Programme lead, analyst, admin** (export): create, schedule, run now, pause, delete.
- **Everyone with view**: download reports. "Creating and scheduling reports needs the export permission."
- **Administrator**: add the channels that can receive reports under Administration, Alerts.

Keyboard

- `j` and `k` move through reports and schedules, `Enter` opens a schedule.
- `Esc` closes the inspector.

When something is wrong

- **"No report has been made yet. Create one from a run, or schedule one to arrive on its own."**
- **"No one can receive reports yet."** Invite people, or add an email or Teams channel under Administration, Alerts.
- **Partly delivered**: one recipient or channel refused it. Check the channel under Administration, Alerts; email needs an email backend, which Administration, Jobs and system update reports under Health checks.
- **Failed**: open the schedule and choose **Run it now** after fixing the channel; the job appears under Runs, Jobs.

Waves and gates

A migration wave is a set of objects moved together in one cutover, with a target date, a stage and a gate: the conditions every object must meet before the wave may go. The landscape's **Waves** view shows each wave with its verdict counts and gate state; the **Migration waves** panel is where the programme lead plans waves and sets gates. Open the view with **Waves** on the landscape (`/?view=waves`) and the panel with **Edit** on a wave (`?open=waves`).

The Waves view

Each wave heading reads, for example, "Wave 1: finance foundations, due 16 Nov 2026, 1 go, 0 at risk, 3 no-go, Gate not passed, 2 of 3 not met", with **Edit**. Under it, the wave's objects with their ruler, score, verdict, records and rules. The programme lead's next-step button follows the nearest wave: "Clear the 3 no-go objects in Wave 1: finance foundations" with "Wave 1: finance foundations: due 16 Nov 2026, in 17 days. 1 go, 0 at risk, 3 no-go."

An object's own **Readiness** tab lists its waves with the verdict, blocking gaps, score, readiness and records blocked, and the wave threshold.

The Migration waves panel

For each wave: **Name**, **Target date**, **Stage** (Plan, Mock load 1, Mock load 2, Dress rehearsal, Cutover) and its sign-off state ("Not signed off.").

Gate: up to three conditions, each with what fails it now:

- **Min score**: "Every object at 85.0 or more", with the objects below it.
- **Zero criticals**: "No critical rule over tolerance", with the objects capped by a critical rule.
- **Contracts met**: "Data contracts met", or "No active contract names these objects."

Save the gate stores it. "No gate yet: a wave without one moves on its verdict alone."

Readiness certificate: "not yet (the signed certificate arrives in step 9)". The link shows the certificate's current checks; a signed PDF is not produced in this version.

Add a wave creates one; **Objects** assigns each object to a wave or to none; **Save changes** or **Discard**. **Delete** removes a wave.

The numbers and where they come from

- **Go, at risk, no-go** counts are the verdicts of the wave's objects in the run in the top bar, from the score bands.
- **Gate state** checks each condition against the same run: the shown scores, the critical rules over tolerance, and the active contracts that name the objects.
- **Due in N days** counts from today to the target date.

What you can do, by role

- **Programme lead** (analyse): plan waves, set gates, assign objects.

- **Data owner**: read the gate state of their objects' waves.
- **Everyone**: read the Waves view.
- **Sign off** (approve): on each wave in the Migration waves panel, once its verdict is go. The panel shows who signed it off and when; any later change to the wave clears the sign-off.

Keyboard

- `g` then `l`, then the **Waves** view switch.
- `Esc` closes the panel.

When something is wrong

- **A wave shows every object no-go but scores look close to target**: the cap holds them. The gate's **Zero criticals** line names the objects; clear their critical rules first.
- **Contracts met says no contract names the objects**: create one in People, Data contracts, or remove the condition.
- **The certificate says "not yet"**: expected in this version.

Administration and settings

Administration

Administration is everything that sets up the tenant: who can sign in, what the licence covers, how scoring and alerts behave, and the state of the platform. It opens from your account menu, **Administration**, for anyone who holds an administrative permission, at `/admin`. "Everything that sets up this tenant: who can sign in, what the licence covers, how scoring and alerts behave, and the state of the platform. Each section opens in this column." You see only the sections your role opens; the rest say which permission they need.

People and access

- **Users**: "7 of 25 seats. A person's roles are grants: open a row to add or remove them." **Invite a user**, **Import users**, filters by role and status (Invited, Invite expired, Active, Locked, Suspended). A banner warns about invitations older than 7 days. Opening a person shows their grants, what they can do now and every action on the account. Needs manage users.
- **Roles and permissions**: "Built-in roles cannot be changed. Copy one to start a custom role, then set exactly what it may do. Give a role for one system or object only, from a person's grants." The matrix "What each role may do", **Access requests** (grant, or decline with a reason the person reads) and recent changes.
- **Single sign-on**: OpenID Connect with your identity provider (Microsoft Entra ID, Okta). Issuer address, client ID, client secret ("Stored encrypted and never shown again."), button name, the redirect address to register, **Test sign-in** and **Turn off single sign-on**. "People who sign in for the first time get the Viewer role and no access grants; give them access under Users."
- **Sign-in page**: the button name, allowed email domains, **Require single sign-on** ("Password sign-in stops for everyone except administrators, who keep it as a way in if the provider is down."), and the fixed password rules: "At least 12 characters, not the email name, and not a common password"; "10 failed sign-ins lock the account for 15 minutes".
- **View as a role**: "See the product as one role sees it, to answer a support question or check a role before you grant it. Everything is read-only while you view as a role, and starting and stopping are in the audit log." The top bar shows "Viewing as" until you stop.
- **Share links**: "Read-only links to one page, opened as a viewer by people of this tenant until they expire. Revoke one and it stops working at once."
- **Out of office**: "Who is away and who covers their approvals and assigned work. Each person sets their own from Preferences; an administrator can end one."
- **API keys**: "A key acts as its own service account with only the roles you give it, and shows in the audit log under its name. The key itself is shown once, when you create it."

Tenant

- **Licence**: company, status, tier, seats, validity, last checked, platform version and the licensed modules and features. **Check with Meridian HQ now**. "Plan changes and invoices are handled in Meridian HQ; this deployment keeps its data either way."
- **Billing**: resolved exceptions by billing tier on top of the base fee; read only.

- **Settings**: the tenant; **Dimension weights** ("The score is each measured dimension times its weight. The weights must add up to 100." **Use the DAMA defaults**); **Score bands** ("At or above target is on target (go), at or above warning is at risk, below that is failing."); **Cap rule and tolerance**: "One critical rule failing beyond its tolerance holds that object at 85; two or more hold it at 70. The caps are fixed. The tolerance says how many failing records count as beyond.", with **Failing share, percent** and **Failing records** and the severity weights and multipliers; **Cost model and currency**; **Planner assumptions**; **Retention** (keep runs, reports and audit rows for so many days, keep the baseline; purges rehearse first and the purge log records each one); **Delete all data**, which issues a token that works after 24 hours.
- **Stewardship**: "One accountable owner per object and one steward per company code, plant or organisation." The template per object (owner role, owner, stewards by slice, steward, teams), **Save the template**, **Check what applying does** (a dry-run diff), **Apply**, and **Gaps**: "Every object needs one owner and at least one steward per active slice", for example "Vendor master (AP): No steward for purchasing organisation 1000".
- **Configuration bundle**: "Rules switched on or off, field mappings, waves and their gates, settings, SLA policies and the escalation chain as one JSON file." **Download the bundle**; import with a diff first: "Import adds and changes; it never deletes what the file leaves out. Check the diff first."

Data

- **Field mappings**: "Name the column in your files that holds each standard SAP field, then save the row." **Reset all to defaults**.
- **Reference data**: postal code and bank code (SWIFT BIC) lists per system. "Until a list is loaded those rules check the format only." An upload replaces the system's current list.
- **Rules import and export**: "Every rule of the tenant as YAML or Excel, and many rules at once back in. Imported rules are validated against the data dictionary and land as drafts for review, never live." **Check the file**, then **Create drafts**.
- **Audit log**: every change, who made it, the object and the result, filters by person, action, object kind and dates, **Export CSV** and **Export Excel**, and **AI calls**: "Every call to the language model: which service, how long, and whether a deterministic rule answered first. Prompts and records are never stored here."

Notifications

- **Alerts**: channels (email, Microsoft Teams, signed webhook) with their digest, board pack and critical alerts; **When to alert** (critical and high failing rules, score drop in points) and **Score floors per object**. "Alerts carry counts, rule ids and links, never record values."
- **Notifications**: the tenant's default digest. Each person sets their own under Preferences.

Operations

- **Runs and jobs**: start an extraction (objects, company codes, plants, document dates, analyse straight after), start an analysis of a stored extraction, load the configuration, scheduled runs, **Upload a file** (a CSV, TSV, Excel, JSON or Parquet file of one object becomes a run and is

analysed at once; its columns are matched by the field mappings; needs upload) and recent jobs with **Retry**.

- **Jobs and system update**: jobs by state and kind, the platform version and **Update to** the latest when Meridian HQ knows one, and **Health checks** (database, cache and queue, background workers, licence, email delivery, file storage): "Each check says what it found and what needs fixing. They refresh every 30 seconds."

Provisioning, theming and compliance

These are set through the API in this version; there is no screen for them yet:

- **SCIM provisioning**: SCIM 2.0 under `/scim/v2`, authenticated by a per-tenant token issued, rotated and revoked at `POST /api/v1/admin/scim/token`. Map identity-provider groups to roles at `GET/PUT /api/v1/admin/scim/group-roles`. Only grants SCIM made are removed, and the last active administrator is never left without admin.
- **Theming**: `GET/PUT /api/v1/settings/theme` sets the tenant's colours and logo (PNG or SVG), with contrast checks.
- **Compliance**: `GET/PUT /api/v1/settings/compliance` records data residency, retention days per table family, encryption flags and the last backup.

The numbers and where they come from

- **Seats** count active and invited users against the licence.
- **Settings** hold the numbers every score uses: the dimension weights, the score bands, the tolerance (failing share and failing records) and the severity weights. A change applies to the next analysis; scores already stored with a run are kept as they were.
- **Health checks** are measured by the platform itself and refresh every 30 seconds.
- **Billing** counts resolved exceptions by tier for the month, on top of the base fee, from Meridian HQ.

What you can do, by role

- **Platform admin**: every section.
- **Manage users**: users, roles, single sign-on, sign-in page, view as, share links, out of office, API keys.
- **Manage settings**: settings, billing, stewardship, configuration bundle, alerts, notifications.
- **Manage system**: licence, jobs and system update.
- **View audit** (auditor): the audit log, and a person's activity in People.
- **Manage field mappings**, **manage systems**, **manage rules**: field mappings, reference data, rules import and export.
- **Trigger sync** or **analyse** (SAP Basis, analyst, programme lead): runs and jobs.
- Anyone who holds view reads the licence, settings, mappings, reference data, alerts and notifications, read only.

Keyboard

- In every section table, `j` and `k` move and `Enter` opens the row in the inspector.
- `Esc` closes the inspector; the section index stays on the left.

When something is wrong

- **"... needs the ... permission. Ask an administrator to grant it."** Your role does not open that section.
- **"These settings are read only for your role."** You can read but not change them.
- **Health check warning "No email backend configured"**: digests and email alerts are not sent until the deployment's email settings are set.
- **Single sign-on test fails**: check the redirect address registered at the identity provider matches the one shown, character for character.
- **A purge request**: purges are irreversible. The token only works after 24 hours so a mistake can be caught; the purge log keeps the record.

Preferences

Preferences is how Meridian looks and reads for you, where it starts, and what it tells you about. It belongs to you, not to the tenant. Open it from your account menu, **Preferences**, or with `?open=preferences`. Every change is saved at once.

What is in it

- **Display**: **Theme** (Light, Dark or System) and **Row density** (Compact, Default or Comfortable).
- **Language and region**: **Number and date format** (Browser default or a locale such as en-ZA or de-DE) and **Time zone** (UTC by default). A sample line shows the result: "Dates read 16 Oct 2026, 05:17, scores 71.8, counts 1,234,567."
- **Start screen**: **Open Meridian as** your role's default or another role's start screen. "It changes the next step on the landscape and the work it shows first. Your access does not change."
- **Keyboard**: "Show keyboard hints beside menu items and places".
- **Notifications**: the digest of your open items (Off, Daily, Weekly), "Tell me at once about critical findings" and "Also email me". **Inbox and triage settings** opens triage.
- **What you hear about**: the events that reach your inbox, such as run finished, score drop, new critical breach, batch waiting and mentions, each on or off.
- **Out of office**: set from People, Out of office, or `?open=away`.

The numbers and where they come from

Nothing here changes a number. The locale changes how numbers and dates are written; the time zone changes the times shown; scores and counts stay the same.

What you can do, by role

Everyone sets their own preferences. The tenant's default digest and alert channels are set by an administrator under Administration, Notifications and Alerts.

Keyboard

- `Tab` moves through the settings, `Space` toggles a switch.
- `Esc` closes the panel.

When something is wrong

- **Dates show in UTC**: set your time zone here.
- **No email arrives**: "Also email me" is off, or the deployment has no email backend (Administration, Jobs and system update, Health checks).
- **The start screen did not change**: the start screen is kept in this browser; another browser keeps its own.

Keyboard and the command line

Meridian can be driven from the keyboard: a whole ledger can be worked without the mouse, and every canvas and action is reachable by typing. `?` shows the full sheet at any time; the keys are the same on every screen.

The command line

Open it with `⌘K` (Ctrl+K on Windows and Linux) or `/`. "Search or type a command." Type:

- a **record key** such as a material number to open the record;
- a **rule id** such as `MM014` to open the rule;
- an **object and system**, such as "material master ECC", to open the object;
- a **run** date to open the run;
- a word: **approve** opens the batches waiting for approval; **export** saves the fix sheet of what is on screen.

With nothing typed it offers **Go to**: Landscape, Work, Runs, People, Rules and Connect a system. `↑` and `↓` move through the results, `Enter` opens one, `Esc` closes.

Everywhere

- `⌘K` or `/`: open the command line.
- `?`: show the shortcut sheet.
- `⌘Z`: undo the last change (when focus is not in a text field).
- `Esc`: close the top layer: menu, then inspector, then the selection.

Go to

- `g` then `l`: Landscape.
- `g` then `w`: Work.
- `g` then `r`: Runs.
- `g` then `p`: People.

Tables and lists

- `j` or `↓`: next row. `k` or `↑`: previous row.
- `Shift+J` or `Shift+↓`, `Shift+K` or `Shift+↑`: extend the selection.
- `Enter`: open the row in the inspector.
- `x`: select or clear the row.

Work ledger

- `a`: accept the suggested value on the selected rows.

- `n`: mark the selected rows as not an issue.
- `r`: accept the risk on the selected rows.
- `e`: export the fix sheet of what is on screen.
- `f`: move to the filter chips.
- `⌘D`: fill the New value of the first selected row down the selection.
- Paste several lines into a New value cell to fill the selected rows.

Inspector

- `]`: open the next row. `[`: open the previous row.
- `Tab` cycles inside the inspector.
- `Esc` closes it and returns you to where you were.

Where the numbers come from

Nothing typed in the command line is guessed: results come from the tenant's own records, rules, objects and runs, grouped by kind.

What you can do, by role

Every role can use every key. A key that acts (`a`, `n`, `r`) does nothing visible when your role does not hold the action; the ledger says "Read-only for your role".

When something is wrong

- **A key does nothing**: focus is in a text field. Press `Esc` or click the page, then try again. `⌘K` and `Esc` work everywhere.
- **"Select rows first: x selects the highlighted row, Shift+J extends the selection."** A ledger key was pressed with no row selected.
- **The command line finds nothing**: try a shorter word, or the record key as SAP stores it, with its leading zeros.
- **Keyboard hints are missing beside places**: turn them on under Preferences, Keyboard.

Walkthroughs by role

Walkthrough: programme lead

You run the data programme: you need to know how ready each system and object is for the next wave, what holds it back, who is working on it, and what to tell the steering committee. Your start screen is the **Landscape** with the next-step button for your nearest wave.

Your first five steps

1. **Landscape**: read the sentence and press the next step, for example "Clear the 3 no-go objects in Wave 1: finance foundations".
2. **Waves**: switch the landscape to **Waves**, open **Edit** on the wave and set its gate: min score, zero criticals, contracts met. **Save the gate**.
3. **An object**: open a no-go object and read its receipt: measured, capped, by which critical rules.
4. **People**: see who holds the open work and **Balance work** if one person holds most of it.
5. **Reports**: schedule the board pack for the steering committee.

Your jobs and where they are done

- **Landscape readiness by system and wave**: Landscape, Coverage and Waves views.
- **Plan waves (objects, dates, stages)**: Migration waves panel, **Add a wave**, **Objects**.
- **Wave gate criteria**: Migration waves panel, **Gate**, **Save the gate**.
- **Gate status**: the Waves view, "Gate not passed, 2 of 3 not met".
- **Readiness certificate**: the certificate link shows "not yet (the signed certificate arrives in step 9)"; a signed PDF is not produced in this version.
- **Sign off a wave**: **Sign off** on the wave in the Migration waves panel (needs approve); only a wave whose verdict is go can be signed off, and any later change to the wave clears the sign-off.
- **Balance work between stewards**: People, **Balance work** (even counts, whole objects or by the assignment template), applied with undo.
- **Read data contracts and their breaches**: People, **Data contracts**.
- **Watch objects**: **Watch** on the object header for run, score and breach notices.
- **Export the ledger or any grid**: **Share and export**, **Export to Excel**.
- **Share a canvas with the steering committee**: **Share and export**, **Share this page as a link** (read-only, expires).
- **Start extraction and analysis, schedules**: System inspector; Administration, Runs and jobs.
- **Reports and report schedules**: Reports.

The numbers you report

- **The shown score** is capped: 85 at one critical rule over tolerance, 70 at two or more. Report the measured score beside it when it matters; the receipt states both.

- **Value at risk** is from the cost model named in the footer; it is an estimate from your own cost settings, not a forecast.
- **Go, at risk, no-go** come from the score bands under Administration, Settings.

When something is wrong

- **The steering committee sees different numbers**: check which run each of you reads. The run is in the top bar and in every report.
- **A wave cannot pass its gate**: open each object named under **Zero criticals**; the critical rules over tolerance are listed in its receipt.
- **Contact**: SAP Basis for extraction failures, your administrator for access, the object owner for approvals that wait.

Walkthrough: data steward

You clean the data: you decide the right value for each failing field, accept what Meridian suggests where it is right, sign off what is not a problem, and hand your batch to a second person. Your start screen is the **Landscape** with "Accept the suggested values" as your next step.

Your first five steps

1. **Landscape**: press **Accept the suggested values**; Work opens on your rows.
2. **Work**: read the top rows, ranked by value at risk. Select rows and press `a` to accept their suggested values; the notice has **Undo**.
3. **A record**: press `Enter` on a row without a suggestion. Read the possible values with their meanings, type the New value and press `Enter`.
4. **By rule**: switch the ledger to **By rule** and decide a whole rule at once, or **Export fix sheet** (`e`) to work offline and **Import filled sheet** when done.
5. **Waiting for approval**: check your draft batch under the chip; a data owner approves it.

Your jobs and where they are done

- **Work the ledger: fix, propose, not an issue, accept risk**: Work ledger and record inspector.
- **Accept suggested values in bulk**: **Accept all** on a rule, or the next-step button.
- **Build a fix batch for approval**: every accepted or entered value goes into your draft batch.
- **Undo a decision**: the notice's **Undo**, or `⌘Z`.
- **Comment and mention**: Inspector, **Comments**; type `@` and a name.
- **Watch a rule or record**: Inspector, **Watch**.
- **Fix the upstream object first**: the fix-order hint, for example "Fix Material master first: 3,640 upstream failures. Records there fail before this object can pass."
- **Out of office with a deputy**: Preferences or People, **Out of office**.
- **Saved and shared views of the ledger**: **Views**, **Save current view...**, shared with everyone if you choose.
- **Balance work in the team**: People, **Balance work** (needs assign).
- **Author rules**: Rules, **New rule draft** (needs manage rules).
- **Import rules in bulk**: Administration, Rules import and export; there is no link from the rules library yet.
- **See your own activity**: your row in People, **Activity**.

The numbers you work with

- **Value at risk** ranks the rows; it is the cost-model estimate of what each failing value costs.
- **Possible values** come from the rule; a value outside them still fails and is refused.

- **Suggested value** is the rule's proposal, "Blank when no single value is right."

When something is wrong

- **The ledger says "Read-only for your role"**: you need the apply permission. Ask your administrator.
- **A batch comes back**: the approver sent it back with a reason; it is in the batch inspector.
- **A fixed record returns to To fix**: the next run still found it failing after it was sent to SAP. Check with whoever loads the file.
- **Contact**: the object's owner for approvals, SAP Basis for extraction gaps.

Walkthrough: data owner

You are accountable for one or more objects and you are the second pair of eyes: you approve or send back the fix batches stewards prepare, sign off exceptions, approve contracts and glossary terms, and decide who covers for you when you are away. Your start screen is the **Landscape** with the line "2 batches wait for your approval, SLA met on 92% of your open work" and the next step "Approve batch B-241".

Your first five steps

1. **Landscape**: read the line under the sentence and press **Approve batch B-241**.
2. **Batch inspector**: read the facts, the changes ("`MARA.MEINS` STK to EA") and the approval trail. **Approve** or **Send back** with a reason.
3. **Waiting for approval**: in Work, the chip lists every batch waiting, grouped by batch.
4. **Your object**: open an object you own; **Watch** it and apply the steward template to it.
5. **Out of office**: before you are away, set who covers your approvals.

Your jobs and where they are done

- **Approve or send back a batch (four eyes)**: Batch inspector, **Approve**, **Send back**. You never approve a batch you built.
- **Out of office: approvals route to the deputy**: Preferences or People, **Out of office**; the batch and person inspectors show the cover.
- **Escalation chain after an SLA breach**: Work or People, **Triage and SLA**, **Escalation chain** (changing it needs manage triage).
- **Create and edit a data contract**: People, **Data contracts**, **New contract**; the contract inspector.
- **Approve a contract (never your own)**: Contract inspector.
- **Contract versions and revisions**: Contract inspector, **Versions**, **Revise as a new version**.
- **Contract breaches**: Contract inspector, **Breaches**.
- **Apply the steward template to your objects**: the object canvas, steward template.
- **Watch your objects; subscriptions**: Object canvas, **Watch**; Preferences, What you hear about.
- **Gate status of your objects' wave**: Landscape, Waves view.
- **Sign records off**: Work, **not an issue** or **accept risk** with a reason.
- **Merge pairs and approve glossary terms**: an object's Golden records and Glossary tabs.

The numbers you approve

- **Records** and **Value at risk** on a batch are the records in it and their cost-model estimate.
- **SLA met** is the share of your open work resolved within the SLA days set under Triage and SLA.

When something is wrong

- **Approve is missing**: you built the batch, or your role lacks approve.
- **Approved values do not appear in SAP**: approval does not touch SAP. The batch is exported and loaded by your team's process; the next run confirms each record.
- **Contact**: the steward named on the batch for its content, your administrator for access.

Walkthrough: SAP Basis

You keep the plumbing healthy: connections, extractions, schedules, configuration loads and the jobs that run them. Meridian reads from SAP and never writes to it. Your start screen is the **Landscape** with the system that needs you first as the next step, for example "Extract SuccessFactors", and **Your next work** as a list of system tasks.

Your first five steps

1. **Landscape**: read the task list: "Extract ECC production again", "Run analysis on S/4HANA sandbox", "Extract SuccessFactors".
2. **System**: open the system and **Test connection**. A passing test reads "RFC connection succeeded (client 100, 41 ms)".
3. **Dry run**: in the system inspector choose the objects, run **Dry run** and read the rows and duration; then **Extract**, or **Extract changes only** where available.
4. **Schedules**: add a schedule with a rhythm and the mode **Changes only** for routine runs.
5. **Runs**: watch the job line and open a failed job for its stages and **Retry**.

Your jobs and where they are done

- **Test a connection, health**: System inspector, **Test connection**.
- **Extract objects**: System inspector, **Extract**.
- **Extraction dry run: row counts per table and a duration estimate**: System inspector, **Dry run**.
- **Delta extraction (changes only)**: System inspector, **Extract changes only**; schedules in **Changes only** mode.
- **Schedules, retry, cancel**: the system's Schedules view; Administration, Runs and jobs.
- **Field mappings**: Administration, Field mappings.
- **Load configuration and compare with the target**: the system's Configuration view.
- **Custom objects**: the system's Custom objects view.
- **Watch a system**: possible through the API only; the System inspector has no Watch button yet.

The numbers you watch

- **Records** per extraction and per table, from the job registry.
- **Areas loaded**, for example "11 of 12", and per area what could not be read and why.
- **Health checks** under Administration, Jobs and system update.

When something is wrong

- **HTTP 401 or "Sign-in refused"**: rotate the credentials in Connection settings and test again.
- **"No authorisation to read this table. The user needs read access to it."** Grant the RFC user read access, then **Load again**.

- **A scheduled extraction did not start**: due runs start within 5 minutes; check Runs, Jobs and the background workers health check.
- **Contact**: your administrator for permissions on a system; the programme lead for which objects to extract.

Walkthrough: process analyst

You read the data through the business process: which steps of Procure to Pay, Plan to Produce or Order to Cash the data blocks, what fixing a rule would release, and whom to hand a blocking rule to. Your start screen is the **Landscape** with "Open PTP step 3" as your next step.

Your first five steps

1. **Landscape**: press **Open PTP step 3**.
2. **Processes tab**: read "2 of 4 process steps are blocked by Material master: PTP step 3, PP step 1." and the fields behind each blocked step.
3. **What if**: tick the rules to fix under **What if these rules were fixed** and see which steps would be ready.
4. **Step inspector**: **Hand off** the blocking rule to a steward with a note.
5. **Process designer**: map the fields each activity reads, then publish the model.

Your jobs and where they are done

- **Process readiness by step**: Object, **Processes** tab.
- **What if these rules were fixed (simulation)**: Processes tab; computed in the browser, nothing changes and no data is read.
- **Hand a blocking rule to a steward**: Step inspector, **Hand off** (needs handoff rule).
- **Export the as-is process**: Processes and the designer (needs export process).
- **Comment on a rule or a step's rule**: Rule inspector, **Comments**.
- **Related objects and fix order**: Object header, related objects; the fix-order hint in Work.

The numbers you read

- "A step is blocked when a field it writes fails a critical or high rule, or passes under 70%."
- Failing counts are records, and a record can fail several rules.

When something is wrong

- **A step you expect blocked is Ready**: its fields are not mapped to the activity. Map them in the designer; it shows on the next run.
- **Hand off is missing**: your role lacks handoff rule.
- **Contact**: the object's owner for priorities, the steward the rule was handed to for progress.

Walkthrough: auditor and viewer

You read: every canvas, every number and its receipt, and, as an auditor, the audit log and each person's activity. Nothing you do changes data. Your start screen is the **Landscape**, the same one the programme lead sees, read-only. A viewer may also arrive through a share link, which opens one page read-only with the bar "Shared view. Read-only, as a viewer."

Your first five steps

1. **Landscape**: read the sentence and the coverage matrix; hover a cell for its reading.
2. **An object**: open the receipt and every **i**: the arithmetic behind each figure.
3. **Approvals**: in Work, **Waiting for approval** and **Exceptions** show the four-eyes evidence: who drafted, who approved, who signed off and why.
4. **Audit log** (auditor): Administration, **Audit log**; filter by person and dates and **Export Excel**.
5. **Person activity** (auditor): People, open a person, **Activity**.

Your jobs and where they are done

- **Read every canvas**: all canvases, read-only.
- **Audit log and export (auditor)**: Administration, Audit log.
- **A person's audit trail (auditor)**: Person inspector, **Activity**.
- **Open a read-only share link**: `/share/{token}`; the bar says "Shared view", and **Leave the shared view** returns you.
- **Export a grid (auditor)**: **Share and export**, **Export to Excel**.
- **Read comments**: Inspector, **Comments**.
- **Rule change trail joining versions and runs**: not joined yet; read the rule's **Versions** and **History** side by side.

The numbers you check

- Every object score has a receipt: dimension points, measured, cap and shown. The formula line states the weights, the severity weights, the tolerance and the cap rule.
- Comparisons state their method: the common set of rules and the change threshold.
- The audit log's **AI calls** list every call to the language model, which service made it and whether a deterministic rule answered first; prompts and records are never stored there.

When something is wrong

- **A share link says it expired or was revoked**: ask the person who shared it for a new one.
- **The audit log is missing**: it needs the view audit permission.
- **Contact**: your administrator for access; the programme lead for the run a report was built on.

Walkthrough: administrator

You set the tenant up and keep it right: who signs in and with which roles, how scoring and alerts behave, the licence, the platform version and its health, and the configuration you move between tenants. Your start screen is the **Landscape**; **Administration** is in your account menu.

Your first five steps

1. **Users**: invite people and give each a role; scope a role to one system or object from their grants.
2. **Single sign-on and sign-in page**: connect your identity provider, test it, then decide whether to require it.
3. **Settings**: check the dimension weights, score bands, cap rule and tolerance, cost model and retention.
4. **Stewardship**: name an owner per object and stewards per slice, **Check what applying does**, then **Apply**; close the **Gaps**.
5. **Jobs and system update**: read the health checks and update the platform when a new version is available.

Your jobs and where they are done

- **Invite users, set roles**: Administration, Users.
- **Custom roles and scoped grants**: Administration, Roles and permissions.
- **Single sign-on**: Administration, Single sign-on and Sign-in page.
- **API keys**: Administration, API keys.
- **Audit log and export**: Administration, Audit log.
- **Export and import the tenant configuration with a dry-run diff**: Administration, Configuration bundle.
- **View the product as a role (support)**: Administration, View as a role; the bar shows "Viewing as".
- **Import and export rules in bulk (YAML and Excel)**: Administration, Rules import and export.
- **Govern share links**: Administration, Share links, **Revoke**.
- **See everyone's out of office**: Administration, Out of office, **End it**.
- **Steward assignment template, apply with a dry-run diff, audit the gaps**: Administration, Stewardship.
- **Licence, platform update, health**: Administration, Licence; Jobs and system update.
- **Tenant settings (scoring, cost, retention)**: Administration, Settings.
- **SCIM provisioning, theming, compliance**: through the API in this version (see the Administration chapter).

The numbers you set

- Dimension weights must add up to 100.
- The caps are fixed at 85 and 70; you set the tolerance: the failing share and the failing records that count as beyond.
- Score bands: target, warning and critical.

When something is wrong

- **Someone cannot see a section**: compare their role with the matrix under Roles and permissions, or use View as a role.
- **Email is not delivered**: Health checks, Email delivery.
- **The licence shows fewer modules than expected**: modules come from the licence in Meridian HQ; **Check with Meridian HQ now**, then contact Meridian support.
- **Support**: open help, copy the diagnostics under **Contact support** and send it with your request. It holds no record values and no secrets.