Meridian

User guide

Meridian guide

Every screen, what its numbers mean and where they come from, what each role can do, the keyboard, and what to do when something is wrong.

SAP data quality, measured

Edition of 12 November 2026

Contents

The screens

- Meridian in three minutes
- The menu
- Objects
- Landscape
- Ask Meridian
- An object and its tabs
- The rule canvas and the rules library
- The Work ledger
- Approvals and exceptions
- Post-clean monitoring
- People and inbox
- Runs
- Run and Compare
- Systems
- Golden records
- Glossary
- Lineage
- Data catalogue
- Processes and the designer
- Reports and schedules
- Migration
- Waves and gates

Administration and settings

- Administration
- Preferences
- Keyboard and the command line

Walkthroughs by role

- Walkthrough: programme lead
- Walkthrough: data steward
- Walkthrough: data owner
- Walkthrough: SAP Basis
- Walkthrough: process analyst
- Walkthrough: auditor and viewer
- Walkthrough: administrator

The screens

Meridian in three minutes

Meridian checks your SAP data against fixed rules, scores every object, lists every record that fails and routes each fix through a second person before it goes back to SAP. It puts all of that on three surfaces: the **rail** on the left, the **canvas** in the middle and the **inspector** on the right. Every screen is a place in the rail's menu, every piece of work is a row in the ledger, and anything you select opens in the same inspector.

The three surfaces

- **The rail** is a narrow strip of seven places: **Home**, **Objects**, **Runs**, **Fix**, **Rules**, **Master data** and **Migration**. Click one to open its landing page. Point at one, or tab to it, and a flyout opens beside it with that place's pages; pin the flyout to keep it open. **Fix** holds the Action centre, Work, Aletheia verified fixes, Cleansing, Requests and People. **Rules** holds the rules, rule promotion, the rule dry run, anomaly rules, learned monitors, verified survivorship and the evidence ledger. **Master data** holds MDM health, the entry gate, reference data, record links, Material 360, issues, data mining and process discovery. **Migration** holds the migration analysis, S/4HANA readiness, go-live, forecast, value, the score lens, cost and causes, and Configuration. **Objects** lists each connected system with its objects in the flyout, and each object row carries its score; its page is a board of every object, one table per system, and a run can be chosen for any object. Below the seven sit small links to **Reports**, **Atlas** and, for administrators, **Administration**. Under 1100 px the rail becomes tabs along the bottom of the screen. The top bar shows where you are as a breadcrumb.
- **The canvas** shows the place you chose: the landscape, the Objects board, one object with its tabs, one rule, the Work ledger, a run or a comparison of two runs, a system, People, Reports or Administration.
- **The inspector** opens beside the canvas when you select a row: a record, a rule, a run, a job, a batch, a person, a system. It always has the same order: title, status line, facts, actions, history. `Esc` closes it and returns you to where you were.

Every main screen follows the same layout. A **banner** on top says the state in one plain sentence, such as the go-live verdict on the landscape or "25,712 records to fix" on Work, with a short line explaining it, up to three figures and one button for the next thing to do. Under it sits the screen's **chart**, horizontal bars that all start from zero so lengths compare directly, or the landscape's heat map of systems and objects. **Aletheia**, the grounded assistant, is docked on the right of the page; on a narrow screen it moves under the page. Every table shows the description beside each code: the rule name beside the rule id, the field text beside the SAP field, and so on. Meridian opens in a light theme; dark is a choice under Preferences.

The top bar holds the breadcrumb, the command line (⌘K or `/`), **Ask Meridian**, the inbox count, help (?) and your account menu with Preferences, Reports, the keyboard shortcuts and, for administrators, Administration. The run you are reading is shown by the status strip under the page.

The numbers on every screen

Every score in Meridian is drawn on one **ruler**: six segments whose widths are the dimension weights (completeness 25, accuracy 25, consistency 20, timeliness 10, uniqueness 10, validity 10), each filled to its own score, with a tick at the **measured** score. The number beside the ruler is the score **shown**.

When a critical rule is over tolerance the shown score is capped and reads with both numbers, for example "70.0 capped, 71.8 measured"; objects are ranked by the measured score.

The cap rule is fixed: one critical rule over tolerance caps a score at 85, two or more at 70. A critical rule counts as over tolerance when it fails at least the share of records or the number of records set under Administration, Settings, Cap rule and tolerance. The receipt under every object score states it in one line, for example: "Measured 71.8, the tick on the ruler. Capped at 70 by MM003 and MM035, two critical rules over tolerance."

Verdicts come from the score bands: **Go**, **At risk** and **No-go**. Value at risk comes from the cost model named in the footer, such as "cost model ZAR", and every amount is marked as an estimate: **est.**, or **default estimate** until the tenant saves its own cost model under Administration, Settings. Nothing is estimated by a language model: every figure is a rule result over stored data, and each one has an **i** beside it that shows the arithmetic.

Where each role starts

Everyone lands on the **Landscape**, shaped for their role by one **next-step button**:

- Programme lead: "Clear the 3 no-go objects in Wave 1: finance foundations".
- Data steward: "Accept N suggested values for" the top rule.
- Data owner: "Approve batch B-241".
- SAP Basis: "Extract SuccessFactors" (the system not yet extracted), with the Basis home above the matrix: connections, extractions, failed jobs, authorisation errors and schedules.
- Platform admin: "Open Administration".
- Process analyst: "Open PTP step 3".
- Auditor and viewer: the same landscape, read-only.

Under **Preferences**, **Start screen**, you can open Meridian as another role's start screen. Your access does not change.

Keyboard

- `⌘K` or `/` opens the command line. Type a material number, a rule id such as `MM142`, an object, a run, or a word such as approve or export.
- `g` then `h`, `o`, `r`, `f`, `w`, `u`, `m`, `i` or `p` goes to Home, Objects, Runs, Fix, Work, Rules, Master data, Migration or People.
- `?` shows every shortcut. `⌘Z` undoes the last change. `Esc` closes the top layer.

When something is wrong

- **A grey dash** means the number is unknown, never zero. Hover it for the reason.
- **"Not extracted yet. Extract"** on a system means nothing has been read from it; **"Extracted, not analysed. Run analysis"** means no scores exist yet.

- **The red line at the bottom of the screen** names a failed job, such as "Extraction of Employee central failed: OData refused credentials (HTTP 401)". Open it for the stages and a retry; SAP Basis owns connection errors.
- **"... needs the ... permission. Ask an administrator to grant it."** means your role does not hold that action. Your administrator grants it under Administration, Users.
- For anything else, open help (?), copy the diagnostics block under **Contact support** and send it to your administrator. It holds no record values and no secrets.

The Demo data bar

The demo build carries a bar at the top of every page: "Every figure on this page is sample data built into the interface, not read from your SAP systems." It cannot be dismissed, and a customer installation never shows it.

The strip at the foot of every page

The small strip with the Aletheia mark under every page is a licensed attribution, not a second product: the licence of the Aletheia mark requires it on every screen, so it cannot be hidden. Hover it for the same words in the interface's language.

The menu

The rail on the left is the menu. It has seven places: **Home**, **Objects**, **Runs**, **Fix**, **Rules**, **Master data** and **Migration**. Point at one, or move to it with the keyboard, and its menu opens beside the rail.

What is on it

- **Home** (`/`) is the landscape: every system and object with its score.
- **Objects** (`/objects`) lists the systems and their objects as a tree, with a link to connect a system. See the Objects chapter.
- **Runs** (`/runs`) is the timeline of every measurement.
- **Fix** holds the Action centre, Work, verified fixes from Aletheia, Cleansing, Requests and People. A count on Fix, and on Work in its menu, shows the records still to fix.
- **Rules** holds Rules, rule promotion from Aletheia, the rule dry run, anomaly rules, learned monitors, verified survivorship from Aletheia and the evidence ledger.
- **Master data** holds MDM health, the entry gate, reference data, record links, the data catalogue, Material 360, issues, data mining and process discovery.
- **Migration** holds Migration, S/4HANA readiness, go-live, the forecast and plan, value delivered, the score lens, cost and causes, and configuration. See the Migration chapter.
- **More**, below the rail, links to Reports, the Atlas and Administration.

Each menu has a **Pin the menu open** button and a **Close menu** button. On a narrow screen the rail becomes a bar at the bottom with the same seven places, and each opens a sheet.

What you can do, by role

Everyone sees the same seven places on the rail. Inside the menus, a few links show only for roles that can use them: rule promotion from Aletheia and the rule dry run show only for people who can manage rules. Under More, Administration shows only for people who can administer. The screens themselves still check the role. What each person can do on the screens depends on the role, as the chapter for each screen says.

The numbers and where they come from

The count on Fix adds up the records still to fix in the work lists, with each object at its latest run. Only runs whose work list has been built are counted. When no work list has been built, the count is not shown. It is never estimated and never replaced by a count of cleansing proposals.

When something is wrong

- **A menu does not open**: point at the place again, or focus it with the keyboard. Selecting the place itself opens its main page.
- **The menu closes when you move away**: pin it open with **Pin the menu open**.

Keyboard

- `g` then `h`, `o`, `r`, `f`, `u`, `m` or `i` goes to Home, Objects, Runs, Fix, Rules, Master data or Migration.
- `g` then `w` goes to Work, and `g` then `p` goes to People.
- `?` shows every shortcut.

Objects

Objects is the board of every object on every system. Open it from **Objects** on the rail, or press g then o. Its address is `/objects`. Use it to find the object you want, then open it for its score, its rules and its runs (see the An object and its tabs chapter).

What is on it

- **The sentence and the button**: one sentence says which object has the lowest measured score on the board, and the one button opens that object.
- **One table per system**, with the object that scores lowest first. Each row shows the **Object**, the **Score** out of 100, the **Trend** over the object's finished runs, the **Open findings** (records that fail at least one rule), **At risk** (the priced cost of those findings from the cost model, or a dash when nothing could be priced) and the **Run** the figures come from.
- **Held-down scores**: when failing critical rules cap an object, the row says so, for example "held down from 71.8". The score shown is the capped one; the measured score is the figure before the cap.
- **A system with no finished run** shows a plain sentence and a link to open the system, so you know to extract and analyse it first.

Which run you are reading

Each object is shown at its own newest finished run. To read every object at one run, pin that run: the board then says "Pinned to one run" and offers **Show each object at its newest run** to go back. Inside an object, the run picker sets the run, and the run stays in the address, so a link shows the same run.

Finding other things

- **Runs** on the rail lists every measurement on a timeline (see the Runs chapter).
- The **Objects** menu on the rail lists the systems and their objects as a tree, with a link to connect a system.
- The command line (Command K or slash) opens any object by typing its name.

What you can do, by role

Everyone who can read a system sees its objects. Acting on an object, such as fixing records, depends on the role, as the Work chapter says.

The numbers and where they come from

The score is the measured score out of 100, or the capped score when failing critical rules hold it down. Open findings come from the run shown. At risk is priced from the cost model, and reads as an estimate while the shipped costs are in use.

When something is wrong

- **An object is missing from its system's table**: no finished run scored it. The row says so, and the object page shows the earlier runs.
- **A system shows a sentence instead of a table**: it has no finished run. Open the system, extract it, then run an analysis.
- **A dash under At risk**: nothing could be priced for that object from the cost model.

Keyboard

- `g` then `o` opens Objects.
- `Esc` closes the topmost layer.

Landscape

The landscape is the start screen for every role. It answers three questions in order: how ready is the data across every system and object, what holds it back, and what is the one thing you should do next. It opens at `/`, from **Home** on the rail, or with `g` then `h`.

What is on it

- **The Demo data bar**, only in the demo build: "Every figure on this page is sample data built into the interface, not read from your SAP systems." A customer installation never shows it.
- **The top sentence**: one plain line above everything that says what state the data is in and what to do next, for example "This page shows every connected system at its newest run: 2 systems and 15 objects. 11 cannot reach 85 yet. Start with the fix queue, or open any object to see its run." Before any run it says so and tells you to extract a system and run an analysis.
- **The next-step button**, straight under the top sentence: one per role (see below), the one strong button of the page, with the wave and its due date beside it for the programme lead.
- **One strip per system**, at that system's own newest run (there is no run picker on Home): the system name and whether it is active (an inactive system that still holds finished runs stays listed and is labelled **Inactive**), its newest run as "{System} run N" with the date on hover, the mean measured score of its objects out of 100 with its verdict, a small trend of its lowest-scoring object across runs, its open findings and their priced cost from the cost model. Every connected system has a strip, so the strips add up to the count in the top sentence: a system never extracted or never scored says so and links to Systems; a strip whose runs cannot be read says so with a link.
- **Three columns under the strips**: **Worst movers**, the objects that fell furthest since their previous run (select one to open it); **Fix queue**, the top of the work ledger, each row with its object and record key so rows of the same rule are told apart; and **Found by Aletheia**, measured facts nobody has confirmed yet. Confirm one before you act on it.
- **The title and scope**: "Landscape" and the systems in scope, such as "All systems".
- **The go-live banner**, at the top: one headline that gives the verdict of the next go-live. The small line above it names the next migration wave ("Next go-live: Wave 1: finance foundations"), or reads "All objects in scope" when no wave is planned. The headline is the worst verdict among the objects in that wave: **Go**, **At risk** or **No-go**. The line under it says why, for example "3 objects block go-live: Material master, Customer master, Vendor master", "Nothing blocks go-live, but some objects are below target." or "Every object in scope meets its target." Three figures follow: **Days to go-live** with the wave's date ("no date set" without one), **Cost of waiting** per month from the cost model ("no cost model" without one) and **Objects ready**, such as 4/15, with the critical issues still open. The banner has one button, chosen by the state: **Fix 3 blockers** opens Work on the no-go objects, **Clear 12 critical issues** opens Work on the critical rows, and **Sign off the wave** opens the waves when nothing is left.
- **The sentence**: one line that says what holds the landscape back, for example "11 of 15 objects are held under 85 by 30 critical rules over tolerance; fixing MM003 and MM035 would release 1 of them."

- **The provenance line**: records checked, rules measured and the run, for example "400,662 records checked, 794 rules measured, run 16 Oct 2026, 05:17".
- **The landscape map**: a heat map with one row per system (its name, type and client) and one tile per object, the lowest score first. The colour, the mark and the word on each tile all give the verdict, so it reads without colour too. Each tile shows the object, its score and either the records to fix, the critical rules failing ("2 critical rules failing") or "clean". The legend at the top counts the tiles of each verdict. Select a tile to open the object.
- **Three cards under the map**: **Readiness over time**, the average object score over the last runs ("Two runs are needed to draw a trend." until there are two); **Top blockers**, the five critical rules failing the most records, each with its message beside the rule id and the SAP field, and a link to Work filtered to that rule ("No critical rule is failing." when none is); and **Cleaned and watched**, the records fixed this week with a small weekly chart of records fixed against records that failed again, and **Open monitoring**.
- **Aletheia**, docked on the right: "Grounded answers from this run. Every number is checked against a fact, with the rules cited." Ask a question in plain words; see the Ask Meridian chapter. On a narrow screen the panel moves under the page.
- **The coverage matrix**, under the panel: objects as rows, systems as columns. Each cell is a small ruler with the shown score, the word **capped** and the measured score beside it where a cap applies ("70.0 capped, 71.8 measured"), the verdict (Go, At risk, No-go), records checked and rules measured. A cell flagged **low coverage** has fewer than 70% of its rules producing a result. A system never extracted shows "Not extracted yet. Extract"; one extracted but not scored shows "Extracted, not analysed. Run analysis".
- **The order**: objects are ranked by the measured score, so capped objects still sort by how good they are.
- **Coverage, All objects and Waves**: the view switch. **Waves** groups the objects under each migration wave with its due date, its go, at-risk and no-go counts and the gate state (see the Waves and gates chapter).
- **Work**: the fix queue in the strips above shows the top rows; open **Work** for the full ledger. SAP Basis sees system tasks in the Basis home instead.
- **The Basis home** (SAP Basis only), above the matrix: **Connections**, **Last extraction**, **Failed jobs to triage**, **Authorisation errors** and **Schedules**. See the SAP Basis walkthrough.
- **Objects from uploaded files**: when no system is connected but files have been uploaded, the landscape lists those objects with their records to fix, says they came from uploaded files, and offers Work, the runs and connecting a system.
- **Ask Meridian**, in the top bar on every page, opens the same grounded answers in a panel of its own.

The landscape atlas

Open the atlas on the landscape opens a mosaic of every object of every system. Each tile is one object: its size follows the records it checks, its shade follows its score band (95 and up, 85 to 95, 70 to 85, below 70) and it prints the score and the change the shown run made. A tile with a heavy border moved in that run.

The **Runs** strip under the banner has one bar per stored run, oldest first, counting the objects at 85 and up after that run. Select a bar to see the landscape as it stood after that run, or press **Replay the runs** to step through them. Objects the run did not score keep their last stored score; nothing is interpolated between runs. Select a tile to open that object at that run.

The numbers and where they come from

- **The banner's verdict** is the worst verdict among the objects of the next wave: one no-go object makes it No-go. Without a planned wave it covers every object on the landscape.
- **Days to go-live** counts the days to the wave's target date. **Cost of waiting** is the monthly opportunity cost from the cost model, an estimate. **Objects ready** counts the objects of the wave whose verdict is go; the critical issues are the open critical rows in Work.
- **The tiles** are ordered by score, lowest first; a tile's records to fix are the open rows of that object.
- **Readiness over time** plots the average object score of each run. **Top blockers** counts the failing records of each critical rule.
- **The score in each cell** is the object's shown score in the run in the top bar: the weighted dimension scores, capped at 85 by one critical rule over tolerance and at 70 by two or more. Open the **i** beside the object name for the receipt.
- **The verdict** comes from the score bands under Administration, Settings: at or above target is go, at or above warning is at risk, below that is no-go. A critical rule over tolerance keeps an object out of go.
- **The sentence's counts**: "held under 85" counts the objects whose shown score is capped; "critical rules over tolerance" counts the rules behind those caps; "would release" names the rules whose fix lifts the most objects off their cap.
- **Value at risk** in the fix queue is the cost model named under the table: "Rows ranked by value at risk, cost model ZAR." Every amount is an estimate and says so: **est.**, or **default estimate** while the tenant has not saved its own cost model, with a link to Administration, Settings.
- **The footer** repeats the run, the records checked, the rules measured, the scoring time, the cost model and the licence validity, and shows a failed job in red when one needs attention.

What you can do, by role

- **Programme lead**: follow the next step ("Clear the 3 no-go objects in Wave 1: finance foundations"), switch to **Waves**, open any cell, share the page from **Share and export**, export the grid to Excel.
- **Data steward**: "Accept N suggested values for" the top rule opens Work on that rule's rows; the fix queue lists the highest-value records.
- **Data owner**: the line under the sentence says how many batches wait for you and your SLA, for example "2 batches wait for your approval, SLA met on 92% of your open work"; the button is "Approve batch B-241".
- **SAP Basis**: the Basis home above the matrix; the button extracts or analyses the system that needs it first; the task list shows every system's state.
- **Process analyst**: "Open PTP step 3" opens the first blocked process step.

- **Auditor**: the next-step button is "Open the rules library".
- **Viewer**: sees the programme lead's start screen. Buttons that need a right the role does not have, such as connecting a system, uploading or starting an analysis, are not shown.

Selecting a tile in the heat map or a cell in the matrix opens the object page. Selecting a system heading opens the system in the inspector.

Keyboard

- `g` then `h` returns here from anywhere.
- ⌘K opens the command line ("Search, ask, or type a command"); type an object name and system, such as "material master ECC".
- In the matrix and in the fix queue, `j` and `k` move, `Enter` opens.

When something is wrong

- **The banner reads "no date set"**: the next wave has no target date. Set it on the wave (see the Waves and gates chapter).
- **Cost of waiting shows a dash**: there is no cost model. Set one under Administration, Settings.
- **Many objects read 70.0 capped**: they have two or more critical rules over tolerance, so the cap holds them at 70. The measured score beside it differs, and the objects are ranked by it.
- **A system row says "Not extracted yet"**: ask SAP Basis to extract it, or extract it yourself if you hold trigger sync.
- **A cell shows a dash**: no run has scored that object on that system. "Not analysed" and "Not extracted" say which step is missing.
- **Low coverage**: some rules could not run because a table, field or value list was missing. Treat the score as indicative and ask SAP Basis to check the extraction scope.
- **The footer shows a failed job**: open it from the footer link; the job inspector says why it stopped and offers a retry.

Ask Meridian

Ask Meridian answers questions about the current run in plain words. It is grounded: the answer may use only the facts Meridian measured, and every number in it must match one of those facts. The same answers are in two places. **Aletheia** is docked on the right of every main screen: the landscape, an object, a rule, Work, Runs, a run and a comparison, People, Requests, the rules library, Reports, a system and Administration. Its header reads "Grounded answers from this run. Every number is checked against a fact, with the rules cited." On a narrow screen the panel moves under the page. **Ask Meridian** in the top bar opens the same answers in a panel of its own, and you can also type a question in the command line and choose Ask Meridian.

What is on it

- **Your question**, typed in the box ("Ask about this run") and sent with **Ask**, and four suggested ones: "What is holding the score down?", "Which objects are under 85?", "What is the money at risk?" and "What did Aletheia find?".
- **The answer**, marked either "Grounded, every number matches a fact" or "Not answered" when the draft quoted a figure that is not a fact.
- **Facts used**, with "Coverage" saying how much of the question the facts cover, and **Rules cited**, so you can open the rule or object behind each sentence. The intent Meridian read from your question is shown, and "Reworded by model" when a language model phrased the sentence.
- **Was this answer helpful?** Under each answer, press **Helpful** or **Not helpful**. One vote per answer is recorded with the run it came from, so the team can see which answers worked; after the vote the buttons give way to a thank you.
- The Aletheia mark, which identifies the grounded intelligence that checks the answer.

Aletheia found

At the top of the landscape and of every object, the **Aletheia found** band lists what Aletheia found without being asked. It ranks the failing rules of the latest run by their impact on the score. Each finding shows the rule number with its text, the object and the SAP field it fails on, and why it matters: how many failing records it gained or lost since the previous run, and, where Meridian traced where the bad values come from, what share of them come from one origin. Beside it are the failing records with a sparkline over the last runs, the money at risk when a cost model covers the rule, and a confidence. The confidence is the share of reviewed issues stewards confirmed as real, or the share of failing values from the dominant origin; when neither has been measured it shows a dash.

Each finding carries one action. **Draft cleansing batch** drafts a cleansing batch over the rule's open issues; it still needs approval before anything is written back. **Draft tuned version** drafts a tuned version of a rule that produces too many false alarms, for review in the rules library. **Open rule** opens the rule when there is nothing to act on directly.

Under the findings, **What changed since the last run** shows each object's score change against its previous run and the rules that moved it. **Proof of value** follows every finding someone acted on: it waits for the next run, then shows the score gain, how many fewer records fail, and the proven cost before and after for the cost metrics that rule feeds.

Ask "What did Aletheia find?" in the dock to hear the same findings as grounded sentences. The answer uses only rule numbers, field names and counts, never record values.

The numbers and where they come from

Meridian collects the run's facts first: scores, failing counts, money at risk and the rules behind them. A language model only rewords those facts. Before the answer is shown, every figure in it is compared with the facts; if one does not match, the answer is withheld rather than shown with a wrong number.

What you can do, by role

Every role can ask. The answer only covers objects and systems your role can already see.

Keyboard

- `⌘K` opens the command bar; type a question there or use Ask Meridian in the top bar.
- `Enter` sends the question; `Escape` closes the panel.

When something is wrong

- **"Not answered"**: the question needs facts the run does not hold. Ask about scores, rules, objects or money at risk.
- **"The numbers in the draft answer did not match this run's facts, so it was withheld."** The model's draft did not match the facts. Rephrase, or open the screen the facts point to.
- **The answer reads as from an older run**: answers use the latest finished run. Check Runs for one still in progress.

An object and its tabs

An object is one area of SAP data scored as a whole on one system, such as Material master on ECC production. The object canvas is where you read its score with the receipt that proves it, and where every view of the object lives as a tab: Overview, Rules, Records, Fixes, Trend, Readiness, Processes, Lineage, Profile, Glossary and Golden records. Open it from a cell on the landscape, from **Objects** on the rail (the board lists every object of every system with its score, or choose it from the Objects menu), or by typing its name in the command line. Its address is `/o/{system}/{object}`, for example `/o/ecc-production/material-master`.

The header

On Overview and Rules, under the tabs, the **S/4 readiness** banner says where the object stands; the other tabs put their own content straight under the tabs. The small line names it ("Material master on ECC production: S/4 readiness"), the headline is its verdict (**Go**, **At risk** or **No-go**) and the line under it is the receipt sentence, such as "Capped at 70 by MM003 and MM035, two critical rules over tolerance." Two figures follow: **Records to fix**, "of 61,816 checked", and **Money at risk**, per month from the cost model ("estimate, default costs" while the shipped costs are in use, "no cost model" without one). The one button is **Fix 2 blockers** while critical rules are over tolerance (it opens Work on this object's critical rows), otherwise **Fix 25,712 records**, and **Review the runs** when nothing is left to fix.

Aletheia is docked on the right of the page, as on every screen: ask about this object in plain words and the answer is checked against the run's facts (see the Ask Meridian chapter).

The tabs sit under the title: **Overview**, **Rules** (with its count), **Records** (with the records to fix), **Fixes**, **Trend**, **Readiness**, **Processes**, **Lineage**, **Profile** (the fields of each table, with type, how many are blank and the range or common values), **Glossary** and **Golden records**. The object opens on Overview: one sentence on what you are reading, the one next step with its reason, the dimension bars, the top failing rules and the score at each finished run. Rules carries the ruler and the receipt.

- **Watch**: "You hear about runs, score drops, new critical breaches and comments on Material master." Watching is per person; what reaches you is set under Preferences, What you hear about.
- **Data contracts**: the contracts that name this object.
- **Related objects**: each related object with its relation, **upstream** or **downstream**. Fix the upstream object first: records there fail before this object can pass.
- **The facts line**: the run, records checked, rules measured, records to fix, value at risk, S/4 features blocked and the wave verdict, for example "Run 16 Oct 2026, 05:17, 61,816 records checked, 80 rules measured, 25,712 to fix, ZAR 4.2 m est. at risk, estimated from your cost model, 3 features blocked, No-go for wave 1". While the tenant has not saved its own cost model the amount reads **default estimate** and links to the cost model in Administration.

The ruler and the receipt

The large ruler has the 0 to 100 axis with the 60 critical, 75 at risk and 90 go marks, and two cap lines, **cap 1 critical** (85) and **cap 2 or more** (70). Each of the six segments is filled to its dimension score, which is printed under it. The tick is the measured score; the number beside it is the score shown, with

capped and the measured score when a cap applies: "70.0 capped, 71.8 measured". Each segment is named by the formula's letter (C, A, Cs, T, U, V), with the full name on hover.

Under it the receipt says it in words, exactly as the product states it: "Measured 71.8, the tick on the ruler. Capped at 70 by MM003 and MM035, two critical rules over tolerance."

The receipt table then adds the score up: one row per dimension with its weight, score, points, rules measured, records checked and coverage; then **Weighted**, **Measured**, **Cap, 2 critical rules over tolerance** and **Score shown**. The formula line under it is the full rule: the six weights, "severity weights critical 4, high 3, medium 2, low 1; tolerance 10% of rows or 25 rows, as set on each rule; cap rule: 85 at one critical rule over tolerance, 70 at two or more." Every row has an **i** that explains the figure from the run itself.

Policy classes breached follows when an applied policy pack's class has a failing critical rule of this object: the class, its authority and the rules. A filled shield beside the score says the same on the ruler and the landscape cell; the diamond beside it is a data contract.

The **run picker** beside the object's name chooses which of this object's runs you read. Each run is named by its number, such as "SAP ECC run 12", with its date beside it. An earlier run shows a band that says so, the ruler and the rules are that run's, and actions are off with the reason "Open the current run to act"; the band's button returns to the current run. A one-run object shows a note instead of a chart; the trend starts at the second run. The choice is written in the address as `?run=`; an old `?as_of=` address is redirected to it.

Score by dimension, on the Rules tab, is a bar chart of the six dimensions on one scale from 0 to 100: each bar is the dimension's score, with its share of the score beside the name ("25% of the score") and the value at the end of the row, or "not measured" when no rule of that dimension produced a result. The shortest bar is the dimension that costs the most.

Top failing rules, beside it, lists the five failing rules with the highest value at risk: each with its severity and name, the rule id, the SAP field with its description beside it (`MARA-MEINS` Base unit of measure), a small trend of its failing rows over its last runs ("Failing rows over the last 8 runs") and "2,904 failing", with a bar sized to the worst of the five. Select a rule to open it.

Critical rules over tolerance lists the rules behind the cap with their failing records and share, for example "MM003 Material type is mandatory - determines procurement and storage behaviour, 2,904 failing, 4.7%" and "MM035 Valuation class is mandatory for FI account determination, 1,420 failing, 2.3%".

The tabs

- **Rules**: every rule of the object, ranked by value at risk, with field, dimension, severity, tolerance, over tolerance, failing rows, failing %, coverage and value at risk. Each rule says what goes wrong in SAP while it fails, for example "Goods movements and MRP cannot post for a material whose base unit is not an ISO unit." Where suggested values exist, **Accept all** accepts them for that rule. Pivot the list **By rule**, **By field**, **By dimension**, **By severity** or **By S/4 feature**; **Learned** and **Drafts** list rules waiting for review. The line under it reads "Ranked by value at risk; total ZAR 4.2 m across 80 rules measured. Failing % is of 61,816 records checked." Where the tab shows only part of the object's rules, the last cell ("and 905 more rules") opens the rules library filtered to the object.
- **Records**: the Work ledger scoped to this object. See the Work ledger chapter.

- **Fixes**: the fix batches of this object. "No fix batches for Material master yet. A batch starts when a steward accepts suggested values or imports a filled fix sheet."
- **Trend** (the old Runs tab; `/runs` still opens it): the measured score over the last eight runs with the cap line and the baseline, a table of measured and shown scores per run ("Scores are as stored with each run; the measured score is before the cap, the shown score after it."), **Compare with** to open two runs side by side, and **Why the score moved** against the baseline: rules newly failing, fixed, worse and better.
- **Readiness**: the migration waves this object belongs to with the verdict, blocking gaps, score, readiness and records blocked, and "What the failing rules do in S/4HANA": the features blocked and degraded, each with the rules behind it. **Blockers** turns that round: one row per rule of this object behind a blocked or degraded feature, with its severity, the features it holds back and its failing rows, most failing first. **Readiness by object** lists every object in the same waves with its ruler, verdict and blocking gaps, the most blocked first, with this object marked; select one to open it.
- **Lineage**: the drawing from source table to S/4 feature. Select a node to pin its facts, with its **Upstream** and **Downstream** nodes listed under them; select one of those to move to it. See the Lineage chapter.
- **Profile**: what the run found in each field the object reads, table by table: the field with its dictionary text, its type, the blank share, distinct values, length, and the commonest values or the range of a number or date. Values stay withheld for personal data and free text. **Mined dependencies** lists where one field decides another in almost every record, for example "MARA-MTART Material type decides MARA-MBRSH Industry sector, holds in 99.3%", with the rows breaking it. A person who manages rules picks a severity and selects **Make it a rule**; the rule runs from the next analysis.
- **Processes**, **Glossary**, **Golden records**: each has its own chapter.

The numbers and where they come from

- **Score by dimension** reads the dimension scores of the run in the top bar; the shares are the dimension weights under Administration, Settings.
- **Records to fix** in the banner is the open rows of this object in Work; **Money at risk** is the cost-model estimate for its failing rules.
- **Top failing rules** are ranked by value at risk among the rules with failing rows in the run; the trend is read from each rule's own history.
- **Blockers** comes from the feature impact of the run: the failing rows of each rule that blocks or degrades a feature.
- **Readiness by object** takes the verdict and blocking gaps from the waves' readiness and the ruler from the run in the top bar.

What you can do, by role

- **Everyone**: watch the object, read every tab, export any grid with **Share and export**, comment on a rule in its inspector.
- **Data steward**: **Accept all** suggested values on a rule (needs apply), open the records to fix.

- **Data owner**: apply the steward template to objects you own, from the object's steward template action.
- **Programme lead**: share the page as a link, compare runs.
- **Process analyst**: the Processes tab and its what-if simulation.
- **Rules manager**: turn a mined dependency on the Profile tab into a rule.

Keyboard

- In every tab table, `j` and `k` move, `Enter` opens the row in the inspector, `x` selects.
- `[` and `]` step to the previous or next row while the inspector is open.
- `Esc` closes the inspector.

When something is wrong

- **"... has no object called ..."**: the address names an object Meridian does not know; the page says so instead of showing a real-looking page. Open **Objects** on the rail and choose an object. A record key that is not in the latest run says so the same way.
- **"No run yet"** and an empty ruler: this object has never been analysed on this system. Ask SAP Basis to extract and analyse it.
- **The measured and shown scores are the same, but the receipt names a critical rule**: "A critical rule is over tolerance, but the score is already below its cap."
- **The Records tab says the ledger could not be loaded**: open Work from the link under the message; the ledger reads the latest run of each object.
- **Low confidence**: fewer than 70% of the rules attempted produced a result. The receipt row **Coverage** shows which dimension is short. SAP Basis can widen the extraction.
- **A figure looks wrong**: open its **i**. If the arithmetic there does not match, copy the diagnostics from help and send them to your administrator.

The rule canvas and the rules library

Every number in Meridian comes from a rule: a fixed, deterministic check with an id such as `MM142`, a field, a dimension and a severity. The **rule canvas** shows one rule across every object and system it runs on. The **rules library** lists every rule of the tenant, with drafts and suppressions. Open a rule by typing its id in the command line, from any rule row, or at `/r/{id}`. Open the library from the command line ("Rules") or at `/rules`.

The rule canvas

The title is the rule name, for example "Base unit of measure MEINS must reference a valid entry in T006", with the rule id (`MM142`) under it and **All rules** to go back to the library. The banner says "This rule" and counts the failures in its headline ("4,212 records fail this rule", or "No record fails this rule"); the line under it says what the rule checks, in words and with the SAP field ("It checks Base unit of measure `MARA-MEINS`"). Two figures: **Failing records**, with the share of records checked under it ("4.99% of records checked"), and **Objects affected**. The one button is **Show the failing records**. Aletheia is docked on the right.

The facts line under it reads like "High, Consistency, 4,212 failing across 2 objects, 4.99% of records checked, ZAR 1.18 m at risk, estimated, within tolerance, blocks 3 S/4 features".

Seven tabs follow. The rule opens on **Overview**:

- **Overview**: why the rule exists (its authority and what goes wrong in SAP), a chart of its failing records over the last runs, and **Failing records**. Where the failing records key on more than one field, such as material and plant, the newest run's failing records are counted by each part of the key, for example by **Plant**: one chip per value with its count, and **All**. Select a chip to narrow the ledger under it to those records; select it again, or **All**, to show every failing record.
- **Blast radius**: where the rule's bad data goes. Four columns: the fields the rule reads, the process steps those fields feed (with the transaction), the S/4HANA features the steps enable, and the KPIs they drive. Each step shows the open documents the failing records touch, joined in SAP on their keys; each blocked or degraded feature shows its cost a month. Select any field, step or feature to light its path end to end and dim the rest; select it again to clear. The figures above the map total failing records, documents in flight, steps touched, features blocked and the cost a month. A dash means Meridian has not measured that value, for example before the rule has run.
- **Definition**: the check class (such as "Exists in check table"), the field (`MARA-MEINS`), the expression (`MARA-MEINS in T006-MSEHI`), its parameters such as the check table `T006` and check field `MSEHI`, the source ("Shipped with Meridian" or a custom rule), the authority ("Best practice", or "SAP hard constraint") and what goes wrong in SAP. **Versions** lists every version with who changed it and why; the shipped definition names its file and content hash.
- **Failing records**: the scope, for example "Scope: 2 objects, 1 system", with each object, its run and its failing count. **Open failing records** opens the Work ledger filtered to this rule.
- **Prove in SAP**: how to reproduce the rule's count by hand in SE16 or SE16N (see below).
- **Impact**: the S/4 features this rule blocks or degrades and the lineage from field to rule to feature.

- **History**: each run's failing count next to the version changes, so a jump in failures can be read against a change to the rule.

Contract shows the data contracts that name this rule.

Prove in SAP

Your SAP team can check any number Meridian shows without trusting Meridian. **Prove in SAP** gives, for the rule and the run of one object (choose the object when the rule runs on several):

- **Record grain**: the table the rule judges and its key, for example `MARA`, one record per `MATNR`.
- **SE16 selection**: one row per selection field, marked **Population** (the records the rule judges) or **Failing** (added to find the failing ones), with the SE16 option and values. **Initial value** means leave the value empty with the option "=", which selects blank, 0 and 00000000. The population always leaves out what SAP itself marks as dead, such as records flagged for deletion (`MARA-LVORM` is initial), and the facts say how many records that left out.
- **What SE16 cannot express**: a pattern, a lookup in a check table (for example "MARA.MEINS has no entry in check table T006: list T006 in SE16 and compare"), a duplicate, a join to another table, a check digit. The selection above is then the closest one, and this list says the rest in plain words. When one selection is enough, the section says "One SE16N selection on MARA gives the failing count".
- **Population** and **Failing**: the counts Meridian measured in that run, to compare with the number of entries SE16 shows.
- **Five failing records**: their keys as SE16 shows them (`MATNR 232`, without the leading zeros of the conversion exit), with the key as stored (`MATNR=0000000000000232`) beside it.

To prove a run in one go, **Download the reconciliation workbook** (needs the export permission): a summary sheet with one row per rule and empty **SE16 count** and **Match** columns for your team, and one sheet per rule with its selection, steps, counts and sample keys, for the ten rules with the most failing records in that run. The same pack is `GET /api/v1/rules/{id}/reconcile?run={run}` and the workbook `GET /api/v1/runs/{run}/reconciliation.xlsx?top=10`.

Cloud sources (SuccessFactors, Concur, Ariba) have no SE16: the section says so and the same selection applies in the source's own reporting.

The rule inspector

Selecting a rule anywhere opens the inspector with the same facts: severity, dimension, field, **Tolerance** (for example "25 rows" on `MM003`; only critical rules carry a tolerance, so a high rule such as `MM142` reads within tolerance), **Over tolerance**, failing records by object, **Failing %** ("4.99% of 84,456 records checked") and **Value at risk**. Below come the definition, scope, the features it blocks, lineage, versions, contract, **Prove in SAP**, history, **Watch** and **Comments**.

The rules library

The banner on top says "Rules library". Its headline counts the enabled rules that fail ("312 rules fail on live data", or "No enabled rule fails") and the line under it names the worst, the rule with most failing records, with its id beside its message ("The worst is MM003 Material type is mandatory, 2,904 failing

records"). Three figures: **Rules** ("across 15 objects"), **Enabled** ("run on every analysis") and **Suppressed** ("kept out of the score"). The one button is **Open the worst failing rule**; with nothing failing a rule manager sees **New rule draft** instead. **Failing rules by object** under it is a horizontal bar per object, the eight with most failing rules, with "41 failing rules" at the end of the row. Aletheia is docked on the right.

The library title says what is there, for example "794 rules across 15 objects, 794 enabled". It reads every page of the rule list, opens filtered to an object from that object's Rules tab, and its **Share and export** offers Excel and a CSV of the columns and rows on screen. Three views: **All rules**, **Drafts** and **Suppressed**. Filter by object, dimension, severity, check class and state; the table shows rule, object, severity, dimension, class, failing records, pass rate, state and source. **New rule draft** starts a rule in authoring; a draft runs nowhere until it is reviewed and enabled.

Migration scenarios sit above the table: a card titled "Which rules matter for your S/4HANA path" with one row for each way to reach S/4HANA, each with a plain description: **New implementation** (a fresh system loaded through the migration cockpit), **System conversion** (the existing ECC system converted in place), **Selective data transition** (chosen objects and organisational units copied across) and **Carve-out** (one company code or business unit extracted). Each row counts the enabled shipped rules that are **Required**, **Recommended** or **Not applicable** for that path, and how many are migration blockers. Disabled rules are not counted. The **Migration scenario** filter, and the **Show rules** button on a row, limit the table to one path and add three columns: **Migration relevance**, **S/4 area** and **Why**, the reason in words. A second filter, **Relevance**, narrows the list to required, recommended or not applicable rules. The relevance comes from each rule's S/4HANA authority, its impact (blocking or warning) and its S/4HANA area, through one fixed mapping that is the same for every tenant. A rule that records no S/4HANA area is recommended for every path.

Rules can also be imported and exported in bulk under Administration, Rules import and export: imported rules are validated against the data dictionary and land as drafts for review, never live.

Aletheia rule promotion

Meridian ships about 2,700 rules that were generated from the SAP data dictionary. They are all switched off, because a rule that is sound in general is not yet proved on your data. The Aletheia rule promotion page, under Rules, Aletheia rule promotion, or at `/rule-promotion`, switches them on for your tenant, and only on evidence. No AI is used and no record leaves Meridian.

The page lists each object with its generated rules, how many are promoted and how many fall in each bucket. Open an object and press **Preview**. Meridian runs each switched-off rule on the latest extraction, changes nothing, and sorts the rules into four buckets:

- **Ready to promote.** The field is filled in at least 20% of records, the rule fails at most 2% of them, and at least 100 records are checked.
- **Needs review.** The field is filled, but the rule fails more records than the threshold. The row shows the failing count, the rate and up to five sample record keys, never a field value.
- **Needs a value set.** The rule checks a value against a list, and your system's live configuration has no such list. It is never switched on automatically, because it would flag every value you added yourself.
- **Skipped, unused.** The field is empty, filled in too few records, or was not extracted.

Each row shows the rule with its description, the field with its description, the check, the failure rate, how much of the field is filled and how many records were checked. The thresholds can be set per tenant in the `rule_promotion.thresholds` setting.

Users who may manage rules also see three buttons. **Apply** promotes every rule in the Ready bucket. **Promote anyway** on a Needs review row asks for a reason of at least five characters and promotes that one rule. **Revert** puts a promoted rule, or all of them, back to the shipped state. Every promotion and every revert is written to the audit log with its evidence, and a promoted rule runs from the next extraction.

Promoting is not the end. A promoted rule is checked again each time you press Preview. When its failure rate has risen five points or more over the rate it was promoted at, the page shows it in a warning above the table, with the old rate, the new rate and the sample records.

Rule dry run

Before you switch a rule on, edit it or add one, the Rule dry run page shows what it would do. Open it from Rules, Rule dry run, from the Dry run button on a rule, or at `/rule-dry-run`. It needs the manage rules permission.

Name the object and the rule, switched on or off. To test an edit, give the fields to change as JSON, for example a new pattern. To test a new rule, leave the rule id empty and give the whole draft. Meridian runs the rule on the latest stored extraction, on at most the number of rows per table that you choose. It never calls SAP and writes nothing to the findings.

The page shows how many records were checked and how many fail, the failure rate, the projected object score before and after with the same scoring a run uses, and the failing values grouped by shape, never as values. For an edit it also shows how many records start to fail and how many start to pass against the current rule. When the row cap cut a table, the page says so and the counts are scaled to the full table. A strict rule can leave the score unchanged when a critical breach already caps it.

Aletheia verified fixes

Many rules can tell you a value is wrong but cannot say how to correct it. A rule has an automatic fix only when it carries `auto_fix` steps. The Aletheia verified fixes page, under Fix, Aletheia verified fixes, or at `/fix-drafts`, closes that gap for one rule at a time and never applies anything by itself.

The top of the page charts, for each object, how many rules have a fix and how many do not. Below it, each drafted fix is a card with the rule and its description, the field and its description, and the aggregated shape of the failing values: the most common masks (letters as A or a, digits as 9), lengths, case, and how many values carry leading, trailing or inner spaces. The drafter sees this shape and nothing else. No field value ever reaches the model.

The drafter picks steps only from the fixed set of fix operations, such as trim, upper case, pad on the left or map values. Each operation is shown with its description. Meridian then proves the steps on your latest extraction and shows the result as a bar: how many failing values are fixed, how many the rule still rejects after the fix, and how many are still failing. It also counts passing values the steps would change, which should be none. When the model is switched off or fails, Meridian drafts from a short list of safe candidates instead and marks the draft as drafted by rule.

Every draft waits as pending. Users who may manage rules can press **Approve**, which makes the steps the rule's fix with low confidence, or **Reject**, which keeps the draft in the audit log and changes nothing. A rule that already has a fix, or has no extraction to prove against, cannot be drafted.

The numbers and where they come from

- **Failing** is distinct records failing the rule in the latest run of each object in scope.
- **The Overview trend** plots the failing records of each run that measured the rule, up to the last 30.
- **The chip counts** are distinct failing records in the rule's newest run with that key value.
- **Failing %** divides by the records checked in those runs.
- **Over tolerance** applies to critical rules only: a critical rule is over tolerance when it fails at least the share or the count of records set as its tolerance. "Only critical rules count: when one is over tolerance it caps the score (85 for one rule, 70 for two or more); a critical rule within tolerance still lowers its dimension but never caps."
- **Value at risk** is from the cost model of the run, in the tenant currency.
- **Pass rate** is the share of checked records that pass.

What you can do, by role

- **Everyone**: read, watch, comment, narrow the Overview ledger by a chip, export the library, read the SE16 selection of a rule.
- **Export permission**: download the reconciliation workbook of a run.
- **Data steward and rule maintainers** (manage rules): **New rule draft**, **Edit in authoring**, **Suppress** with a reason, **Hand to a steward**.
- **Process analyst** (handoff rule): hand a blocking rule to a steward from the step inspector.
- **Platform admin**: enable, disable and change severity; bulk import and export.
- **Auditor**: read the versions and the history; the rule's change trail across versions and runs is read side by side, there is no single joined trail yet.

Keyboard

- ⌘K, then a rule id, opens the rule.
- In the library table, `j` and `k` move, `Enter` opens the rule in the inspector.
- `Esc` closes the inspector.

When something is wrong

- **A rule fails far more records than last run**: compare **History** with **Versions**. A version change on the same day usually explains it; otherwise the data changed.
- **Coverage below 100%**: the rule could not read some records because a table, field or value list was missing from the extraction.

- **A rule is wrong for your business**: suppress it with a reason, or mark the affected records not an issue in Work. Both are audited. Ask the rule's maintainer or your administrator to change its severity or tolerance.
- **Your edit buttons are missing**: authoring needs the manage rules permission.

The Work ledger

Work is the one list of everything that needs a person: every failing record and field, with its current value, the values the rule accepts, a suggested value, a New value cell and the decision, on one row. The row is the fix. Open it from **Work** at the top of the map, with `g` then `w`, or at `/work`. On an object, the **Records** tab is the same ledger scoped to that object.

What is on it

- **Ledger and Monitor**: the two views of Work, as tabs at the top. **Ledger** is the list described here; **Monitor** (`/work/monitor`) shows whether records a batch fixed stay fixed: see the Post-clean monitoring chapter.
- **Exceptions** (`/work/exceptions`): what went wrong in a business process or an SAP transaction, raised by an exception rule or by hand and worked under an SLA set by severity (critical 8 hours, high 24, medium 72, low 168). The banner counts the open exceptions, how many are past their SLA, the share resolved within SLA, the average time to resolve and the SAP transactions raised in the last 24 hours. Under it, open exceptions by severity, the status tabs (Open, Investigating, Waiting for approval, Resolved and so on) and a filter by severity and source. Each row shows the title and description, what raised it (for example `MIGO 5000412233`), severity, status and tier, owner, SLA date and value at risk. Open a row to read the linked finding and the comments, assign it (moves it to Investigating), escalate it one tier with a reason (needs approve; the SLA halves at each tier, up to tier 4) or resolve it with what was done and the root cause (needs approve). **Rules that raise exceptions** lists each rule with its object and condition; a role that can manage rules switches a rule on or off. A new rule starts off.
- **Triage** (`/work/triage`): the ranked queue of issues and stewardship items for one person, a team, the unassigned or everyone (pick it in **Show the work of**). The banner counts the items waiting and those past their SLA, with the share resolved within SLA over the last weeks, the average time to acknowledge and to resolve, and the breaches. Under it, the weekly backlog burn (resolved minus opened) and the open items by owner and by team. The queue splits into **Overdue**, **Due today** and **Due later**; each row shows the finding, severity, owner and team, priority, SLA state, due date and rank score (severity times affected records times the features the rule blocks, raised as the SLA runs down; hover it for the parts). A role that can assign ticks rows and changes them together: assign or reassign to a person or a team, release, set a priority from P1 to P5, acknowledge, snooze until a date with a reason, wake up, pause the SLA while waiting for SAP or the requester, or resume it. **Teams** share work by taking turns or by fewest open items; **Rules that route new work** assign new items by object, severity and other scope to a person or a team, first match wins. Teams and rules need manage triage, settings or rules.
- **The title and scope**: "Work" and the objects in scope, such as "All objects".
- **The banner**: the small line says "Open work", the headline counts it ("26,600 records open") and the work sentence follows. Two figures: **Critical** ("fix these first") and **With a suggested value** ("ready to accept"). The one button depends on your role and the list. A steward who can accept fixes sees **Accept 112 suggested values on MM142**, which accepts the suggested values of the rule with the most of them after you confirm. The button names the rule, so the count is that rule's, not the

3,418 suggested across all rules. An approver with rows waiting sees **Review the 12 waiting for approval**. Otherwise the button is **Start with the 2,904 records failing MM003**, which filters the ledger to the rule with the most open records.

- **Rules failing the most records**: under the banner, the eight rules with the most failing records, each with its description, the object, severity, failing count, the share of checked records that fail, a sparkline of failing records over the last runs, how many have a suggested value, and the object owner. Narrow Work to one object and the board shows that object's rules.
- **Fix pipeline**: records at each step from to fix to waiting for approval, approved, waiting for the SAP load, sent to SAP and fixed, then the same split per object.
- **The breakdown chart**: **Open records by object**, a horizontal bar per object (the eight with most open records), all on one scale from zero, with "1,420 open" at the end of each row. With an object in scope it becomes **Open records by rule** (the six worst rules, each rule id beside its message). Select a bar to filter the ledger to it. Beside it, **Cleaned and watched** shows the records fixed this week and how many failed again (see the Post-clean monitoring chapter).
- **Aletheia**, docked on the right: ask about the open work in plain words; every number in the answer is checked against the run's facts (see the Ask Meridian chapter).
- **The sentence**: what the current filter holds, for example "25,712 records to fix, ranked by value at risk." and, under it, "Rows are read from the latest run of each object. Value at risk is from cost model ZAR."
- **The state chips**: **To fix**, **Fix proposed**, **Waiting for approval** and **Exceptions**, each with its count. The chips are the filter.
- **The Policy chip**: narrows the ledger to the records failing a rule that evidences an applied policy pack or one of its classes, for example "Policy: Vendor bank data (SOX financial controls)" (`policy=sox:vendor_bank`).
- **The controls**: **Flat** or **By rule** grouping; **Sort** by value at risk, due, age or owner; **Compact** or **Standard** rows; **Views** (saved and shared views); the object filter; **Export fix sheet** (e) and **Import filled sheet**.
- **The rows**: by default the record key (for example Material number and Plant), name, rule, severity, **Current value**, **Suggested value**, **New value** and **Decision** beside the rule, then SAP field (`MARA-MEINS`, written TABLE-FIELD as in SAP), problem, owner and value at risk with its **est.** marker. The field description, possible values, first seen and days open are in **Columns**. Numeric record keys show as SAP shows them (`100` for `0000000000000100`), the full key on hover.
- **The fix-order hint**: where an upstream object still fails, a line says so first, for example "Fix Material master first: 3,640 upstream failures. Records there fail before this object can pass."

As of run reads the ledger as it stood when an earlier run finished: each object at its newest run up to then. A band says "As it stood at run 2 Oct 2026, 05:09; statuses are today's." with **Back to today**. Grouping by rule and the fix sheet read today's rows only, so they are off while a policy or a past run narrows the ledger.

The ledger loads 100 rows at a time ("Showing 100 of 25,712. Sorted by value at risk within the rows loaded.") with **Show 100 more**.

The numbers and where they come from

- **The chip counts** are distinct records in each state in the latest run of each object in scope, whatever the rows loaded.
- **Value at risk** is the cost-model estimate of what the failing value costs, in the tenant currency, marked **est.** (or **default estimate** while the shipped assumptions are in use).
- **Possible values** are the values the rule accepts; open the row for the full list with meanings.
- **Suggested value** is the value the rules propose. "Blank when no single value is right."
- **Days open** counts from the day the oldest failing issue of the row was first seen.

Editing a row

Type in **New value** and press `Enter`; until you do the cell reads "Not saved yet: press Enter". Or choose a **Decision**: **fix** (accept the suggested or entered value), **not an issue** or **accept risk**. A record in the inspector shows every failing field with the same cells, plus its facts, feature impact, provenance and history. Its **History** tab lists every run that measured the record with the values of its failing fields in that run, then every decision on it in order, with who and when, each linked to its run. On a phone, editing happens in the inspector.

Open full record in the record inspector opens the record on its own page, `/o/{system}/{object}/records/{key}`, for example `/o/ecc-production/material-master/records/100`. The title is the record key with its name and status under it. **Findings** lists each failing rule with its severity, message, state and suggested value. **Facts** groups every field of the record by its SAP table, with the field description, the SAP field, the value ("blank" when empty) and the rules that reject it. **Change log** lists every run and decision on the record, and **Possible duplicate** links to open duplicates and the golden record when one is named. Decisions stay in the inspector.

On a phone

Under 1100 pixels wide the rail becomes seven tabs along the bottom of the screen (Home, Objects, Runs, Fix, Rules, Master data and Migration), and each opens its menu as a sheet; Work is under Fix, and the inbox stays in the top bar. At 760 pixels wide and narrower the ledger's chips, controls and filters in force move into a sheet: **Filters** (with the count of filters in force) opens it from the bottom, and **Done** or a tap outside closes it.

Every change shows a notice at the bottom of the screen, such as "Accepted 1 suggested value", with **Undo**. `⌘Z` takes the last change back while focus is not in a field: "Undone: accepted 1 suggested value". The chip counts move at once.

Working in bulk

- **More than 25 rows at once**: a change of more than 25 rows, from a key (`a`, `n`, `r`) or an **Accept all** button, waits on a confirmation bar with the count, for example "Accept 112 suggested values for MM142? Enter confirms, Escape keeps the rows as they are."
- **Select rows** with the checkboxes, `x`, or `Shift+J` and `Shift+K` to extend. The bulk bar offers **Value for all selected rows**, accept the suggested values, not an issue, accept risk and **Assign…**.

- **Fill down**: select the rows, then press ⌘D in the New value you want to copy: "Filled EA down 12 rows".
- **Paste a column** of values from Excel into a New value cell to fill the selected rows.
- **Accept all suggested** in a rule group, or the next-step button "Accept 112 suggested values on MM142".
- **By rule** groups the rows to fix under their rule, so you can decide a whole rule at once.

Accepted and entered values go into your draft fix batch. They wait there for a second person: see the Approvals and exceptions chapter.

Export and import

Export fix sheet (e) saves the rows on screen as an Excel sheet with the same columns, in business words. Fill in New value and the decision offline, then use **Import filled sheet**: "Drop the filled fix sheet here, or choose a file (.xlsx). Meridian checks keys, SAP fields and values before anything changes." The result lists what was accepted and what was refused, row by row.

Share and export, above every grid, also offers **Export to Excel** (the visible columns, as a branded workbook), **Save these columns as a view** (shared with everyone if you choose) and **Share this page as a link** (read-only, for people of your tenant, until it expires).

What you can do, by role

- **Data steward** (apply): edit values, accept suggestions, bulk actions, import, assign to yourself.
- **Data owner** (approve): sign records off as not an issue or accept risk, approve batches.
- **Programme lead** (assign): assign rows, balance work from People.
- **Everyone**: open the full record page and the Monitor view.
- **Auditor, viewer, analyst**: read and export; the ledger says "Read-only for your role".

Keyboard

- `j` or ↓, `k` or ↑ move; `Shift+J` and `Shift+K` extend the selection; `x` selects; `Enter` opens the row.
- `a` accepts the suggested value on the selected rows; `n` marks them not an issue; `r` accepts the risk.
- `e` exports the fix sheet of what is on screen; `f` moves to the filter chips.
- ⌘D fills down; ⌘Z undoes; `[` and `]` step through rows in the inspector.

When something is wrong

- **"Sandbox mode is on: decisions apply to the sandbox run only."**: the tenant is in sandbox mode and the row belongs to the real run. Exit the sandbox under Administration, Settings, or decide on the sandbox copy.
- **"The work list of the latest run is not built yet, so no rows can be listed. It is built after each run finishes."** Wait for the analysis to finish, or ask whoever started it.
- **"No rows in this view."** The filters exclude everything. **Reset to the default view** from Views.

- **"Select rows first: x selects the highlighted row, Shift+J extends the selection."** A bulk key was pressed with nothing selected.
- **"The file could not be read as a fix sheet."** Export a fresh sheet and keep its columns and key fields unchanged.
- **A value you typed is refused**: it still fails the rule. Pick one of the possible values.
- **A record you fixed is back in To fix**: the next run still found it failing after it was sent to SAP. Check the load reached the right client, then decide again.

Approvals and exceptions

Every value decided in Work waits in a fix batch until a second person approves it. That is the four-eyes rule, and it has no bypass: whoever drafted a batch cannot approve it. Exceptions are records signed off without a change, as **not an issue** or **accept risk**, with a reason. Both live in the Work ledger, under their own chips: **Waiting for approval** and **Exceptions**. Open approvals from the command line by typing "approve", or at `/work?filter=approvals`.

Waiting for approval

The sentence reads, for example, "612 records wait for approval: the values are in a draft fix batch that a second person approves." The rows are grouped by batch. Each batch heading names the batch (for example **B-241**, "Unit of measure and material type"), who built it, the records, the value at risk, the due date and the state, with **Approve** and **Send back** on the heading.

Open a batch to read it in the batch inspector:

- **Facts**: batch, name, status, owner, steward, created, due, records and value at risk.
- **Records**: each change as the SAP field and the pair of values, for example "`MARA-MEINS` STK to EA". The first 50 show in the inspector; **Open all 328 in the ledger** opens the rest.
- **Approval trail**: who drafted it and who proposed the values, with times.
- **Actions**: **Approve**, **Send back** (with a reason the steward reads), **Download sheet** and **Reassign…**. "Approving acts on the whole batch. You can take it back from the notice that follows."
- **Watch** and **Comments**.

After approval

Approval does not touch SAP. It is Meridian's record that the values are right. The approved batch is exported as a load file for your SAP process, and records then read **Sent to SAP**. Where live write-back is switched on for the deployment (`MERIDIAN_LIVE_WRITEBACK_ENABLED`, material master, customer, business partner and vendor on ECC, S/4HANA on premise and EWM), an approved change can go to SAP through standard BAPIs instead: each record is simulated first and committed only when SAP's RETURN messages hold no error, those messages are kept per record, and each request needs someone other than the requester. A vendor goes live only where SAP holds it as a business partner with the supplier role (always on S/4HANA, on ECC where customer/vendor integration is active): its name, search term, address and bank account holder go through the business partner, and Meridian asks the system per record before anything is sent. Vendor company code and purchasing data, every vendor on a release without that integration, and every other object go out as a load file, each with its reason. Either way the work is not finished until a run says so: the next finished run reads the records again, and those that pass become **Fixed**. A record that still fails returns to **To fix**.

With live write-back switched on, the inspector of an approved batch (open it from **Work** or the inbox) carries a **Live write-back** section, and its load file button reads **Export load file**. **Check with SAP** comes first and sends nothing: a table lists each field of the batch with the record, whether it goes by **BAPI** (and which BAPI) or by **Export**, and the reason a field cannot go by BAPI. For vendor fields the check asks SAP which vendors are business partners there, so a vendor field reads **Export** with the reason when the system cannot take it. Then **Send to SAP**, for an approver who did not build the

batch, asks once with the number of records and starts a run. The request returns at once; a background worker writes the records in chunks and saves the result of each chunk before it starts the next, so a large batch does not time out and nothing SAP has committed is lost if the service restarts. A second **Send to SAP** while a run is open shows that run instead of starting another.

While the run is open the section refreshes by itself. It shows a progress bar and the counts of records committed, failed, to check in SAP and still pending, plus how many fields stay for the export. Each failed record is listed with its key and description, its state and the RETURN messages SAP gave, each with type, message class, number and text. When the run is done the batch reads **Sent to SAP** and the audit log holds one event with the counts; fields without a BAPI still go out with **Export load file**.

A run stops by itself with a readable reason when the connection to SAP fails several times in a row; it is then **Paused**. Fix the connection and press **Resume**. A run that shows no progress for a long time is marked stale and can be resumed the same way. **Resume**, **Cancel run** and **Retry failed** need the approve permission and, like the send, someone other than the requester. **Retry failed** queues only the records SAP refused. A record that was being sent when a worker died is never sent again on its own, because Meridian cannot know whether SAP committed it: it reads **To check in SAP**. Check the record in SAP, then press **Checked in SAP, send again**, which queues the refused records and the ones to check; a record already committed in an earlier run of the batch is never written twice.

The record statuses, in order: **To fix**, **Fix proposed**, **Waiting for approval**, **Approved**, **Sent to SAP**, **Fixed**. Exceptions are **Not an issue** and **Accept risk**.

Exceptions

The sentence reads "48 records are signed off as not an issue or an accepted risk." The columns change to **Resolution**, **Reason**, **Signed off by** and **When**, for example "Accept risk: Replaced by the new plant in wave 2; the risk is accepted until then". **Reopen** "Puts the record back to rows to fix. Needs the apply permission."

The numbers and where they come from

- **The chip counts** are distinct records in each state in the latest run of each object in scope: "612 records wait for approval" counts records, not batches.
- **Records** and **Value at risk** on a batch are the records in it and their cost-model estimate, in the tenant currency.
- **Due** comes from the SLA days for the batch's worst severity, set under Triage and SLA.
- **Fixed** is set only by a run: a value waiting for approval, approved or sent to SAP never counts as fixed.

Change requests

Requests, under Fix on the rail, lists change requests: a request to create a material, customer, vendor or business partner in SAP, or to change one. The views are **Mine**, **To approve** and **All**, with filters by status, object and kind. Each row opens the request on the right (`?open=request:{id}`).

The banner on top says "Change requests". Its headline counts what waits ("3 requests wait for approval", or "No request waits for approval") and the line under it says whether any failed ("1 request failed to reach SAP and need a new look." or "Nothing has failed in this list."). Three figures: **Waiting**

("submitted or in approval"), **Failed** ("did not reach SAP") and **Drafts** ("not yet submitted"), all counted over the requests in the current view and filters. The one button is **New request**. Under the banner, **Requests per person** is a horizontal bar per requester, the eight with most requests, with "4 requests" at the end of the row. Aletheia is docked on the right.

- **New request**: choose the object and the kind, give the SAP key when you change a record, a title, and the fields, one `TABLE.FIELD=value` per line. **Check** saves a draft and runs the rule pre-checks and the duplicate search: failed rules link to the rule and likely duplicates to the golden record, with how alike they are. **Submit for approval** sends it into the approval workflow once the check shows no problem. Needs the master data write permission.
- **The request inspector** shows the proposed values, the checks, the approval steps with who decided each and when, the history, comments and attachments. The requester can cancel a request that is not finished; a request that failed in SAP can be retried once you have checked in SAP that the record does not exist; a request with a load file can be exported.
- **Your decision**: when a step waits for you, the inspector of the request shows **Approve**, and **Reject** with a reason the requester reads. The step, its due time and whether you decide on behalf of someone are stated.

Ageing and due time

An open request shows how long it has waited. The list has **Age** and **Due** columns: the age in days or hours, and the due time of the approval step that waits. The banner adds an **Overdue** figure, and when a request is past the due time of its waiting step the banner turns to no-go and says "2 requests are past the due time of the waiting approval step." The inspector states **Waiting for** (the step), **Due** and **Age**. The due time comes from the hours set on the step in the approval workflow.

Field history

The inspector has a **Field history** table. Every field that was set, changed, removed, defaulted or derived on the request is kept with the old value, the new value and the person who did it, with the field description beside each code. A request without a field change says "No field change is recorded for this request."

Entry defaults

Entry defaults on the Requests screen lists values a new **create** request gets in the fields its requester leaves empty. A change request is never defaulted. Each default is a field as `TABLE.FIELD` and a value; the history shows them as **Defaulted**. Adding (**Save default**) or removing a default needs the manage workflows permission.

Mass change

Choose the kind **Mass change** in **New request** to change the same field or fields on many records in one request, up to 10,000. Paste the record keys, one per line; a bare number takes the key field of the object (`?keys=` in the address fills them in). Give the field and the value. The **preview** shows, per record, the current value from Meridian's own data beside the proposed one, with the counts found, not found, would change and already hold the value; it shows the first rows only. The request then goes through the same approval steps and permissions as a single change. On approval the history holds one row per record and field that changes, and a load file is delivered. Nothing is written to SAP by Meridian, and no record data goes to an AI.

Derivation rules and critical fields

The button **Derivations and critical fields**, beside Entry defaults, opens two lists; every code shows its description. Adding and removing needs manage workflows.

- **Derivation rules.** A rule says: when a field of the request holds one of the listed values ("Is one of", separated by commas), set another field. The new value is either a **constant value** or a **copy from field**: the request's proposed value of a source field, skipped when that source is empty. For example, when `MARA.MTART` is `ROH`, set `MARC.BESKZ` to `F`. Rules run in the order they were made when a request is created or updated, before the check, so a later rule sees what an earlier one set. A derived value is authoritative and shows as **Derived** in the field history.
- **Critical fields.** A list of fields per object. A request that touches one is flagged critical in the list and the inspector and gets an extra **Critical field** approval step after the usual steps, decided by a second person who is not the requester or an earlier approver. Administrators can also use the critical fields of a request in the conditions of their own workflow steps.

Out of office and escalation

- When an approver is away, their approvals go to the deputy they named under Preferences, Out of office. The batch inspector says so, for example "You cover for T. Mokoena until 30 Oct 2026: their approvals and assigned work come to you as well." Four eyes still holds: nobody approves their own batch.
- A batch past its due date is escalated by the escalation chain under **Triage and SLA** (see People and inbox).

What you can do, by role

- **Data owner** (approve): approve or send back batches you did not draft; sign records off as exceptions.
- **Data steward**: draft batches, read their state, reopen exceptions; a steward who also holds approve can approve other stewards' batches, never their own.
- **Programme lead**: approve, reassign batches.
- **Auditor**: read batches, trails and exceptions as four-eyes evidence; export them.

Keyboard

- `⌘K`, then "approve", opens the approvals.
- `j` and `k` move, `Enter` opens the batch or the row; `n` and `r` sign the selected rows off.
- `⌘Z` undoes an approval or a sign-off from the notice.

When something is wrong

- **"No value is waiting for approval. Accepted and entered values land here until a second person approves them."** Nothing is pending.
- **Approve is missing on a batch**: you drafted it, or your role does not hold approve. Ask another approver.

- **"Sent to SAP" records come back as To fix**: the load did not take. Check the file was loaded completely, into the right client, and that the run you read finished after the load.
- **A sign-off is missing its reason**: the reason is required when you sign off; older records may show "The API does not send the reason yet".

Post-clean monitoring

Cleaning a record is not the end of the job. Monitoring, under Work, shows whether records that a batch fixed stay fixed, and which rules keep pulling them back. Open it from the Monitor tab at the top of the Work ledger, or from **Open monitoring** on the **Cleaned and watched** card of the landscape and Work.

Cleansing

The Cleansing page, under Fix, Cleansing, shows what the cleaning engine did. The banner gives the records cleansed in SAP and rolled back at the latest weekly snapshot. Under it:

- **Proposals to decide.** The proposals waiting for a decision, grouped by object and fix, with the number of records, the average confidence and how many are at 85% confidence or more. Users who may approve can approve the confident ones of a group in one step, or open a group and approve or reject each proposal. A rejection needs a reason, which is kept in the audit trail and sent back to the AI as feedback. **Review the proposals** in the banner jumps here.
- **Duplicate records to merge.** Each pair of records the cleaning engine found to be the same, with the object, the fields that matched and the match score. Open a pair to compare the two records field by field, with the value the merge would keep. Users who may approve choose the record to keep, give a reason and press **Merge**; the reason is kept in the audit trail.
- **Approved, waiting to be written to SAP.** Each approved proposal with the change it makes. Users who may apply fixes press **Write to SAP**; the change then appears under the changes you can still undo.
- **Where every proposal stands.** One bar split by status: waiting for a decision, approved, cleansed in SAP, rolled back and rejected.
- **Records cleansed, week by week.** The cleansed count at each weekly snapshot. The snapshots are counts at a point in time, so the weeks are never added together.
- **Cleansed by object.** Records in SAP per object, with the share of that object's proposals.
- **Applied changes you can still undo.** Each change applied to SAP, with the fields it changed (before and after) and the deadline. Users who may apply fixes can roll a change back for 72 hours; after that the button is disabled.
- **Audit trail.** The last 25 decisions with who made them. Roles without the export permission see a dash instead.

Data mining

The Data mining page, under Master data, Data mining, shows what the mining tasks found by comparing records and values rather than by running the rules. Every list keeps each object at its latest complete run, so it is the current state. The banner counts the likely duplicate pairs and outliers, with the patterns of the last 30 days, the share of analysed objects mined and the runs mined beside it. Under it:

- **Findings by object.** Duplicate pairs and outliers per object, most first.

- **Patterns.** The outliers grouped by field and the duplicates grouped by object, with severity, how often each was found and the confidence.
- **Likely duplicates.** Each pair of records with the match score and the fields whose values are the same, each with its description.
- **Outliers.** Each value far from the rest of its field, with the usual range or how many standard deviations it sits from the average, and how it was found.
- **Fields that depend on each other.** Pairs of fields whose values move together or where one decides the other, strongest first. These are candidates for a cross-field rule.

Users who may change master data can press **Mine the latest runs** to mine the latest run of every object again. A role that cannot read the mining results sees a dash in each list.

Anomaly rules

The Anomaly rules page, under Rules, Anomaly rules, turns an outlier that mining found into a proposed rule. It lists each field that has outliers, with the field description, how many outliers it has and the usual range. Users who may manage rules press **Propose a rule** on a field. The page then shows:

- **The proposed rule.** A deterministic rule that flags values outside the usual range of the field. It is a template, so no AI is used. The AI drafts a rule only when no template fits, and then it sees the field, the statistic and counts, never a record.
- **The evidence.** The outliers found, how many records the rule would flag when it runs on the same data, and the confidence.

Press **Save as draft** to keep the rule as a draft in the rule workbench, where it is reviewed and published as any other draft.

Null spikes and unexpected codes

Under the outliers, the panel **Null spikes and unexpected codes** compares each run with the run before it of the same system and shows counts only, never a record. Meridian reads it from the field statistics of the runs (row count, blank count, distinct values and the most common values of code fields); no AI is used. Each row shows the field with its description, the kind, the runs compared ("Run 12 against run 11"), the figures before and now, and the severity. By default the panel shows the latest run of each object that has shifts.

- **Null spike.** The share of blank values in a field rose by at least 10 points against the earlier run, and both runs hold at least 100 rows. The severity is medium, high from a rise of 25 points and critical from 50 points. The figures read "12% empty" before and now.
- **Unexpected code.** A code appears in the field that none of up to five earlier runs listed. Meridian judges this only when the current run and the one before it list every value of the field, so a new code is certain. Fields with many distinct values, masked fields and sensitive fields get no such check, because their statistics do not list every value. The severity is high when the new codes make up 5% or more of the filled rows, otherwise medium. The figures read "3 of 48,000".

A run with no earlier run to compare gets no shifts, and earlier runs are not checked afterwards. When nothing shifted, the panel says "No null spike or unexpected code was found."

Press **Propose rule** on a row to draft a rule from the shift. A null spike drafts a null check on the field. An unexpected code drafts a domain value check on the values the earlier run listed, and only when it listed every value. The draft is saved in the rule workbench like **Save as draft** above, so nothing is enforced until a person reviews and publishes it. Pressing it again for the same shift does not create a second draft; the page says "This draft already exists. Find it in the rule workbench."

Learned monitors

The Learned monitors page, under Rules, Learned monitors, compares the latest run of each object with that object's own earlier runs. It needs no setup: material master, customer master and SuccessFactors get it as soon as they have runs. Every figure comes from the stored run results, so no record value is read and none is sent to a model.

Each object keeps a baseline of up to 30 earlier runs of the same system for these figures: records checked, score, and the fail rate of each rule. It also keeps the empty share and the distinct share (distinct values as a share of the filled rows, for fields with at least 100 filled rows) of each field that has field statistics. The usual range is the median of those runs plus or minus a robust spread (the median absolute deviation), so one bad run does not move it much. When an object has at least 28 earlier runs, the baseline also compares like weekdays with like. The table **What is learned** shows, per object, the latest run, how many earlier runs the baseline uses against how many it needs, and whether the object is **Watching** or still **Learning**. An object with fewer than 5 earlier runs only learns: it raises nothing and shows a dash for breaks.

After each analysis run Meridian checks six kinds of break and stores each with its severity, the usual range, the observed value and the runs behind it:

- **Volume drop** and **Volume spike.** Records checked fell below or rose above the usual range.
- **Fail rate jump.** A rule failed on a larger share of records than usual. The row shows the rule code with its description.
- **Score drop.** The score fell below the usual range.
- **Field emptied.** A field that is normally populated is now empty or far emptier than usual. The row shows the field code with its description.
- **Distinct values shifted.** The share of distinct values among the filled rows of a field moved by at least 20 percentage points either way, for example a code field that collapsed to a few values or a key field that began to repeat. The row shows the field code with its description. Only counts are compared; no value is read.
- **Value shapes shifted.** The share of filled rows whose shape falls outside the field's usual shapes rose by at least 10 percentage points. A shape is the layout of a value (letters count as A, digits as 9), and the usual shapes are those that hold at least 5% of the rows in at least half of the earlier runs. A field whose usual shapes cover under 80% of its rows (free text, keys) is skipped, and so is a field with fewer than 100 filled rows. Only the profile's five commonest shapes are read, so a field with many layouts is not judged.
- **Format broke.** The share of values that do not parse (numeric fields) or are not valid dates (date fields) rose by at least 5 percentage points. It is judged for any field with at least 100 filled rows, including personal fields, because it is a count.

- **Range shifted.** The mean of a numeric field moved by at least 50% of its usual level in either direction, the largest value of a numeric field rose or its smallest value fell by that much, or the latest date of a date field fell back, or the earliest date moved earlier, by at least 365 days. The row names the figure that moved (**Mean**, **Largest value**, **Smallest value**, **Latest date**, **Earliest date**) and shows how far it moved, as a percentage or in days; it does not show the value. One field raises at most one range break per run, the strongest.

Limits of the three. A latest date that moves forward as new records arrive is normal and never a break, so a dummy far-future date such as 31.12.9999 is not caught by **Latest date**; it does not parse as a date and shows up in **Format broke** instead. A profile taken from a sample of the rows gives no largest, smallest, latest or earliest figure, because a sample's extremes are not the table's; the mean and the shares are still watched. Personal and sensitive fields, and fields masked for privacy, are never watched for shapes or ranges, only for the share that does not parse. No value, shape, mean or date of a field is ever sent to a model: the wording is built from the share or the size of the move alone.

The panel **Breaks in the latest run** lists them, most severe first. Press a break to see **Why Aletheia flagged this**: a sentence built from the figures by a fixed template, and the runs it rests on ("Run 12 (this run)"). Press **Ask Aletheia to explain** to have the model reword it. Aletheia sees only the aggregate figures and the rule or field description, and Meridian drops any sentence that states a number outside those figures. When the model is not available, the fixed wording stays.

Between full runs

A full analysis takes minutes to hours, and its figures (score, fail rate, empty share, records checked) exist only after the checks have run. One figure is known earlier: the number of rows a sync extracted, which Meridian stores as soon as the extraction lands. Every 15 minutes it compares that count with the row counts of the earlier synced runs of the same object, with the same baseline and the same learned minimum as the run monitors. The panel **Between full runs** lists a **Volume drop** or **Volume spike** found this way, with **Rows extracted**, the usual range, **Earlier syncs used** and the status **Analysis pending**. Press one to read the explanation, which states that the break comes from the extraction and not from a finished run. Steward decisions work as for any other break.

The check is limited on purpose:

- It covers volume only. Score, fail rate and empty share cannot be known before the checks run, so nothing is shown for them.
- It covers full extractions of a profile with one module. A delta or sampled extraction does not carry the size of the object, and a profile with several modules sums their rows, so both are left out.
- It looks at the latest sync of the last 72 hours whose analysis has not finished. When the analysis finishes, the run monitor judges the run on records checked and replaces the interim row.
- It reads stored counts only. It starts no extraction and does not run while the analysis is working.

Users who may change master data can decide on a break:

- **Acknowledge.** You have seen it. The baseline does not change.
- **Mark expected.** Choose **Leave this run out of the baseline** if the run was a one-off, or **Accept as the new normal** if the new level is right; the baseline then restarts at that run.
- **Snooze.** Hide the break for 1, 7 or 30 days. When the time passes and the break still stands, it reads as open again.

Each decision is written to the audit log with the person who made it. Other roles can read the page and ask for an explanation but see no decision buttons.

Learned minimum effect

A break is raised only when its effect reaches the minimum effect of that monitor. A monitor is one object, one figure and one rule or field. The effect is the share of the median for records checked, points for score, percentage points for fail rate, empty share, distinct share, shape share and unparseable share, the share of the median for the mean, largest and smallest value of a number, and days for the latest and earliest date. Each figure has a fixed floor: 30 percent of the median for records checked, and fixed point values for the others (45 points for empty share, which is exactly where the monitor fires today, 20 points for distinct share, 10 for shape share, 5 for unparseable share, 50 percent for number ranges and 365 days for date ranges).

The panel **Learned minimum effect** shows, for every monitor that decisions touched, the **Minimum now**, the **Fixed floor**, the **Ceiling** and **Decisions that moved it**, newest first, each with the run, the person and the effect.

- **Mark expected** raises the minimum to 5 percent above the largest effect marked expected in the last 90 days.
- The ceiling is 3 times the floor, so a monitor is never silenced. A break above the ceiling is always raised.
- **Acknowledge** on a break smaller than the current minimum lowers the minimum to that effect. It never goes below the floor.
- Decisions older than 90 days stop counting, so a raised minimum relaxes by itself.
- Severity cut points scale with the minimum divided by the floor, at most 1.5 times, so a total loss of records stays critical however high the minimum.

The panel **Below the learned minimum** lists breaks that are outside the usual range but smaller than a raised minimum. They raise no alert, send no notification and are not counted as open. **Confirm as real** records the decision as an acknowledgement, and the minimum drops to that effect, never below the floor. A later run replaces the list.

Limits. The minimum is one number per monitor, learned from decisions alone. The distinct-share and shape monitors compare counts and shares only: they cannot tell which values changed. The shape monitor reads only the five commonest shapes of a field, and the range monitors need an unsampled profile for the largest, smallest, latest and earliest figures. The rule is a fixed calculation on aggregates, not a trained model, and a language model takes no part in it.

Entry gate

The Entry gate page, under Master data, Entry gate, shows the data quality firewall: the check SAP calls before it saves a master data record. It covers the licensed objects, Material master, Customer master (sales) and SuccessFactors employee data. The page shows only what the firewall's audit trail holds. Where the trail did not record a figure, the page shows a dash and not a zero.

- **Records checked and blocked.** Two charts of the last 7 days, with a table of each day. A record is blocked when it fails a critical rule. A record that passes but fails a lesser rule is warned. The warned and passed counts were not kept before this release, so earlier days show a dash.

- **Rules that block most.** The critical rules that stopped the most records, each with its code and what it requires.
- **By object.** Records checked, blocked, warned and passed for each licensed object.
- **Try a record.** Paste records as JSON, choose the object and press **Check records**. The page runs the same check SAP calls. Tick **Show a fix where the rule gives one** to see a suggested fix. A fix appears only when the rule itself defines one, either a corrected value from the rule's auto-fix steps or the values the rule accepts, and each fix names its rule. When no fix is grounded in the rule, the cell shows a dash. No language model sees the record.
- **Send to a steward.** Tick **Send blocked or warned records to a steward as a change request** before you check. Each blocked, warned or duplicate-flagged record then opens a create request and sends it to the approval workflow, and a table shows the request's status and an **Open the request** link for each record. Checking the same record again shows the request that is already open and adds no second one. When no request is opened, the note says why in words: SuccessFactors employee data has no change request, only the first 25 flagged records open one, the change request refused the record, or opening failed while the verdict still stands. The request carries the record values like any other change request, so avoid it for personal data you do not want held there.
- **Hold the SAP save.** Choose **Blocked** or **Flagged** in the **Hold** list before you check. It also sends the records to a steward, as above, and answers whether the SAP save may go on. Each covered record shows its hold state: **Held** while the request waits for approval, **Released** once it is approved (the save may go on, even for a blocked record, for 24 hours and once), **Rejected**, **Written** when Meridian writes or has written it to SAP or the load file was already taken, and **Not held** with the reason when no request could be opened. A line under the table says **The SAP save stops.** or **The SAP save may go on.** A released request cannot also be exported as a load file, and the Requests screen shows it as "Released to the SAP save". The sample SAP code for this has not been run on a live SAP system.
- **Holds panel.** Above the requests, the number held now, and released, rejected and written over the last 7 days, the age of the oldest held request, the median time from opened to decided, and the latest held requests, each with its object, status and a link. A dash or a sentence replaces a figure when nothing was held or decided.
- **Requests opened from the gate.** The latest ten requests the gate opened, each with its title linking to the request, its object and its status in words, and the date it was opened.

The page is a view of the firewall. The SAP side, the API key and the choice to let a save go on or stop when the check cannot run are described in the point-of-entry guide in the repository (docs/point-of-entry.md). The approval workflow runs in Meridian, not in SAP. The duplicate search at entry warns and never stops a save. It also finds a tax or part number that is one character different, or has two neighbouring characters swapped, and the page names the identifier field that matched under Matched on. A near value is a prompt to look: it can be another taxpayer or another part, and a difference of two or more characters is not found.

MDM health

The MDM health page, under Master data, MDM health, shows the daily snapshot of how well master data is governed. The banner gives the health score out of 100 and its change over the last 90 days, with the golden records, the open stewardship items and the share of systems synced beside it. Under it:

- **What makes the score.** The four shares the score is built from, each with its weight and the points it adds: golden record coverage (40%), match confidence (30%), steward SLA (20%) and source consistency (10%). The points add up to the score. A small line beside each share shows its last 90 days.
- **Health score, day by day.** The score at each daily snapshot on a 0 to 100 scale.
- **Stewardship backlog, day by day.** The open stewardship items at each snapshot.

The snapshot runs once a day. Until the first one has run the page says so, and the charts need two days of snapshots before they draw a line. Viewers and process analysts see no figures.

Record links

The Record links page, under Master data, Record links, shows which records are linked across objects: a business partner and the customer it is, a material and its plants, a vendor and its company codes. Relationship discovery reads these links from the SAP link tables at each sync, and the AI impact scorer adds links it infers. The banner counts the links, with the link types in use, the objects they join and the share that was inferred rather than read from SAP. Under it:

- **Find the links of one record.** Pick an object and type a record key to see every link to and from that record. Tick **Include ended links** to see links that no longer exist in SAP.
- **Links by type.** How many links of each type, most first. The list holds the latest 200 links; when there are more, the bars cover only those, and the banner says so.
- **Links.** Each link with the record on each side, what the link means, the SAP table it was read from (or the confidence of an inferred link) and the day it was found.
- **Link types SAP defines.** The link types with the SAP tables on each side.

Roles without permission to see AI confidence see a dash instead of the confidence of an inferred link.

Record issues

The Issues page, under Master data, Issues, is the steward work list. Each issue is one record failing one rule, carried from run to run with its status, the person it is assigned to, its due date and its history. The banner counts the issues still to close, with those in progress, those waiting on SAP or on the requester, and those closed. Under it:

- **Where the issues stand.** How many issues are in each status.
- **Status buttons and filter.** Show one status or all of them, then narrow by severity, by who the issue is assigned to (you, no one or anyone) and by record key. Users with the export permission can export the list to Excel with the same filters.

- **Issues.** Each issue with the record, the rule and the field it failed on, the severity, the status, the assignee, the due date with whether it is on time, and how often it has been reopened. The list is ordered most severe first, a hundred to a page.
- **Change the selected issues.** Tick issues to move them to another status, assign them or add a note. Accepting an issue needs the approve permission and a resolution (risk accepted or false positive); a false positive needs a reason. Resolving needs the resolution fixed in the source. Other moves and assignments need the assign permission.
- **One issue.** Open a record to see the runs that found it failing or clean, and who did what to it. Users who may analyse can add a comment.

Action centre

The Action centre, the first page under Fix and the page the Fix item opens, lists everything that needs you, most urgent first: fixes waiting for your approval, systems that cannot be reached, records failing again, overdue stewardship items, failed jobs, rules waiting for review and, for administrators, access requests and expiring keys. You see only what your role can act on. The banner counts the items by severity and says when the list was last checked; it refreshes every minute. Under it:

- **Filter.** Narrow the list to one kind or one system. The choice stays in the address, so a link to it opens the same view.
- **What needs attention, by kind.** How many items of each kind.
- **Act on these.** Each item with what happened, its severity, the system and object, when it started and how many records or items it covers. The button opens the screen where you resolve it.

Material 360

The Material 360 page, under Master data, Material 360, shows one material across every view and organisational level. Type a material number; leading zeros may be left off. The banner gives the description, the material type, group and base unit, with the rules it fails, the views missing and the plants phasing it out beside it. Under it:

- **Views by level.** A grid of the material views (basic data, sales, purchasing, MRP, storage, accounting and the rest) against each level that carries them: client, plant, sales area, valuation area and storage location. Each cell says whether the view is maintained, missing or not needed there, and marks it failing when a rule fails on that view at that level. A view the material type does not expect says so.
- **Rules failing, by view.** Each failing rule with its description, the view and level, the field, the value found, the severity and a link to its issue when one is open.
- **Follow-up chain.** For each plant phasing the material out, the chain of follow-up materials with their MRP type and discontinuation date, and what is wrong with each link: a successor that does not exist, is blocked or not planned, a loop or a dead end. Under it, the bills of material that still use the material.
- **Materials that look like it.** Other materials whose description, EAN or attributes match, with the match score. Open one to see it the same way.

If the material is not in the material master, or no run has data yet, the page says so.

Work analytics

The Analytics tab, third at the top of the Work ledger, shows how fast the steward work list is being worked down. Every figure covers each object at its latest complete run. The banner counts the issues still open, with the records that carry an open issue, the issues closed in the last seven days and the pace closed a day. Its sentence says when the list will be clear at that pace, or that it is not getting shorter. Under it:

- **Open issues, day by day.** The open issues at the end of each day, with those found and those closed.
- **What the latest run changed.** Per object, the issues the latest run found new, fixed and still failing against the run before, with the rules newly failing and the rules recovered.
- **Time to fix.** From found to fixed over the last 180 days: the median and the slowest tenth, by severity, by object and for the slowest rules.
- **Decisions per person.** Each steward's decisions over the last weeks, split by kind, with a small line for each day. Under it, the open work per assignee with the critical and overdue items.
- **Cleaning engine pace.** What the cleaning engine has found and applied over every day it has recorded: fixes applied a day, fixes still waiting, the share approved without review, the share rejected at review or rolled back after applying, and the average review time.
- **Issues that came back.** The rules and the systems whose issues reopen most after they were closed, with the rate.
- **Where the failing records sit.** The records with an open issue by organisational unit, such as plant or company code, with the share of each value.
- **How good the proposals are.** For each source of proposals (rules, the cleaning engine, steward edits and simulations), how many were approved and how many were still fixed at the next run. A source with nothing proposed shows a dash.
- **Rules stewards call wrong.** Each rule with the steward verdicts on its flagged records: how many were called a false positive, how many were confirmed real, and the false-positive rate, the rule called wrong most at the top. **Show why** lists the notes stewards left when they called that rule wrong, with the records behind each note and the date it was last given. A rule with a high rate is a candidate for tuning or for a scope change.

A role that cannot read a list sees a short note in its place.

Cost and causes

Cost and causes, under Migration on the rail, shows what bad data costs and where it comes from. Money is always shown per business area, such as procure to pay or order to cash, and never added across areas. The banner names the area with the most value at risk, with the largest value at risk, the largest proven cost, the SAP features at risk and the number of root causes. Under it:

- **By business area.** The proven cost and the value at risk of each area side by side.
- **Money already lost.** Costs measured in posted documents, such as late purchase orders or duplicate vendor payments, with the documents counted and the rules behind them. A cost in more than one currency lists each currency on its own.

- **SAP features at risk.** The features the failing rules block or degrade, with the records they touch, the value a record and the rules that cause it.
- **Root causes.** The failing rules grouped by the field they fail on, with the severity, the records, the share that can be fixed automatically and the program or user the bad values mostly come from. A figure that has not been worked out shows a dash; an estimated record count says "about".
- **Score run by run.** The score of the last runs, with the rules each run fixed and newly found.

Before the first scored run the page says so instead of showing figures.

Forecast and plan

Forecast and plan, under Migration on the rail, shows where each object's score is heading, what to fix next and whether the data stewards can keep up. The banner counts the objects below their floor and the objects heading below it, with the sprints needed to clear the open work and the score once they are done. Under it:

- **Where each score is heading.** Each object with its floor, its score now, its last four weeks, its trend and its forecast in 30 and 90 days, each with a 95% range. A forecast needs scores from at least three days; until then the page says how many days there are.
- **Next best work.** The failing rules, cleaning items and exceptions ranked by the severity-weighted records fixed an hour, with the records, the effort and the cost where a cost per record is set.
- **Sprint plan.** The ranked work filled into sprints of the stated hours, with the records fixed, the critical findings cleared and the score before and after each sprint.
- **Can the stewards keep up.** The stewards needed against the stewards working, the new items a day, the items one steward clears a day, and the queue stages where items wait longest. Items already applied or rejected are not counted. Until stewards have recorded work the page says capacity cannot be planned instead of guessing.

The effort and cost assumptions are shown above the sprint plan.

Value delivered

Value delivered, under Migration on the rail, prices what the latest finished run found. Each failing record is given the unit cost of the loss it can cause, taken from your cost model, and the mitigation share is what fixing it removes. The banner shows the annual risk found, the risk mitigated a year, the return on the subscription and the payback in months. When no subscription cost is set in the cost model the return and payback show as a dash with a note instead of an invented fee. Under it:

- **Where the risk sits.** The eight kinds of loss (duplicate payments, missed warranty claims, compliance penalties, inventory write-offs, blocked invoices, failed postings, contract violations and steward time saved), each with what counts towards it, how many findings or exceptions feed it, its annual risk, the amount mitigated and the formula used to price it.
- **Cost avoided month by month.** The monthly report records the subscription, the risk mitigated and the return to date on the first of each month. Until the first month is recorded the page says so.

Only the latest run is priced, so the figures never add the same failing records twice across runs.

Survivorship picks and their accuracy

Aletheia survivorship lists the golden-record winners the model picked and whether the server proved each one. A candidate is shown as its source and the shape of its value, never the value. A steward can act on a proved pick:

- **Accept the pick.** Records that a steward kept it. If the approval policy held the pick back, accepting applies it.
- **Override.** Choose one of the other offered sources, never free text, and optionally a reason: wrong source, out of date, incomplete, wrong format or another reason. The change goes through the steward override path, so the field stays pinned until you clear it on the merge explain screen.

Each decision stores who decided, when, the earlier source and the new source, and writes an audit entry. The accuracy panel shows how often stewards keep the pick: the accepted, overridden and waiting counts, the weekly trend of accepted picks and a bar per field. A pick that is ingested again becomes waiting again.

To require approval before any AI pick is applied, use the Steward approval of AI picks panel on the verified survivorship screen. It shows the current value and the tenant setting behind it, `survivorship_ai_approval`: `required` (the default) holds each proved pick in the approval queue until a different steward approves or rejects it, `not_required` is the opt-in auto-approve and applies it at once. Changing it needs the manage settings permission; without it the panel shows the value and says why the control is off. Each change writes an audit entry.

Reference data

Reference data, under Master data on the rail, manages the code lists of an ECC to S/4HANA migration: country, unit of measure, material group, payment terms and, for SuccessFactors, the picklist values. Lists come only from the check tables Meridian has already extracted. A list with no extracted check table is not shown, and a code with no extracted text shows the code alone.

- **Versions.** Each completed load of a check table is a version. Pick one to see the codes it held, and read the trend of codes per version.
- **Source codes.** Every source code with the target it goes to: the same code on both sides, a mapping, or no target yet. Each code sits beside its description.
- **Mappings.** A person with the analyse permission proposes a source to target mapping. A second person, never the proposer, approves or rejects it. Only an approved mapping reaches the migration. Each decision stores who decided and when, and writes an audit entry.
- **Ask Aletheia.** With the trigger AI permission, Aletheia suggests a target from code values and descriptions only. The server checks that the target exists in the list before showing it, and a person still has to approve it.

Post-load reconciliation

Post-load reconciliation, on a migration analysis (Migration on the rail, select an analysed run), compares the load files Meridian prepared with what the target system holds afterwards. Meridian does not run the load. You load the files with your own tool, extract the target system again, and then run the comparison. Until the target has been extracted, nothing is compared.

- **Extract the target system.** When the target has no completed analysis, the panel says the target system has not been extracted yet and offers **Extract the target system**. It needs the trigger sync permission and starts the extraction now, for the run's objects that the target system type supports. It reads whole business objects, not only the tables the load files carry. If an extraction on that system is already running, Meridian refuses and starts nothing: it never queues behind a running extraction or interrupts it. The run must name a target system. When the extraction has been analysed, the comparison starts by itself.
- **Delete an analysis. Delete** in the Analyses list needs the analyse permission and asks for confirmation. It removes the analysis with its gaps, load files and reconciliations and the stored copies of the files exported, and it cannot be undone. A running analysis cannot be deleted. If the stored copies cannot be removed, the analysis is kept so you can try again. Files you already downloaded are not covered.
- **Run reconciliation.** Needs the analyse permission. The run must be analysed, must name a target system, and must not be reconciling already. The comparison runs in the background and the panel refreshes while it runs.
- **Scores.** One row per object with three figures side by side: **Source**, **Projected** (what the analysis expected for the target) and **Actual in target**. The actual figure is the data quality score of the target system's own latest analysis. When the target has no completed analysis, or none for that object, the cell is a dash with the reason. Meridian never estimates it. Under the scores, **Rules that moved the actual score** lists up to ten rules whose pass rate differs between source and target, the largest change first, each with its rule id beside its description.
- **Records per target table.** For each target table: expected (records in the load files), found in the target extraction, missing, and held back but present. Records are matched on the table's dictionary key. A record the blocking gaps held back is reported only when it is present in the target. A target record that is not in the load files is not reported, because the target holds records from other sources. A table that cannot be judged shows **Not compared** with the reason code beside its text: no key in the data dictionary, key field not in the load files, target extraction incomplete or partial, table not extracted, or key field not in the target extraction.
- **Fields compared.** Per field: how many loaded values were compared with the target, how many match and how many differ. A field the load left blank is not compared, because the target may fill its own default. Values are compared after trimming and, for numeric keys, padding to the SAP internal form. A difference in number format or case counts as a difference.
- **Records.** Pick **Missing in target**, **Different in target** or **Held back but present**, and optionally one table, to see the records with the expected and the found value. The list keeps the first 100 records of each table and kind. The counts are complete, and the panel says when a list is cut. The rows hold the compared values, so they follow the access rules of the migration data. A personal field is counted like any other, but its values are not kept: the record shows **Withheld** in place of the expected and the found value.
- **Out of date.** When a target extraction finishes after the comparison, the panel says so with the date of the newer extraction. Run the reconciliation again to compare with it, or wait: when the target system finishes a new analysis, Meridian starts the comparison itself for every run whose

latest comparison used that system, unless one is already running. The panel then says **Started automatically when the target system was analysed again** and shows how the missing, differing and found counts moved since the previous comparison.

- **Download as Excel.** Needs the export permission. Sheets: Tables, Fields and Records.
- **What it was compared against.** At export, Meridian keeps a copy of exactly the files you handed over. The panel says **Compared against the files exported on** a date when it used that copy, so a mapping changed after the load is not a difference. When no copy exists, or it could not be read, the panel says the expected side was **Rebuilt from the current field mappings**, and then a mapping changed after the load does show as a difference. The copy holds record values, personal fields included. It is deleted with its run, with a version or tenant purge, and a data-subject erasure rewrites it without the person. Backups and files you already downloaded are not covered.
- **Columns cut from the extraction.** A table that lost non-key columns but has every record is compared on the fields present; the lost fields show as not in the target extraction. A table read for a date window or scope, or an older extraction recorded as partial, still shows **Not compared**.

Score lens

Score lens, under Migration on the rail, scores every run two ways. The first is the score as it was recorded when the run finished, with the weights in force then. The second re-weights the same run under today's weights, so a change of weights never reads as a change in the data. The banner shows the latest run both ways, how many points the weight change accounts for, and the config-aware score. Under it:

- **Score under today's weights, run by run.** A trend line over the runs, oldest on the left. The table under the chart gives both scores for each run.
- **Each run, then and now.** Each run with its score then, its score under today's weights, the change the weights make, and whether the weights were recorded with the run. Runs scored before weights were recorded are re-weighted from the default weights, and the table says so.
- **Each module on the six dimensions.** A heat map of the latest run: one row per module, one column per dimension (completeness, accuracy, consistency, timeliness, uniqueness, validity) and the score. Each cell is coloured by your score targets: on target, below target, or below the warning line. A dimension no rule measures shows a dash. The last column counts the critical rules failing, and a note under the map says which module scores are capped by a critical failure.
- **Each module across the runs.** One row per module with a sparkline of its score, the latest score, the score as recorded then and the change since the first run.
- **Config-aware score.** The latest run counted only over the rules that your loaded SAP configuration says apply. A rule for batch management does not count when batch management is not active in any plant. The panel shows when the configuration was loaded, the score per process and per module, and the reasons rules do not apply with how many rules each covers. When no configuration is loaded every rule counts by default and the panel says so.

What is on it

- **The banner.** One sentence at the top: how many records failed a rule again since the last run and how many were fixed, with both figures beside it. Aletheia is docked on the right, as on every main screen.
- **Per system.** One row per connected system with the score then (the run before the clean), the score now, the change, how many records have failed again and how many were fixed since.
- **Failing again after a clean, by week.** A weekly bar chart over the last twelve weeks: records that failed again after a batch fixed them, set against records that were fixed again.
- **Drift by rule.** The rules whose records fail again most often, with the count for each. Open a rule to see which records came back and why.
- **Still clean, by batch.** Each approved batch with the records it fixed, how many are still clean and how many have failed again.

The numbers and where they come from

Every figure comes from the record issue history that each run writes, joined to the remediation items of approved batches. A record counts as failing again when a run after its batch finds the same rule failing on it. Nothing is estimated: if no run has followed a batch yet, the screen says there is nothing to monitor.

What you can do, by role

- **Data stewards and owners** open the drifting rules and send the records back to the ledger.
- **Programme leads** watch the weekly chart to see whether cleaning holds before a wave gate.
- **Auditors** read the still-clean table as evidence that approved fixes persisted.

Keyboard

- ⌘K, then "monitor", opens this tab.
- `Tab` moves between the tables; `Enter` on a rule or a batch opens it.

When something is wrong

- **"Nothing to monitor yet."** No run has finished since a cleaning batch was loaded. Run the system once the load is in SAP.
- **A rule keeps coming back**: the source process creates the bad value again. Fix the process or the interface that writes the field, not only the records.
- **"The monitor could not be read."** The API did not answer. Check the system health, then reload.

People and inbox

People lists the owners and stewards with their open work, approvals waiting, overdue records and SLA, and is where work is balanced, contracts are read and out of office is set. The **inbox** holds what Meridian tells you: approvals, SLA warnings, breaches, finished runs, mentions and digests. Open People from the map, with `g` then `p`, or at `/people`. Open the inbox from **Inbox** and its count in the top bar.

People

The banner on top says "Overdue work" and answers it: "2 people have overdue work", or "Nobody has overdue work". The line under it names who holds the most, for example "Thandi Nkosi has the most, 21 records past the due date.", or says "Everything assigned is still inside its due date." Three figures: **Open records** ("held by 4 people"), **Overdue** ("past the due date") and **Unassigned** ("nobody holds them"). The one button, under the sentence at the top of the page, opens the work of the person with most overdue ("Open Thandi Nkosi's overdue work"), with who holds the most beside it; when nothing is overdue there is no button.

Open work per person is a horizontal bar chart of the ten people with most open work, each with name and role and "48 open, 21 overdue" at the end of the row. "Red bars belong to people with overdue work." People carries no run, so Aletheia is not docked here and the table takes the full width.

The sentence sums the team up, for example "4 of 6 people hold open work: 270 records, 34 overdue. 2 batches wait for your approval. 26,330 of the 26,600 open records are assigned to nobody." and "Counts are read from the work assigned in the latest state."

The table: person and role, objects (owned or stewarded), open work, approvals waiting, overdue, **SLA met**, breaches this month and last action. Your own row reads "you". The table ends with an **Unassigned** row ("Assign them on Work to give each an owner and an SLA.", linked to Work's unassigned filter) and a **Total** equal to the work list's open records.

Above it: **Triage and SLA**, **Balance work**, **Data contracts** and **Out of office**.

The person inspector

Open a person for their facts (role, objects, open work, approvals waiting, overdue, SLA met, the SLA, breaches this month, last action), **Their objects**, **Open work by object**, what waits for their approval, **Open their work**, **Reassign…** and **Activity**: their own audit trail. Your own activity is always visible to you; another person's needs view audit or manage users.

Balance work

"An assignment plan, applied with undo." Choose how to balance:

- **Even counts**: "Every person ends within one record of the others."
- **Whole objects**: "A person keeps all of one object's records together."
- **By the assignment template**: "Each record goes to the steward named for its object and plant, company code or organisation."

Propose a plan shows the moves before and after. Applying moves the records; the notice that follows has **Undo**. Balancing needs the assign permission.

Data contracts

"What each consumer is promised." Each contract names a producer object, a consumer and a promise, such as "Score at 85 or more", with its compliance (**Breached** or met), when it was last checked, its **Versions** and **Breaches**. Owners and rule maintainers create drafts with **New contract**; a data owner approves a contract, never their own; **Revise as a new version** starts the next version as a draft.

Out of office

"Your approvals, SLA breaches, mentions and assigned work are copied to them between these days. Four eyes still holds: nobody approves their own batch." Choose who covers, the first and last day and an optional reason, then **Set out of office**. **You cover for** lists whom you deputise for, with **End it**.

Triage and SLA

SLA days to resolve by severity (by default critical 1 day, high 3, medium 7, low 14), when to escalate, and the **Escalation chain**: "After a breach the assignee and the team lead hear first. Each level below notifies once more when an item stays past due that long." Changing SLA days and the chain needs the manage triage permission; anyone can choose their own digest here.

The numbers and where they come from

- **Open work** counts the records assigned to the person in the latest state of the work list.
- **Overdue** counts their open records past the due date their SLA gives them.
- **SLA met** is the share of their open work still within its SLA days; a dash means they hold no open work.
- **Breaches this month** counts the records that passed their SLA this calendar month.
- **Approvals waiting** counts the batches waiting for that person to approve.

The inbox

The inbox panel says how many items are unread and how many mentions, groups items under **Today** and **Earlier**, and gives each one its kind, time, sentence and one action: **Approval** ("Batch B-241 waits for your approval. 328 records of Material master, ZAR 412,600 at risk."), **SLA at risk**, **SLA breached**, **Notice** (a finished run), **Data quality** (a rule that rose), **Exception**, **Digest**, **Report**, **Monitor**, **Cleaning**. **Mark read** and **Mark all read** clear them. **Notification settings** opens your preferences.

Comments with an @ mention of you count as unread mentions until you open them.

What you can do, by role

- **Everyone**: read the inbox, set out of office, choose a digest, read People and contracts, see their own activity.
- **Programme lead and steward with assign**: balance work, reassign a person's work.
- **Data owner**: approve contracts, receive escalations, set a deputy for approvals.

- **Administrator and triage managers**: SLA days and the escalation chain.
- **Auditor**: a person's activity trail.

Keyboard

- `g` then `p` opens People.
- In the inbox and the People table, `j` and `k` move and `Enter` opens.
- `Esc` closes the panel.

When something is wrong

- **SLA met shows a dash**: the person has no open work to measure.
- **"Reassigning needs the assign permission, which your role does not hold."** Ask your lead.
- **A deputy cannot approve**: the deputy needs the approve permission themselves, and cannot approve a batch they built.
- **You hear about too much or too little**: Preferences, What you hear about, and per object or rule the **Watch** button.

Runs

A run is one measurement of your data at a point in time: an extraction from a system scored against the rules as they stood. Every score, failing record and report belongs to a run. **Runs** lists every run of every system, so you can see what ran, what failed and what changed. Open it from the rail, with `g` then `r`, or at `/runs`.

Run numbers restart for each system, so outside a system's own page a run is always named with its system: "ECC production run 12". What the product calls the current state is each object at its own newest finished run, not one latest run for everything.

What is on it

- **The sentence**: "Every run of every system, newest first within each system. Open a run to see what it scored, or compare it with the one before to see what changed."
- **The banner**: the small line says "Latest run" and the headline "Finished on 16 Oct 2026" ("No run has finished yet" before the first). The line under it says which way the data moved: "The average score rose by 2.1 points since the run before.", "did not change", or "A second run is needed to say whether the data got better or worse." when there is only one. Three figures: **Runs shown** (with the baseline date or "no baseline set"), **Average score** ("out of 100") and **Took** for the latest run. The one button is **Start a run**, for people who may start an analysis, with the reason beside it.
- **The charts**: **Score over time**, the average object score over the last runs as a line ("Two runs are needed to draw a trend." until there are two), and **How long each run took**, a bar per run on one scale from zero. A long bar is a slow run.
- **Runs and Jobs**: two views.
- **The runs ledger**: one group per system, active systems first, each headed by its name and its count, for example "9 runs, 1 failed". A system that is no longer active keeps its runs, marked Inactive, and they still open. Each group has a date axis with one square per run, then a table: run, status, records checked, objects scored, score range across objects (for example "70.0 to 93.4"), how long it took, and **Compare with previous** on every run but the first of its system. Filter by system and status; **Show archived runs** brings back archived ones.
- **Jobs** table: every background job (extraction, analysis, configuration load, configuration sync, upload, simulation, process discovery) with its object, system, state, progress, start and duration.
- **Trend per object**: small charts, one per object, of the score across runs.

The job line and the job inspector

A running job shows as a thin progress line under the top bar, never a tray. Open it, or any job row, for the job inspector: its stages and tables with row counts, the percent, the error in plain words for a failed job (for example "OData refused credentials (HTTP 401). Check the client id and secret in the system settings.") and, where the job recorded how it started, **Retry**.

Every start form (Administration, Start and schedule; the system inspector's **Extract** and **Run analysis**; **Start a run** here) shows the prediction before you confirm: the duration band and how sure it is, the records it will check, the cost from the cost model's run cost (a dash when none is set) and

where it would queue, for example "Position 2 in the queue; 1 running, 1 queued". A queued job says its place on the job line and in its inspector.

The numbers and where they come from

- **Records** and **Rules** are the records checked and the rules measured in that run, across every object.
- **Score** is the lowest and highest shown score of the run's objects.
- **Took** is from the job registry: extraction and analysis timings are recorded by the workers.
- **The prediction** reads the tenant's last eight finished runs of each object: the median records, the median seconds per record, and a band from the 25th to the 75th percentile; with fewer than four runs the band is the lowest to the highest. Nothing is estimated by a language model.
- The baseline is the run you pinned; every comparison and "since the baseline" figure measures against it.

What you can do, by role

- **SAP Basis and analyst** (trigger sync, analyse): start an extraction or an analysis, retry a failed job.
- **Anyone with analyse**: start a run ("Starting a run needs the analyse permission"), pin a baseline, archive older runs.
- **Anyone with export**: download a run's report.
- **Everyone**: read runs and jobs, open a run, compare two runs.

Starting runs, schedules and retries for one system at a time also live under Administration, **Start and schedule**: "Start an extraction, an analysis, a configuration load or an upload, and schedule them. The runs themselves are on Runs."

Keyboard

- `g` then `r` opens Runs.
- `j` and `k` move through runs or jobs, `Enter` opens one.
- `Esc` closes the job inspector.

When something is wrong

- **A run reads Failed with no records**: open its job. Connection errors (HTTP 401, RFC logon failures) belong to SAP Basis; the job says which credential or table failed.
- **A run is complete but scores are missing for an object**: the extraction read no records for it, or the analysis skipped its rules. Open the object; its receipt says how many rules produced a result.
- **Retry is missing**: the job did not record how it started; start it again from the system or from Administration, Start and schedule.
- **The job line shows nothing but a job is running**: the job registry updates every few seconds; reload the page if it does not appear within a minute, then ask your administrator to check Platform and updates.

Run and Compare

The **run canvas** shows one run: what it scored, object by object, and the jobs that made it. **Compare** puts two runs side by side and explains why each score moved. Open a run from the Runs ledger or an object's Trend tab (`/run/{id}`); open a comparison with **Compare with previous** in the ledger, **See every change** on a run, or **Compare with** on an object's Trend tab (`/run/{earlier}/vs/{later}`).

One run

- **The title**: the system and the run number, for example "ECC production run 12", with its run chip: date, status and whether it is the baseline.
- **Since the run before**: "12 objects scored higher and 3 lower. 14 rules started failing and 37 stopped.", a bar per object, and **See every change**. The first run of a system says there is nothing earlier to compare with.
- **The banner**: the small line says "This run", the headline is the run's worst verdict and the line under it sums up the objects: "15 objects scored: 2 go, 2 at risk, 11 no-go." Three figures: **Ready** ("of 15 objects"), **At risk** ("below target") and **Blocking** ("block go-live"). The one button is **Show scores from this run**.
- **The facts line**: the system, records checked, rules measured, scoring time and trigger, for example "ECC production, 400,662 records checked, 794 rules measured, scored in 11 m 42 s, trigger scheduled".
- **The sentence**: "15 objects scored: 2 go, 2 at risk, 11 no-go."
- **Actions**: **Show scores from this run** (sets the run in the top bar so every canvas reads it), **Pin as baseline**, **Archive the 8 older runs**, **Download report**, **Compare with**.
- **By object**: "Measured score, then the score shown after the cap". Each object with its measured score, shown score (with **capped**), verdict, records checked, rules and open work.
- **Jobs**: "The extraction and analysis that made this run", with state, start and duration.

Compare

Compare works on two runs of one system. The title names both, for example "ECC production run 12 vs ECC production run 9", and a run chip for each sits under it with two lists, **Earlier run** and **Later run**, that offer only runs of that system. If a link pairs runs of different systems the page says so and shows no figures: "Compare runs of the same system. ECC production run 3 and SuccessFactors run 2 belong to different systems, and run numbers restart per system."

The one button is **Open what got worse**, with its reason beside it, for example "Material master lost the most, -3.2 points. This opens its open records as they stood at ECC production run 12." It opens Work as of the later run for that system and object. Work cannot yet filter to what is new in that run, so the list is everything open at the time. When nothing got worse the button is **Open run details**.

The banner says "Compared with" the earlier run and gives the answer as its headline: **Data quality improved**, **Data quality got worse** or **Data quality is about the same**. Three figures follow: **Rules newly failing** (with "37 stopped failing" under it), **Objects higher** and **Objects lower**.

Two tabs sit under the banner. **Scores and rules** holds everything below. **Records** lists the records behind each changed rule.

Change in score by object is a diverging bar chart: one row per object, worst first, with bars to the left for objects that got worse and to the right for objects that got better, the change at the end ("-3.2 points"). The legend reads **Worse** on the left and **Better** on the right. A long bar to the left is where to look first.

Aletheia is docked on the right of both screens; ask why a score moved and the answer cites the facts and rules it used.

The sentence sums it up: "12 of 15 objects scored higher and 3 lower. 14 rules started failing and 37 stopped. 12 objects are still capped by a critical rule, so the shown score hides part of it."

The method is stated under it, exactly as the product states it: "Rules are compared on the common set: the 5167 rules that ran cleanly in both runs. 11 rules added since the baseline of 2 Oct 2026 are excluded. A rule counts as changed when it flips between failing no records and some, or fails at least 10% and 10 records more or fewer. Coverage is the share of records a rule could evaluate."

- **By object**: **Before** and **After** (measured scores out of 100), the change in points, shown before and after, records checked, coverage, and the counts of rules **New**, **Closed**, **Worse** and **Better**. The total row adds them. Scores here are measured, because a verified figure exists per run and not per object.
- **What changed**: "Rule counts match the table above." Four lists: **Newly failing rules**, **Fixed rules**, **Worse** and **Better**, each rule with its object and its failing records before and after, for example "MM201 Base unit of measure not in the check table, Material master, 0 to 316 records failing".

Above the object table, the **Why the score moved** panel explains one object at a time; pick it in the **Object** list, which starts on the object that moved most. A bar chart shows **Points by rule**: each rule with its code and description, and how many points it added or cost, whether its failing records changed, the records checked changed, the rule is new, or the rule no longer ran. Separate bars show the dimension weighting and the critical cap when they moved the score. The contributions are computed from the same scoring code as the run, and the line above the chart states how many of the changed points they add up to; they agree within 0.05 points. The sentences below the chart are written by Aletheia from the largest contributions only, and each sentence is dropped if a number in it is not in the figures. If no model is available, the sentences come from a fixed template and the panel says so.

On an object's Trend tab the same comparison appears as **Why the score moved**, and ends with what the cap did, for example "The shown score moved only 69.8 to 70.0, because the cap holds it at 70.0."

The numbers and where they come from

- Measured and shown scores are as stored with each run; nothing is recomputed when rules change later.
- Only the common set of rules is compared, so a rule added later never shows as "newly failing".
- The change threshold (10% and 10 records) stops noise from counting as a change.

What you can do, by role

- **Everyone**: open runs, compare any two, export the tables.
- **Anyone with analyse** (programme lead, analyst, SAP Basis, admin): pin the baseline, archive older runs.
- **Anyone with export**: download the run's report.

Keyboard

- `j` and `k` move through the tables, `Enter` opens an object or rule.
- `g` then `r` returns to Runs.

When something is wrong

- **The shown score did not move although many records were fixed**: the cap holds it. Read the measured score, and fix the critical rules over tolerance to release it.
- **"Compare runs of the same system."** The two runs belong to different systems, and their numbers mean nothing side by side. Open one of them and use **Compare with previous**.
- **A comparison says few rules were compared**: one of the runs ran fewer rules cleanly (a failed extraction or a narrower scope). Compare runs of the same scope.
- **Pin as baseline or Archive is missing**: changing a run needs the analyse permission; archived runs stay readable from **Show archived runs** on Runs.
- **"This run could not be read. It may belong to another tenant or have been removed."** Check the link, or open the run from the Runs table.

Systems

A system is one SAP or cloud source Meridian reads from: ECC, S/4HANA on premise or Cloud, Extended Warehouse Management, SuccessFactors, Concur, Ariba or BTP. Meridian reads from SAP with a read-only role; live write-back is off unless an administrator turns it on (see the Approvals chapter). The system canvas is where SAP Basis connects a system, tests it, extracts with a scope, schedules extractions, loads the configuration and reviews custom objects. Open a system from its heading on the landscape or the Objects menu (`/s/{id}`); connect a new one from the Objects menu, the command line ("Connect a system") or at `/s/new`. Every system is listed at `/s`.

The systems list

`/s` opens **Systems**, with the sentence "2 systems connected." and either "Every connection answers." or how many "cannot be reached or signed in to". **Connect a system** opens the setup. A drawing shows which systems feed Meridian (on the left) and which receive from it (on the right), each name coloured by its health. Under it, **Connection health** lists every system: name, release or type and environment, the health (**Connected**, **Degraded**, **Unreachable**, **Sign-in failed** or **Not tested**), "Last sync" with its time or "Never synced", and the reason when the connection or the last sync is not healthy. Select a row to open the system in the inspector. With nothing connected the page says "No system is connected yet." with a link to connect one.

Connect a system

`/s/new` opens **Set up Meridian**: "Setup only reads from SAP. Nothing is written to SAP until an approved batch is sent." Six steps run down the left, each ticked when done; a later step opens once the one before it is done. On the right, **What the first run will do** counts the systems, objects, tables and policy packs applied so far.

1. **Connect SAP**: the connection form below. Saving it moves to the next step.
2. **Connect SuccessFactors**: optional. Choose SuccessFactors as the system type, or **Skip for now**.
3. **Choose objects**: the objects the SAP system offers, each with its number of tables. Tick at least one, then **Continue**.
4. **Choose policy packs**: GxP, SOX and GDPR with their authority. Ticking a pack applies it at once; untick to remove it.
5. **Schedule runs**: the rhythm and mode, as on the Schedules view (the default is every night at 02:00, changes only). **Save schedule** creates one schedule for each chosen object.
6. **Invite people**: an email and a role (data steward, data owner, analyst or viewer) per line, **Add another**, then **Finish setup**. Empty lines are skipped.

"Setup complete" follows: "The first run starts on the schedule you set. Open the system to start it now." **Open the system** opens its canvas.

The connection form in step 1:

1. Name the system, choose the **System type** and the **Environment** (Production, Quality, Development).

2. Fill in the connection. For ECC and S/4HANA on premise: **Host**, **Client** (3 digits), **Instance number** (2 digits) and an optional **User** ("Leave empty to use the deployment's default RFC user"), then the password under **Sign-in secrets**. Cloud systems ask for their base URL and OAuth details instead.
3. For ECC, S/4HANA on premise and Extended Warehouse Management choose the **Extraction path** and the **Table reader** (see "Extraction path" below). The defaults, **Table reader** and **Deployment default**, read as before.
4. **Test connection**. A passing test reads like "RFC connection succeeded (client 100, 41 ms)".
5. **Connect system**. "You can connect now and test later. Nothing is read from SAP until you extract."

Extraction path

How an ABAP system's tables are read, chosen in the connection settings:

- **Table reader**: "Reads each table with the table reader below. Needs read access to the tables."
- **ODP, table reader for the rest**: "Reads mapped tables through ODP providers (RODPS_REPL). Fields a provider does not carry are reported per table; other tables use the table reader." The providers are classic extractors such as `0MATERIAL_ATTR` (MARA), `0MAT_PLANT_ATTR` (MARC), `0CUSTOMER_ATTR` (KNA1), `0VENDOR_ATTR` (LFA1) and `0GL_ACCOUNT_ATTR` (SKA1), and on S/4HANA released CDS views such as `I_Product` and `I_BusinessPartner`.
- **Automatic: table reader**: "Reads with the table reader. ODP is used only when you choose it, after confirming with SAP that your licence allows it (SAP Note 3255746)."
- **OData through the SAP Gateway** (S/4HANA on premise only): "Reads through the released APIs your Gateway exposes, such as API_PRODUCT_SRV and API_BUSINESS_PARTNER, with this system's user." Enter the **Gateway base URL**; the system's client is sent as `sap-client`.

A table read with a date window, a scope or by parent keys always uses the table reader, because ODP reads a provider in full. The **Table reader** setting chooses `RFC_READ_TABLE`, `/BODS/RFC_READ_TABLE2` or `BBP_RFC_READ_TABLE`, or **Find the first available**.

Which path read each table is on the extraction job's **Tables** list (**Read through**, for example "ODP, 0MATERIAL_ATTR" or "Table reader") and on the Basis home: **Connections** shows each system's extraction path and **Last extraction** how its tables were read, such as "ODP: MARA; 4 with the table reader".

The system canvas

The header gives the health ("Connected"), the type and protocol ("ECC, RFC"), the last extraction and its records, and the last connection test. **Test, load or extract** opens the system inspector; **Connection settings** opens the connection editor.

The banner under it says "State of the system" and puts the one thing to know in its headline: "An extraction from ECC production is running, 40% done", "The latest extraction from ECC production failed", "ECC production is connected", "The connection to ECC production has not been tested yet", or "The connection to ECC production needs attention". The line under it says when it was last extracted ("The last extraction was on 16 Oct 2026, 05:09." or "It has never been extracted."). Three figures: **Objects extracted**, **Extractions kept** ("to compare against") and **Objects scored** ("by the latest run"). Aletheia is docked on the right, as on every screen.

Five views:

- **Overview**: **Records per object**, a horizontal bar per object: "What the latest extraction, on 16 Oct 2026, 05:09, read from each object." Then "What the latest run scored", object by object, and **Recent jobs** from the job registry.
- **Extractions**: the stored extractions of this system and whether each was analysed.
- **Configuration**: the configuration Meridian read from the system's customising tables, by area (enterprise structure, financial accounting, controlling, materials management and so on), each with "Loaded, 3 of 3 objects" and what it found, such as "4 company codes, 11 plants". **Load again**, **Discover now**, the **Drift log** and **Compare with target** (the target system chosen in the connection settings, or the SAP standard baseline). A table that could not be read says why, for example "T500P: No authorisation to read this table. The user needs read access to it."
- **Schedules**: "Each schedule extracts the object again and re-runs its checks, so trends and alerts see every run. Times are UTC." Add a schedule with an object, a rhythm (daily, weekly, monthly, hourly or a custom cron) and a mode: **Full extraction** or **Changes only**. "Changes only reads the records SAP changed since the last download, with a full download at least once a week." Pause or delete a schedule from its row.
- **Custom objects**: "Custom (Z) objects found across every connected system in this workspace, from the latest analysis." Facts (custom objects, active anomalies, ghost and dormant, mapped to standard), **Detect now**, and the registry, anomalies, ghost, dormant, drift, rules and mapping views. Open a Z object to map it to its standard equivalent or add a Z rule.
- **Source design**: the data dictionary as discovery read it from this system, against the SAP standard. Facts (SAP product and release, last discovery, tables, customer tables, customer fields, configuration read) and **Discover again** (needs the trigger sync permission). **Tables and fields** lists each table with its description, kind and customer fields; open one to see every field here beside its SAP standard definition and check table, with the difference in words, and the standard fields this system does not have. **Configuration against standard** shows each check table's live entries against the standard list the rules fall back to, naming customer entries and missing standard entries. **Drift between discoveries** compares two finished discoveries and lists tables and fields added, removed or changed, with the old and new definition.

Extract with a scope

In the system inspector, **Extract** asks for the objects and the scope: "Choose objects to set the scope: company codes, plants, organisations and a date window." With no filter the scope reads "Scope: everything, no organisational filter." The same extraction form under Administration, Start and schedule adds company codes, plants, a document date window and **Analyse straight after extracting**, which scores the new extraction when it lands.

Dry run shows "What an extraction of 3 objects would read, without reading it": the rows per table last time and an estimated duration, in the form "rows in tables, about a duration (this system's recent speed, rows a second)". Where the system has no recent extraction the speed reads "an assumed speed". From there **Extract changes only** starts a delta extraction that reads the change documents since the last download, where the system supports it; otherwise the reason is shown.

Run analysis on ECC production scores a stored extraction against the rules as they stand now. "Earlier scores are kept as they were."

The numbers and where they come from

- **Systems connected** on `/s` counts every system; the systems that cannot be reached or signed in to are those whose health is Unreachable or Sign-in failed.
- **What the first run will do** counts what the setup has saved so far: tables are the tables of the chosen objects.
- **Records** is what the last extraction stored, across every object.
- **Configuration** counts areas loaded out of those the connector knows, for example "`11` of `12` areas".
- The dry run counts rows from the last extraction of each table; a table never extracted is not counted and says so.

What you can do, by role

- **SAP Basis** (manage systems, trigger sync): run the setup, connect, test, rotate credentials, extract, dry run, delta, schedule, load configuration, remove a system.
- **Analyst and programme lead** (analyse, trigger sync where granted): run analysis, extract.
- **Rule maintainers** (manage rules): edit a custom object's registry entry, add a Z rule.
- **Policy packs and invitations** in the setup need manage rules and manage users; without them the step shows the refusal under it.
- **Everyone else**: read the systems list and each system.

Keyboard

- `⌘K`, then "connect", opens Connect a system.
- In every table, `j` and `k` move and `Enter` opens.
- `Esc` closes the inspector.

When something is wrong

- **"Sign-in refused"** or **HTTP 401**: the user or secret is wrong or expired. Open **Connection settings**, type the new values ("Stored secrets are never shown. Type new values to replace them; an empty field keeps the stored one.") and test again.
- **"No authorisation to read this table. The user needs read access to it."** Grant the RFC user read access to that table in SAP (S_TABU_NAM, or S_TABU_DIS), then **Load again**. The read-only role template is `docs/sap/authorisation-role-template.md`.
- **The Basis team refuses `RFC_READ_TABLE`**: choose another **Table reader** (`/BODS/RFC_READ_TABLE2` or `BBP_RFC_READ_TABLE`), or an **Extraction path** that does not read tables directly: **ODP** or, on S/4HANA on premise, **OData through the SAP Gateway**. The data read is the same; fields ODP cannot deliver are listed per table.
- **"RODPS_REPL_ODP_OPEN is not available on this system"**: the ODP functions are missing or not authorised. Grant `S_RFC` for function group `RODPS_REPL`, or choose **Automatic**, which reads with the table reader.

- **An OData table failed with 404**: the API service is not activated on the Gateway (`/IWFND/MAINT_SERVICE`).
- **"The table does not exist in this release. Nothing to fix."** Expected on some releases.
- **"This system is extracted but no rule has run on it. Run an analysis."** Use Run analysis.
- **"Testing, loading and extracting need the trigger sync permission on this system."** Ask your administrator for a grant on that system.
- **A schedule did not run**: due runs start within 5 minutes of their time; check Runs, Jobs for a failed extraction.

Golden records

A golden record is the one trusted version of a thing your systems hold more than once, such as a material kept twice or a supplier in two systems. The **Golden records** tab of an object lists the match candidates waiting for review, the golden records already built and the survivorship rules that decide, field by field, which value survives. Open it from the object's tabs, for example `/o/ecc-production/material-master/golden`.

What is on it

- **Match candidates**: "148 pairs to review", best score first, each pair with its match score and both keys, for example "0.94, MATNR=5000 and MATNR=6000", and **Review**. "Showing the first 5 of 148, best score first." **Show 20 more** loads more. With the approve permission each pair has a box: tick the pairs a steward should decide, pick a priority (Urgent, Normal or Low) and, if you like, a due date, then **Send N to a steward**. The pairs go on the stewardship queue as merge decisions, due in 48 hours unless you set a date.
- **Golden records**: "1,284 records", with golden key, name, sources, survivorship, confidence, last merged and open conflicts. The line under the table explains them: "Confidence is the match engine's score for the surviving record. Last merged is when the record last changed. Open conflicts are fields where the sources disagree and no rule has settled it."
- **Survivorship rules**: the rule per field, for example "MARA-MEINS Trusted source, ECC production" or "MARA-MATKL Most recent", with **Edit the rules**.

The golden record inspector

Open a golden record for **Golden fields and how each value was chosen** (the value chosen, the rule that chose it, or "No rule decided this field"), its **Source records**, its open conflicts and its history ("No merge has changed this record yet."). **Choose records to split out** separates a record merged in by mistake.

Review a pair

Review opens the pair: Record A and Record B **Side by side**, "The two records side by side, pick a value per field", with the **Fields that differ** and **Keep the key of** to choose the surviving key. "Each value you pick on a field that differs is kept in the merged record; the others follow the survivor." Merge the pair, or mark it **Not a duplicate** with the reason ("Why these are different materials"): "This pair will not be offered again." Every decision is in the audit log.

Survivorship rules

"A rule decides which record's value survives in the golden record, field by field. A field with no rule keeps the survivor's value. Rules are listed by field." The rule types and what each says:

- **Most recent**: "The value from the record changed last survives."
- **Most complete**: "The longest non-empty value survives."
- **Most frequent**: the value most sources hold survives.

- **Trusted source**: "The value from the first trusted system that has one survives." List the trusted systems in priority order.
- **Chain of systems** and **Manual override**.

Each rule has a weight from 0 to 100 and can be switched off. A rule the miner suggested is marked "suggested by the miner" and takes effect only when a person saves it. **Save rules** stores them; they apply at the next match run.

The numbers and where they come from

- **Match score** and **Confidence** come from the deterministic match engine, from 0 to 1 and as a percentage.
- **Sources** counts the records merged into the golden record.
- **Open conflicts** counts fields where the sources disagree and no rule settled it.

What you can do, by role

- **Data owner** (approve): merge a pair or mark it not a duplicate. "Merging, or marking a pair as not a duplicate, needs the approve permission. A data owner can decide this pair."
- **Data steward** (mdm create and edit): read pairs, edit survivorship rules.
- **Everyone else**: read, export.

Keyboard

- `j` and `k` move through candidates and golden records, `Enter` opens.
- `Esc` closes the inspector.

When something is wrong

- **No match candidates**: matching has not run for this object, or found nothing above its threshold.
- **Many open conflicts**: a field without a survivorship rule. Add one for the fields downstream processes depend on.
- **A golden value differs from SAP**: golden records live in Meridian. SAP changes only through an approved batch or an approved write-back.

Glossary

The glossary is the shared language of an object: one business term per SAP field, with its definition, its owner, its status and the rules that cite it. The **Glossary** tab of an object lists its terms and the fields that still have none. Open it from the object's tabs, for example `/o/ecc-production/material-master/glossary`.

What is on it

- **Search terms** and **Add terms from the rule library**.
- **Business terms**: "6 terms", each with term, definition, SAP field (`MARA-MEINS`), owner, status (**Active** or **Under review**) and rules. For example: "Base unit of measure: The unit in which stock of the material is managed. Every other unit is converted to it."
- The line under the table: "Owner is the owner of the material master object. Rules counts the rules that cite the term."
- **Fields without a term**: the fields read by failing rules that have no term yet, each with the rule that reads it, for example `MBEW-BKLAS` and `MM035`. "The fields are those read by the failing rules of this object in the latest run. Terms are drafted from the rule library with Add terms from the rule library, and a data owner approves each."

The term inspector

Open a term for its definition, **Fields this term governs**, **Rules that cite this term**, its review cycle and when it was last reviewed, and its history ("The change is in the term's history."). **Mark as reviewed** keeps an active term current. "Editing and approving a term needs the approve permission. A data owner can do both." Approving moves a term from **Under review** to **Active**.

The numbers and where they come from

- **Rules** counts the rules that cite the term.
- **Fields without a term** comes from the failing rules of the latest run, so it changes as rules start or stop failing.

What you can do, by role

- **Data owner** (approve): edit, approve and mark terms as reviewed.
- **Data steward**: draft terms with Add terms from the rule library, comment on terms.
- **Everyone**: read, search, comment, export.

Keyboard

- `j` and `k` move, `Enter` opens a term.
- `Esc` closes the inspector.

When something is wrong

- **A term stays Under review**: it waits for the object's owner. People shows who owns the object.
- **No owner shown**: the object has no owner yet. Administration, Stewardship lists the gaps.
- **A field is missing from Fields without a term**: only fields read by failing rules are listed.

Lineage

Lineage draws why a failing value matters: from the source table, to the field, to the rule that fails on it, to the S/4 feature it blocks or degrades. The **Lineage** tab of an object draws it for the failing rules with the highest value at risk. Open it from the object's tabs, for example `/o/ecc-production/material-master/lineage`.

What is on it

- **Four columns**: **Source tables** (`MARA`, `MARC`, `MBEW`), **Fields** (`MARA-MEINS`), **Rules** (`MM142`) and **S/4 features** (Migration cockpit: material load, Material ledger actual costing, Embedded EWM product master and so on).
- **Node facts**: "Hover or focus a node to see its facts here. Click a rule or a feature to open it." Click a node to pin its facts; under them **Upstream** lists the nodes that feed it and **Downstream** the nodes it feeds, each with a count, or "None". Select one to move the facts to it.
- **Export as SVG**: the drawing as a file for a slide or a document.
- The line under it: "Drawn from the 8 failing rules of Material master with the highest value at risk. Features come from the impact model of the current run."

Related objects

The object header lists the related objects with their relation, **upstream** or **downstream**. An upstream object's failures pass into this one, so the Work ledger and the record inspector say which to fix first, for example "Fix Material master first: 3,640 upstream failures. Records there fail before this object can pass."

The numbers and where they come from

- The rules drawn are the failing rules of this object in the current run, highest value at risk first.
- The features and whether each is blocked or degraded come from the feature impact rules of the run: the same mapping the Readiness tab shows.

What you can do, by role

- **Everyone**: read, open a rule or a feature, export as SVG.

Keyboard

- `Tab` moves through the nodes; the node facts follow the focus. `Enter` opens a rule or a feature.

When something is wrong

- **The drawing is empty**: no rule of this object fails in the current run, or the object has no run.
- **A feature you expect is missing**: no impact rule maps the failing rules to it. Impact rules are shipped with Meridian; ask your administrator.

Data catalogue

The data catalogue is one place to search and browse every SAP table and field Meridian knows, across every connected system. Open it from the rail, Master data, **Data catalogue**, or type "catalogue" in the command line. The address is `/catalogue`.

What is on it

- **Search**: type a table (`MARA`), a field (`MEINS`), a table and field together (`MARA-MEINS` or `MARA.MEINS`), words of a description ("gross weight") or the name of a business term. **Look in** chooses **Tables** or **Fields**; **System** narrows the list to one connected system.
- **Tables**: every table with its description, its number of fields, the rules that read it, the objects whose rules read it and the systems that use it. Tables your systems use come first. Choose **Columns** to show **Defined by**: SAP standard, SAP standard with the S/4HANA changes, the system's own dictionary (a customer table) or a cloud system's view such as SuccessFactors.
- **One table**: open a row to see its key, where it is defined, its systems, its objects and the owners of those objects, then the grid of its fields: type, key, the number of rules, the business term and the share of blank values in the newest profile. **All tables** goes back.
- **Fields**: with **Look in** set to Fields, the fields that match the search, each with its table, its type, its rule count and its business term.
- **A field in the inspector**: open any field (`?open=field:MARA.MEINS`). The inspector shows its facts (type, key, mandatory, data element, domain, check table), the SAP help text, the allowed values of its domain, its latest profile with a link to the profile of that run, the rules that check it (with the severity, the object, whether the rule is off for your tenant and its latest pass rate), its business terms with their steward, the objects and systems that use it with their owners, and **Where its values go**: the process steps, processes, S/4 features and KPIs the lineage model draws below it.

The numbers and where they come from

- Descriptions, types, lengths, keys, check tables and allowed values come from the SAP dictionary Meridian ships (SAP ECC 6.0 standard with the documented S/4HANA changes, and its views of SuccessFactors, Concur and Ariba). When you choose a system that discovery has read, its own dictionary wins, so customer fields and customer tables appear with their own definitions.
- **Rules** are the shipped and tenant rules that read the field, the same columns the check engine reads. A rule your tenant has switched off says so.
- **Systems** are the connected systems where an object that reads the table is enabled, where an analysed run profiled the table, or whose dictionary holds it.
- **Blank** and the profile come from the newest analysed run that profiled the field; with a system chosen, from that system's newest run.
- Business terms and their stewards come from the glossary; owners from the owners set for objects and systems. Nothing on this screen is estimated by a language model.

What you can do, by role

- **Everyone who can view**: search, browse, open a field, follow a rule, a business term, a profile or a lineage link.
- The catalogue only reads. Rules are changed in the rules library, terms in an object's Glossary tab, owners in Administration.

Keyboard

- `⌘K` or `/` opens the command line; type "catalogue" to come here.
- In a grid, `j` and `k` or the arrow keys move between rows, `Enter` opens the table or the field, `Esc` closes the inspector.

When something is wrong

- **A table you expect is missing**: it is neither in the shipped dictionary nor in a discovered dictionary of your systems. Open the system's **Source design** and press **Discover again** so its own tables are read.
- **A field has no profile**: no analysed run has profiled it yet. Profiles are written for the tables an object's rules read when a run is analysed.
- **Systems is empty for a table**: no object that reads it is enabled on a connected system, no run profiled it and no discovered dictionary holds it.
- **"This system is not connected to this tenant"**: the address names a system that is not yours; choose another in **System**.

Processes and the designer

The **Processes** tab of an object shows which business process steps its data blocks: Procure to Pay, Plan to Produce, Order to Cash and the others, down to the step, the transaction and the field. The **process designer** is where the process model is edited. Open the tab from the object's tabs (`/o/{system}/{object}/processes`) and the designer with **Open the designer** (`.../processes/design`).

What is on the Processes tab

- **The sentence**: "2 of 4 process steps are blocked by Material master: PTP step 3, PP step 1."
- **Each process**: its name and code ("Procure to Pay PTP"), its state ("Blocked, 1 of 3 steps blocked"), its description and its steps. Each step names the transaction (`ME51N`, `ME21N`, `MIGO`), its state (**Ready** or **Blocked**) and, when blocked, the activities and fields behind it with status, pass rate and failing records, for example "Enter Quantity and Unit: `MARA-MEINS` Base unit of measure, Blocked, 94.1%, 3,640".
- The rule under it, as the product states it: "A step is blocked when a field it writes fails a critical or high rule, or passes under 70%. Failing counts are records, and a record can fail several rules."
- **What if these rules were fixed** (simulation): "Tick the rules to fix. The steps of Material master are read again without their failures; nothing is changed and no data is read." Each rule shows "3,640 records at least, in 1 step". The result shows which steps would be ready.
- **Process designer** with **Open the designer**.

The step inspector

Open a step for its transaction, its activities and fields, the rules failing on them and **Hand off**, which hands a blocking rule to a steward with a note.

The process designer

"Process designer, Material master", with the model's state, for example "Draft, version 1, updated 12 Oct 2026".

- **Model**: choose a model, **Create a model** (start from the reference model or the group model), **Publish model**, **Rename model**, **Delete model**.
- **The tree**: levels L1 to L4 (for example Procure to pay, Sourcing, Purchase order, Create purchase order with its activities), each with **Up**, **Down** and **Delete**, and **Add a process at the top level**.
- **The model panel**: "Choose an item in the tree to rename it, reorder it, add what sits under it, or map fields to its activities."
- **What changed**, **Save as a new version**, **Discard changes**, and **Versions**.
- The rule at the foot: "A step is blocked in process readiness when a field it reads fails a critical or high rule. Map the fields each activity reads here, and the Processes tab shows them on the next run."

Process discovery

Process discovery (`/processes`) mines the as-is flow of each L1 process from the event tables of one system. Pick the system at the top. The sentence reads "3 of 7 processes mined on ECC production", with the median fitness, the cases mined and the processes behind their schedule.

- **The catalogue** lists every L1 with its cases, variants, median throughput time, automation rate, fitness against the reference model and the data quality defects on the records it touched. **Mine** starts a run; **Mine again** refreshes one. Starting a run needs the run process discovery permission. A process the system cannot mine says why.
- **Open flow** shows the run. The steps are bars sized by case count; click a step to go down a level, from L2 to L5, and the breadcrumb takes you back up. The table under it lists the as-is hand-overs between steps with their counts and median wait.
- **Variants** lists the case paths, most frequent first. Click one to filter the cases to it.
- **Conformance** shows the fitness of each reference step, how often it ran out of order, and the activities that happened outside the reference.
- **Cases** lists the cases of the run; open one to read its events in order, with the time, the user, whether it was automated and the source table.
- **Health** shows how the fitness of each process moved since the run before. **Schedules** re-mines a process on a cron schedule; pause, resume or remove a schedule there.
- **Exports** writes the run as a Signavio file, and a selected step as BPMN. Both need the export process permission.

The numbers and where they come from

- Step states come from the fields mapped to each activity and the rules that failed on those fields in the run.
- The what-if simulation is computed in the browser over the process document and the receipt; it reads no data and changes nothing.

What you can do, by role

- **Process analyst**: read readiness, run what-if, hand off a blocking rule (handoff rule), export the as-is process (export process), edit and publish the model.
- **Data steward**: receive handed-off rules in Work.
- **Everyone**: read.

Keyboard

- `j` and `k` move through steps, `Enter` opens a step.
- In the designer tree, `Tab` moves between items.

When something is wrong

- **"Choose at least one rule."** The simulation needs a ticked rule.

- **Every step is Ready but the object is No-go**: the failing rules are on fields no activity reads. Map them in the designer.
- **Your edits do not show on the Processes tab**: they show on the next run, and only once the model is published.

Configuration

Configuration (`/config`, on the rail under Migration, Configuration, and in the command line) shows the SAP configuration Meridian reads back out of the transactions of the latest run. Nothing on it is typed in by hand: every setting, step and finding comes from the stored config run.

- **The headline**: "14 configuration findings, R 2.1m at stake", with the average health across the areas, the critical and high findings, and the share of settings that carry transactions.
- **Processes**: each end-to-end flow (Procure-to-pay `PTP`, Order-to-cash `OTC`, Record-to-report `RTR` and the others) as a chain of steps. Each step is sized by the documents that reached it and names its SAP table; the red share is its exceptions, the step with the most exceptions is marked, and the days are the wait before the next step. A step the system never ran is drawn hollow and reads "Not found".
- **Health by area**: 100 less 3 per critical, 2 per high, 1 per medium and half per low finding, per area such as Finance `FI` or Materials management `MM`.
- **Settings in use**: the share of each area's configured values that carry at least one transaction. Low shares are settings that were configured and never used.
- **Findings**: worst first, each with its severity, what was found, the fix, the area, the settings it touches and the amount at stake.
- **Changes since the last run**: settings added, removed or changed between the last two config runs.
- **Read the configuration again** (analysts and above) queues a new config run on the latest analysed version; the figures change when it ends.

When no config run exists yet, the page says "Run an analysis on an SAP system to read its configuration."

Reports and schedules

Reports are files built from a stored run: the board pack, the analyst workbook, the analysis report, the trend report and the rule catalogue. Schedules build a report from the latest run and send it on a rhythm. Open Reports from the command line (type "reports") or at `/reports`.

What is on it

- **The banner**: the small line says "Report delivery" and the headline says what needs you first: "2 schedules failed on the last run", "1 schedule is overdue" (more than an hour past its next run), "1 schedule reached only some recipients", "Every schedule delivered on its last run", or "No report goes out on a schedule yet". Three figures: **Reports made** (with how many of the latest came from a schedule), **Schedules on** ("of 3 in all") and **Need attention** ("failed, overdue or partly delivered"). The one button is **Create a report**, for people with the export permission.
- **Reports made by kind**: a horizontal bar per report kind (board pack, analyst workbook and so on), with "4 made, 2 by schedule" at the end of the row. "Counted over the latest 6 reports in the history."
- **Aletheia**, docked on the right, as on every screen.
- **The sentence**: "6 reports made, 2 of 3 schedules are on. Reports are built from stored runs, so a number in a report matches the run it names."
- **Create a report** and **Schedule a report**.
- **Reports**: report, run ("Run of 16 Oct 2026"), format (PDF or XLSX), created, by, size and **Download**. The rule catalogue is "Not tied to a run".
- **Schedules**: "Each builds a report from the latest run and sends it." Report and format, cadence ("Every Monday at 07:00 (Africa/Johannesburg)"), recipients, next run and last result (**Delivered**, **Partly delivered**, **Failed** or "Not run yet"), and **Paused** for a paused schedule.

Create a report

Choose the report, the run it is built from ("Built from the run you choose") and, optionally, one object. The report is built by a worker and appears in the list when it is ready.

The schedule inspector

Open a schedule for its facts (next run, last run and its result), then: name, report (Board pack (PDF), Analyst workbook (XLSX), Analysis report (PDF), Trend report (PDF), Rule catalogue (PDF)), **Which run** ("The latest finished run", "The latest file upload" or the latest run of one system), an optional object, the cadence (every day, week or month, the day and the time in the tenant's time zone), the recipients (people of the tenant and the alert channels that take reports) and **Schedule is on**. Actions: **Save the schedule**, **Pause**, **Run it now** ("It is built on the worker and delivered to the recipients; the next scheduled run is unchanged.") and **Delete** ("No more reports are built or sent from this schedule. Files it already made stay in the list.").

The numbers and where they come from

Every figure in a report is read from the run it names, so the board pack of 16 Oct says what the landscape said for that run. Money is in the tenant currency from the run's cost model. Alerts and channel messages carry counts, rule ids and links, never record values.

What you can do, by role

- **Programme lead, analyst, admin** (export): create, schedule, run now, pause, delete.
- **Everyone with view**: download reports. "Creating and scheduling reports needs the export permission."
- **Administrator**: add the channels that can receive reports under Administration, Alerts.

Keyboard

- `j` and `k` move through reports and schedules, `Enter` opens a schedule.
- `Esc` closes the inspector.

When something is wrong

- **"No report has been made yet. Create one from a run, or schedule one to arrive on its own."**
- **"No one can receive reports yet."** Invite people, or add an email or Teams channel under Administration, Alerts.
- **Partly delivered**: one recipient or channel refused it. Check the channel under Administration, Alerts; email needs an email backend, which Administration, Platform and updates reports under Health checks.
- **Failed**: open the schedule and choose **Run it now** after fixing the channel; the job appears under Runs, Jobs.

Migration

Migration shows whether your records would load into the new system, and what to fix first. Open it from **Migration** on the rail, or press g then i. Its address is `/migration`.

What is on it

- **The steps**: Start an analysis, Fix what would not load, Load into the target, Check the result. The current step is marked.
- **The banner**: the verdict of the latest analysis. Starting an analysis only checks your records. The page updates by itself while it runs.
- **Gaps by kind**: why records would not load, grouped by the kind of problem, such as a value that is not in the check table.
- **Quality in the source and in the target**: a fold-out that compares the quality of each object before and after.
- **All analyses**: a fold-out that lists the analyses that have run, newest first. Select a date to open that analysis.
- **For the open analysis**: Gaps by kind, then **Load files** (records without blocking gaps, shaped for the target system, as CSV files or Excel), **Post-load reconciliation** (the load files compared with the target system after the load), **Field-level gaps** and **Records**, with a switch between records that would not load and records ready to load.
- **Field map**: how each source field maps to a target table and field, with the value mappings beside it.

The workbook from the new system

The panel **Workbook from the new system** fills the new system's own workbook or load file from a finished run. It has three steps: choose the run and workbook, check what fits, export the data.

1. Choose a finished run, and the file the new system gave you. Excel (.xlsx, .xlsm, .xls, .xlsb), OpenDocument (.ods), Migration Cockpit XML (.xml), delimited text (.csv, .tsv, .txt, .dat) and ZIP files can be read.
2. Select **Check the workbook**. The check reads the workbook and exports nothing yet. The workbook itself is not kept. The one thing the check can change is the extraction, as step 5 explains.
3. Read the result: how many rows are ready and how many are held back, and how many sheets can be filled. A line under the result says which format the file was read as. A table lists each sheet with what it is filled from, the columns filled, the rows ready and the rows held back.
4. Where required columns have no data in the run, the panel names them: the new system will refuse those rows until the columns are filled. Fold-outs list the required columns with no data, the country-specific columns that were not checked, and a sample of the rows held back with the reason for each.

5. Where the workbook asks for columns the extraction does not fetch yet, Meridian adds them to the extraction of that source system, so the next run fetches them. This applies to SuccessFactors workbooks, and only to columns that name a plain property of the sheet's entity. The extraction checks each one against the source system's own metadata, so a column is learned, never invented. The panel says how many columns were added, and the fold-out **Columns added to the extraction** lists them.
6. To fetch the added columns now, select **Update the run**. It starts a new extraction and analysis of that system. When the new run finishes, choose it in step 1 and check the workbook again. Without the right to start extractions, the panel asks you to have someone who can start extractions update the run.
7. Export the data in the same format, as Excel or as CSV files. The export holds the ready rows in the workbook's own columns, and held rows on a Held back tab. In the same format, held rows travel in a separate file, and .xls, .xlsb, .xlsm and .ods files come back as .xlsx. Every export is recorded in the audit log. The workbook itself is not stored.

Without a finished run, the panel says so; a finished run is needed before the workbook can be filled.

What you can do, by role

The buttons depend on the role, and a button the role cannot use is not shown:

- **Start an analysis** and **Delete** on an analysis need the right to analyse. An analysis that is queued or running cannot be deleted.
- The panels **Workbook from the new system** and the load files need the right to export. They are not shown in a shared link or while viewing as another role.
- **Update the run** and **Extract the target system** need the right to start extractions. Without it, the workbook panel asks you to have someone who can start extractions update the run.
- **Run reconciliation** needs the right to analyse, and downloading its result needs the right to export.

A role that cannot read migration analyses at all sees "Your role cannot read the migration analyses."

The numbers and where they come from

Gaps by kind come from the latest analysis of your records against the target system. The workbook figures come from the finished run and the workbook you choose, and nothing is guessed.

When something is wrong

- **The workbook could not be checked**: the page shows the reason in plain words. Check that the file is the one from the new system and try again.
- **Required columns still have no data after Update the run**: the new run must finish first. Choose the new run in step 1 and check the workbook again.
- **Rows are held back**: read the table of required columns with no data, and the reasons per row, then fix the source and run again.
- **No run has finished yet**: finish a run first. The workbook cannot be filled without one.

Keyboard

- `g` then `i` opens Migration.
- `Esc` closes the topmost layer.

Waves and gates

A migration wave is a set of objects moved together in one cutover, with a target date, a stage and a gate: the conditions every object must meet before the wave may go. The landscape's **Waves** view shows each wave with its verdict counts and gate state; the **Migration waves** panel is where the programme lead plans waves and sets gates. Open the view with **Waves** on the landscape (`/?view=waves`) and the panel with **Edit** on a wave (`?open=waves:{id}`), which opens that wave first with its gate on top and the other waves collapsed.

The Waves view

Each wave heading reads, for example, "Wave 1: finance foundations, due 16 Nov 2026, 1 go, 0 at risk, 3 no-go, Gate not passed, 2 of 3 not met", with **Edit** and **Certificate**. At phone width the heading stays above each wave's objects, with the same gate state and both links. Under it, the wave's objects with their ruler, score, verdict, records and rules. The programme lead's next-step button follows the nearest wave: "Clear the 3 no-go objects in Wave 1: finance foundations" with "Wave 1: finance foundations: due 16 Nov 2026, in 17 days. 1 go, 0 at risk, 3 no-go."

An object's own **Readiness** tab lists its waves with the verdict, blocking gaps, score, readiness and records blocked, and the gate's threshold in the same words as the gate ("Gate: min score 85.0"). An object no run has scored yet for its wave shows no readiness or records blocked; the tab says so instead of showing zeros.

The Migration waves panel

For each wave: **Name**, **Target date**, **Stage** (Plan, Mock load 1, Mock load 2, Dress rehearsal, Cutover) and its sign-off state ("Not signed off.").

Gate: up to three conditions, each with what fails it now:

- **Min score**: "Gate: min score 85.0", with the objects below it. The same words appear on the Readiness tab and the Waves view.
- **Zero criticals**: "No critical rule over tolerance", with the objects capped by a critical rule.
- **Contracts met**: "Data contracts met", or "No active contract names these objects."

Save the gate stores it. "No gate yet: a wave without one moves on its verdict alone."

Readiness certificate: the latest issued version and its state, for example "Version 2, waiting for signatures" or "Version 2, signed", or "Not issued". **Open the certificate** opens it: issue a new version, sign it, or download the PDF. When the gate refuses an issue, the certificate lists each object in the way, linked to its page, with the rules that stop it.

Add a wave creates one; **Objects** assigns each object to a wave or to none; **Save changes** or **Discard**. **Delete** removes a wave.

The numbers and where they come from

- **Go, at risk, no-go** counts are the verdicts of the wave's objects in the run in the top bar, from the score bands.

- **Gate state** checks each condition against the same run: the shown scores, the critical rules over tolerance, and the active contracts that name the objects.
- **Due in N days** counts from today to the target date.

Go-live flight deck

The flight deck at `/go-live` reads the next wave with a go-live date against the forecast of every object in it. Open it from **Open the flight deck** on the landscape banner.

- **Runway**: one row per object. The grey line is its daily score history; the blue line is the forecast to go-live with its 95% band shaded; the dashed red line is the wave's minimum score. A row whose forecast lands under the line is short.
- **Line meets target** is the day the slowest object's forecast crosses the minimum; it shows a dash while any object has too little history to forecast.
- **Cost of waiting** is the value at risk the wave's failing records carry today.
- **What if**: pick an object, tick the top fixes, and read the score, gates and verdict it would reach if those rules passed. Nothing is changed in SAP.
- Forecasts need three daily scores; an object with fewer shows **Forecast unavailable** instead of a guessed line.

Migration analysis

The Migration analysis page, the Migration item on the rail, shows how ready the source data is to load into the target system. The banner gives the verdict of the latest finished analysis (ready to load, ready with conditions or not ready to load), with the readiness score, the critical gaps and the records that will not load beside it. Under it:

- **Start an analysis.** Users who may analyse choose the kind of analysis (a dry run against S/4HANA, source to target, or source to source), the systems and the objects, then press **Start**. The analysis runs in the background and appears in the list when it finishes.
- **Analyses.** Each analysis with its kind, objects, verdict, score, critical gaps and records blocked. Press **Open** to see one; the latest finished analysis opens by default.
- **Gaps by kind.** The gaps of the open analysis grouped by kind and severity: values with no target value, values longer than the target field, fields the target no longer has, keys that collide, values outside a check table and fields not yet mapped.
- **Field-level gaps.** Each gap with the record, the source field and value, the target field and value, the severity and why it will not load. Filter by object, kind, severity or text; the object filter stays in the address. Users with the export permission can export the list to Excel with the same filters.
- **Records.** The records that will not load, with the reasons, and the records that are ready to load.
- **Field map.** For each object, each source field as it was read, the S/4HANA field it loads into and both descriptions. The source and target definitions (type, length, decimals, check table) stand side by side, and the **Delta** column says what changes between them: a type change, a field shortened or lengthened, a decimals change, a different check table, or a field not carried to the target. The **Transformation** column shows the note that bridges the change, whether a value map translates the values, and where the mapping came from (same field name, SAP standard conversion or a

steward). A change that can lose or reject data with no transformation yet is marked **Needs a transformation**. Bars above the table count the fields that load as read, those whose definition changes, those still without a transformation, those translated by a value map, those not carried and those not confirmed. Users who may analyse can propose the map from the source, confirm a mapping or remove it. A field that needs a value map has a **Values** button: it lists the source values found in the latest analysis with how many records carry each, and the target value saved for each one.

Every gap comes from the target data dictionary or configuration; a gap that is not grounded in either says so. Nothing is estimated: an analysis that has not finished shows a dash.

Workbook from the new system

In a migration, the new system's workbook or load file sets the load format. The Workbook from the new system panel on the Migration analysis page fills that file from a finished run. It appears only with the export permission.

1. **Choose the run and workbook.** Pick the run, then choose the `.xlsx` workbook, up to 25 MB.
2. **Check what fits. Check the workbook** reads the workbook and stores nothing. It shows the rows ready, the rows held back, how many columns each sheet can fill, and any required column with no data in this run. Open **Rows held back** to see each record, column and reason.
3. **Export the data. Export as Excel** gives one workbook; **Export as CSV files** gives a zip. Ready rows are in the workbook's own columns. Held rows are on a Held back tab. The export is recorded in the audit log.

A row is held back when a required value is empty, a value is longer than the workbook allows, a date or yes or no value is not valid, or a value is not in the workbook's list. A column with no source in the run stays empty and does not hold a row back, so check **Required columns with no data** before loading: the new system refuses those rows until the columns are filled.

S/4HANA readiness

The S/4HANA readiness page, under Migration, groups the conversion-readiness rules by SAP simplification area, such as Business Partner / CVI, Material master or the Universal Journal, over every object at its latest run. The banner says whether the landscape is ready to convert, with the areas ready, the blocking and warning rules failing and the records to fix beside it. Under it:

- **Simplification areas.** Each area with the SAP simplification item it answers, its status, how many of its rules ran, the blocking and warning rules failing and the records to fix, worst first. An area none of whose rules ran shows as not measured, never as ready.
- **What stands in the way.** Every failing rule with its name, its area, whether it blocks the conversion or only needs cleansing, the records it affects and its pass rate, blocking rules first.

A failing blocking rule always makes its area red, whatever its pass rate.

What you can do, by role

- **Programme lead** (analyse): plan waves, set gates, assign objects.
- **Data owner**: read the gate state of their objects' waves.

- **Everyone**: read the Waves view.
- **Sign off** (approve): on each wave in the Migration waves panel, once its verdict is go. The panel shows who signed it off and when; any later change to the wave clears the sign-off.

Keyboard

- `g` then `h`, then the **Waves** view switch.
- `Esc` closes the panel.

When something is wrong

- **A wave shows every object no-go but scores look close to target**: the cap holds them. The gate's **Zero criticals** line names the objects; clear their critical rules first.
- **Contracts met says no contract names the objects**: create one in People, Data contracts, or remove the condition.
- **The certificate says "Not issued"**: no version is issued yet. Open the certificate and issue one once the gate passes; if it refuses, the list under the button names the objects and rules to clear.

Administration and settings

Administration

Administration is everything that sets up the tenant: who can sign in, what the licence covers, how scoring and alerts behave, and the state of the platform. It opens from your account menu, **Administration**, for anyone who holds an administrative permission, at `/admin`. "Everything that sets up this tenant: who can sign in, what the licence covers, how scoring and alerts behave, and the state of the platform. Each section opens in this column." You see only the sections your role opens; the rest say which permission they need.

The banner on top says "Health of the setup" and names the first thing that needs an administrator, in this order: "The licence is not valid" (button **Open the licence**), "The licence ends in 21 days" (from 30 days before the end, no-go from 7), "2 people are waiting for access" (**Answer the access requests**) and "3 invitations need attention" (**Open users**). With none of these the headline reads "Nothing in the setup needs attention" and "The licence is valid, nobody is waiting for access and every invitation is on its way." Three figures: **Seats in use**, **Waiting for access** and **Licence** (the days left). **People by role** under it is a horizontal bar per role: "A person with several roles counts in each." Administration carries no run, so Aletheia is not docked on its pages and each section takes the full width.

People and access

- **Users**: "7 of 25 seats. A person's roles are grants: open a row to add or remove them." **Invite a user**, **Import users**, filters by role and status (Invited, Invite expired, Active, Locked, Suspended). A banner warns about invitations older than 7 days. Opening a person shows their grants, what they can do now and every action on the account. Needs manage users.
- **Roles and permissions**: "Built-in roles cannot be changed. Copy one to start a custom role, then set exactly what it may do. Give a role for one system or object only, from a person's grants." The matrix "What each role may do", **Access requests** (grant, or decline with a reason the person reads) and recent changes.
- **Single sign-on**: OpenID Connect with your identity provider (Microsoft Entra ID, Okta). Issuer address, client ID, client secret ("Stored encrypted and never shown again."), button name, the redirect address to register, **Test sign-in** and **Turn off single sign-on**. "People who sign in for the first time get the Viewer role and no access grants; give them access under Users."
- **Sign-in page**: the button name, allowed email domains, **Require single sign-on** ("Password sign-in stops for everyone except administrators, who keep it as a way in if the provider is down."), and the fixed password rules: "At least 12 characters, not the email name, and not a common password"; "10 failed sign-ins lock the account for 15 minutes".
- **View as a role**: "See the product as one role sees it, to answer a support question or check a role before you grant it. Everything is read-only while you view as a role, and starting and stopping are in the audit log." The top bar shows "Viewing as" until you stop.
- **Share links**: "Read-only links to one page, opened as a viewer by people of this tenant until they expire. Revoke one and it stops working at once."
- **Out of office**: "Who is away and who covers their approvals and assigned work. Each person sets their own from Preferences; an administrator can end one."

- **Approval workflows**: who approves fix batches and change requests, in which order, and when a step applies. Each workflow lists its versions with the active one marked. Edit the definition of the next version as JSON, try it with **Dry run** on sample facts (which steps apply, the approvers and how long each has), then **Save as a new version**, active at once or not. **Activate** an older version to go back to it; approvals already under way keep the version they started with. Needs manage workflows.
- **API keys**: "A key acts as its own service account with only the roles you give it, and shows in the audit log under its name. The key itself is shown once, when you create it."

Tenant

- **Licence**: company, status, tier, seats, validity, last checked, platform version and the licensed modules and features. **Check with Meridian HQ now**. "Plan changes and invoices are handled in Meridian HQ; this deployment keeps its data either way." Every licence is verified against the vendor's Ed25519 key built into the image. When Meridian HQ cannot be reached, the last valid signed answer is honoured for 14 days after it was signed: the page still reads Valid while the API reports the status `offline_grace` with the days left; after that, gated pages lock until Meridian HQ answers again.
- **Billing**: resolved exceptions by billing tier on top of the base fee; read only.
- **Settings**: the tenant; **Dimension weights** ("The score is each measured dimension times its weight. The weights must add up to 100." **Use the DAMA defaults**); **Score bands** ("At or above target is on target (go), at or above warning is at risk, below that is failing."); **Cap rule and tolerance**: "One critical rule failing beyond its tolerance holds that object at 85; two or more hold it at 70. The caps are fixed. The tolerance says how many failing records count as beyond.", with **Failing share, percent** and **Failing records** and the severity weights and multipliers; **Cost model and currency**; **Planner assumptions**; **Retention** (keep runs, reports and audit rows for so many days, keep the baseline; purges rehearse first and the purge log records each one); **Delete all data**, which issues a token that works after 24 hours. **Sandbox mode** is a switch: turning it on (**Enter sandbox mode**) copies the latest finished run into a sandbox run; until it is turned off (**Exit sandbox mode**), decisions land on the copy only, decisions on the real run are refused with the reason ("Sandbox mode is on: decisions apply to the sandbox run only."), and notifications, webhooks, digests, report mailouts and write-backs are held back and counted. A band across every page says "Sandbox: nothing leaves this deployment." with since when and by whom. Entering and exiting need manage system; exiting discards the sandbox runs.
- **Stewardship**: "One accountable owner per object and one steward per company code, plant or organisation." The template per object (owner role, owner, stewards by slice, steward, teams), **Save the template**, **Check what applying does** (a dry-run diff), **Apply**, and **Gaps**: "Every object needs one owner and at least one steward per active slice", for example "Vendor master (AP): No steward for purchasing organisation 1000".
- **Configuration bundle**: "Rules switched on or off, field mappings, waves and their gates, settings, SLA policies and the escalation chain as one JSON file." **Download the bundle**; import with a diff first: "Import adds and changes; it never deletes what the file leaves out. Check the diff first."

Data

- **Field mappings**: "Name the column in your files that holds each standard SAP field, then save the row." **Reset all to defaults**.
- **Reference data**: postal code and bank code (SWIFT BIC) lists per system. "Until a list is loaded those rules check the format only." An upload replaces the system's current list.
- **Match rules**: the field comparisons that decide when two records of the same object are one, grouped by object. Each rule names a field, how it is compared (exactly equal, similar spelling, sounds alike, close in value or same meaning), its weight and the similarity needed. A pair scoring 95 or more is merged by itself, below 30 it is dismissed, and anything between goes to a steward; a steward's do-not-match or always-match decision overrides the score. **Simulate on pending pairs** scores the latest pending pairs with the active rules and shows where they would go without writing anything. Needs the manage rules permission. Under the rules, **Proposed by Aletheia** lists the rules Aletheia drew from steward corrections once stewards correct ten or more pairs of an object in a week: the field, how it would be compared, its weight and similarity, why it was proposed and how many corrections are behind it. **Approve** adds it to the match rules as active; **Reject** takes an optional reason for the audit log. Reviewing proposals needs the review AI rules permission.
- **Rules import and export**: "Every rule of the tenant as YAML or Excel, and many rules at once back in. Imported rules are validated against the data dictionary and land as drafts for review, never live." **Check the file**, then **Create drafts**.
- **Policy packs**: GxP, SOX and GDPR. Each pack names field classes, the authority each answers to (an article or section, such as "SOX Section 404(a) (controls over changes to vendor payment data)") and the catalogue rules that evidence it. A card per pack shows fields classified, rules evidencing, rules measured, failing records and classes breached, and whether it is applied. **Apply the pack** and **Remove the pack** need manage rules. Open a pack for its classes, their fields and the evidence of each rule, for the current state or one run. An applied class with a failing critical rule is a policy breach: a filled shield beside the score on the landscape cell, the ruler and the object's receipt, which names the class, its authority and the rules.
- **Database sources**: read a table, a view or one SELECT from a non-SAP database and check the rows with the same rules as a file upload. **Add a connection** with host, port, database, a read-only user, a password and the encryption level; the password is stored encrypted and never shown again. **Test the connection**, then choose **A table or view** or **One SELECT statement**, **Preview the rows**, name the standard field each column becomes and **Load and check**. Meridian runs a read-only session, one statement, at most 100,000 rows and 60 seconds, and refuses loopback, link-local and cloud metadata addresses and the addresses of Meridian's own services; other private network addresses are allowed. A SELECT that uses a dollar sign outside a quoted string, an `E''` string or a nested comment is refused, because Meridian cannot prove it is one statement. PostgreSQL is supported; SQL Server, Snowflake, Databricks and BigQuery are not. A read-only session does not stop a function that writes on the server, so give the connection a SELECT-only user. Needs the manage systems permission.
- **Data quality firewall**: "An SAP BAdI, middleware or a pipeline posts up to 100 candidate records and gets back, at once, the enabled rules each one fails." Send `POST /api/v1/dq-firewall/validate` with an API key whose role may change master data, the object (for example `material_master`) and the records as TABLE.FIELD values. Each record comes back with its verdict

and its failing rules: rule, what it requires, field, severity, dimension and why it failed. A record with a failing critical rule is blocked, and so is the call as a whole; a failing high, medium or low rule is reported without blocking. An unknown object or field name is refused with a plain sentence. The same enabled rules and tenant overrides as the analysis apply. Add `"route_to_steward": true` to open a create change request for each blocked, warned or duplicate-flagged record; the answer carries the request id, or a reason when none was opened, and the same content returns the request that is still open. Add `"hold": "blocked"` or `"flagged"` to stage the record in the approval workflow too: each covered record carries a hold state (held, released, rejected, written or not held) and the answer says `save` stop or go; the default `"off"` leaves the answer unchanged. The request is made by the key's service account. Only an audit entry is stored (counts and blocking rule ids, never record values), and no language model sees the records. The section shows the endpoint, a sample call to copy, calls in the last 24 hours and 7 days, and a **Try it** box that runs **Check records** on records you paste. Needs the manage settings permission.

- **Data subject requests**: find, export and erase what Meridian holds for one record key (the whole `FIELD=value|FIELD2=value2` key, or a value alone). **Search** counts the rows per table and in the extract files, each table shown with a line on what it holds. **Export what is held** downloads a package of those rows. **Dry run an erasure** lists exactly what would go and deletes nothing. **Erase permanently** appears only after a dry run and needs the word ERASE typed and a reason of at least ten characters; it removes the rows (or clears the key and the values where the row must stay, as in the cleansing audit), rewrites the extract files, and records the request. The request list keeps a keyed hash of the subject, counts, the reason and who asked, never the key itself. What erasure cannot do is printed on the screen: Meridian does not delete from SAP, so erase the record there as well; a completed erasure puts the record key on a suppression list (a keyed hash, never the key) so a later extraction drops it, the screen shows how many records were dropped, and an administrator can lift an entry, which is recorded; the list matches the record key only, so the same person under another key comes back, and a few readers of extract files are not filtered; backups are out of reach; raw uploads, reports and archived files are not searched; extract files are matched on the full key only; the search finds record keys, not names or other personal values; an open change request on the record stops the erasure until it is finished or cancelled; a golden record under another key may keep values merged from this one; and comments, attachments and workflow events on change requests, requests to create a record that have no key yet, and sample key lists are not changed. Erasure also waits while an extraction, analysis, upload or configuration sync is running. If one starts after the database rows have gone, the extract files stay as they were and the request is marked incomplete: run the same erasure again with the full record key once the run has finished. Needs the manage settings permission, and no language model sees any of it.
- **Audit log**: every change, who made it, the object and the result, filters by person, action, object kind and dates, **Export CSV** and **Export Excel**, and **AI calls**: "Every call to the language model: which service, how long, and whether a deterministic rule answered first. Prompts and records are never stored here."

Notifications

- **Alerts**: channels (email, Microsoft Teams, signed webhook) with their digest, board pack and critical alerts; **When to alert** (critical and high failing rules, score drop in points) and **Score floors per object**. "Alerts carry counts, rule ids and links, never record values."

- **Default digest**: the tenant's default digest. Each person sets their own under Preferences.
- **Digests**: scheduled summaries for a role (the programme lead every Monday at 07:00) or a migration wave: the landscape sentence, score moves, the waves and their gates, approvals waiting, SLA breaches and the top value at risk. The form takes the cadence (a preset or a cron line, checked as you type), the time zone, the channel (email with addresses and people, an alert channel, or the `digest.sent` webhook event), **Attach the digest as a PDF** and **Send on its schedule**. Each digest is one row: its name with who it is for, the cadence and time zone, the channel, when it was last sent ("Sent 27 Oct 2026, 07:00" or "Never sent") and the next send, and **On** or **Paused**. Select a row to open it. In the digest's inspector: **Send now**, **Preview** (the email in a sealed frame, or the numbers it is built from), **Pause**, **Edit**, **Delete** and the delivery log. Needs manage settings.

Operations

- **Start and schedule** (the runs themselves are on Runs): start an extraction (objects, company codes, plants, document dates, analyse straight after), start an analysis of a stored extraction, load the configuration, scheduled runs, **Upload a file** (a CSV, TSV, Excel, JSON or Parquet file of one object becomes a run and is analysed at once; its columns are matched by the field mappings; needs upload) and recent jobs with **Retry**.
- **Platform and updates**: jobs by state and kind, the platform version and **Update to** the latest when Meridian HQ knows one, **Health checks** (database, cache and queue, background workers, licence, email delivery, file storage): "Each check says what it found and what needs fixing. They refresh every 30 seconds.", and **Support** with **Download diagnostics bundle**: versions, health checks, recent errors and configuration with every secret removed, and no SAP record values.

Provisioning, theming and compliance

- **Provisioning (SCIM)**, under Users: "Microsoft Entra ID, Okta and other identity providers create, update and suspend people here through SCIM 2.0. Give them the base URL and a token." **Create a token**, **Rotate the token** (the old one stops at once), **Revoke**, and **Group to role**: "A person in one of these identity provider groups gets its role. Saving re-syncs everyone in a SCIM group at once." Only grants SCIM made are removed, and the last active administrator is never left without admin.
- **Appearance**, under Settings: the tenant's accent, logo and name, and a **Dark-mode accent** for people who choose the dark theme, with the contrast of the accent checked before it is saved. Light is the default theme; each person can switch to dark under Preferences, Display.
- **Aggregates only to the model**, under Settings, Language model: a switch, off by default. Aletheia never sends SAP records to the model, only counts, lengths and shapes; a column that holds personal data (found by its saved mapping, its header words or the shape of its values) never sends its shape either. Turn the switch on and no sample from a row reaches the model at all: it gets counts and totals, and matching, scoring and rule drafting use their deterministic answers where a sample would have helped. The change is recorded in the audit log. Needs the manage LLM permission.
- **Compliance**, under Settings: data residency, retention days per table family, encryption flags and the last backup.

Sign-in, audit and tenancy

- The first sign-in of an account that must change its password (the seeded administrator, or an account an administrator reset) goes to the password change before anything else.
- The audit log records sign-in, failed sign-in, lockout, sign-out and password change (`auth.login`, `auth.login_failed`, `auth.lockout`, `auth.logout`, `auth.password_change`), and single sign-on (`auth.sso_login`, `auth.sso_refused`, `auth.sso_user_created`). A failed sign-in keeps the email as typed and the reason, never the password.
- A signed-in request runs in the tenant its verified token names; the deployment's own tenant is set by `MERIDIAN_TENANT_ID`.

The numbers and where they come from

- **Seats** count active and invited users against the licence.
- **Settings** hold the numbers every score uses: the dimension weights, the score bands, the tolerance (failing share and failing records) and the severity weights. A change applies to the next analysis; scores already stored with a run are kept as they were.
- **Health checks** are measured by the platform itself and refresh every 30 seconds.
- **Policy pack coverage** counts the fields of the pack's classes, the rules that check them, the rules a run has measured and their failing records in the current state; a class is breached when one of its critical rules fails.
- **Digest numbers** are the landscape's own: the preview's numbers are what the email is rendered from.
- **Billing** counts resolved exceptions by tier for the month, on top of the base fee, from Meridian HQ.

What you can do, by role

- **Platform admin**: every section.
- **Manage users**: users, roles, single sign-on, sign-in page, view as, share links, out of office, API keys.
- **Manage settings**: settings, billing, stewardship, configuration bundle, alerts, default digest, webhooks, digests.
- **Manage system**: licence, platform and updates, the diagnostics bundle, sandbox mode.
- **Manage rules**: apply and remove policy packs; anyone who holds view reads them.
- **View audit** (auditor): the audit log, and a person's activity in People.
- **Manage field mappings**, **manage systems**, **manage rules**: field mappings, reference data, rules import and export.
- **Trigger sync**, **analyse** or **upload** (SAP Basis, analyst, programme lead): start and schedule.
- Anyone who holds view reads the licence, settings, mappings, reference data, alerts and the default digest, read only.

Administrator training guide

The course for the administrator and the SAP Basis colleague (install and first start, licence, users and access, connecting SAP, extraction and schedules, scoring and money, write-back, notifications, operations and troubleshooting) is the printable administrator training guide, `docs/admin-training/meridian-administrator-training.pdf` in the Meridian repository. Ask your Meridian contact for it if you do not have the repository.

Keyboard

- In every section table, `j` and `k` move and `Enter` opens the row in the inspector.
- `Esc` closes the inspector; the section index stays on the left.

When something is wrong

- **"... needs the ... permission. Ask an administrator to grant it."** Your role does not open that section.
- **"These settings are read only for your role."** You can read but not change them.
- **Health check warning "No email backend configured"**: digests and email alerts are not sent until the deployment's email settings are set.
- **Single sign-on test fails**: check the redirect address registered at the identity provider matches the one shown, character for character.
- **"The tenant is already in sandbox mode."** or **"The tenant is not in sandbox mode."**: someone else entered or exited first; the page reads the state again.
- **A digest shows Held back**: the tenant is in sandbox mode, so nothing was sent; the delivery log keeps the attempt.
- **The interactive API documentation answers 404**: `/docs` and `/openapi.json` are served only on a development deployment or with `MERIDIAN_OPENAPI_PUBLIC=1`.
- **A purge request**: purges are irreversible. The token only works after 24 hours so a mistake can be caught; the purge log keeps the record.

Preferences

Preferences is how Meridian looks and reads for you, where it starts, and what it tells you about. It belongs to you, not to the tenant. Open it from your account menu, **Preferences**, or with `?open=preferences`. Every change is saved at once.

What is in it

- **Display**: **Theme** (Light, Dark or System; Light is the default, and the account menu switches between light and dark too) and **Row density** (Compact, Default or Comfortable).
- **Language and region**: **Number and date format** (Browser default or a locale such as en-ZA or de-DE) and **Time zone** (UTC by default). A sample line shows the result: "Dates read 16 Oct 2026, 05:17, scores 71.8, counts 1,234,567."
- **Start screen**: **Open Meridian as** your role's default or another role's start screen. "It changes the next step on the landscape and the work it shows first. Your access does not change."
- **Keyboard**: "Show keyboard hints beside menu items and places".
- **Notifications**: the digest of your open items (Off, Daily, Weekly), "Tell me at once about critical findings" and "Also email me". **Inbox and triage settings** opens triage.
- **What you hear about**: the events that reach your inbox, such as run finished, score drop, new critical breach, batch waiting and mentions, each on or off.
- **Out of office**: set from People, Out of office, or `?open=away`.

The numbers and where they come from

Nothing here changes a number. The locale changes how numbers and dates are written; the time zone changes the times shown; scores and counts stay the same.

What you can do, by role

Everyone sets their own preferences. The tenant's default digest and alert channels are set by an administrator under Administration, Default digest and Alerts.

Keyboard

- `Tab` moves through the settings, `Space` toggles a switch.
- `Esc` closes the panel.

When something is wrong

- **Dates show in UTC**: set your time zone here.
- **No email arrives**: "Also email me" is off, or the deployment has no email backend (Administration, Platform and updates, Health checks).
- **The start screen did not change**: the start screen is kept in this browser; another browser keeps its own.

Keyboard and the command line

Meridian can be driven from the keyboard: a whole ledger can be worked without the mouse, and every canvas and action is reachable by typing. `?` shows the full sheet at any time; the keys are the same on every screen.

The command line

Open it with `⌘K` (Ctrl+K on Windows and Linux) or `/`. "Search, ask, or type a command." Type:

- a **record key** such as a material number to open the record;
- a **rule id** such as `MM142` to open the rule;
- an **object and system**, such as "material master ECC", to open the object;
- a **run** date to open the run;
- a word: **approve** opens the batches waiting for approval; **export** saves the fix sheet of what is on screen.

Start with `?` to ask in a sentence, such as `?critical vendors assigned to me due this week`. The answer is the ledger filter that will run, shown as chips with the slot each phrase fills (Object, Severity, Owner, Due, Policy, System), the explanation, and the words it did not understand; **Open in Work** (or `Enter`) opens exactly that filter. The words are matched against fixed tables; nothing is guessed.

With nothing typed it offers **Go to**: Landscape, Objects, Work, Runs, People, Rules and Connect a system. Typing a page name such as Reports, Requests or Configuration goes to that page. ↑ and ↓ move through the results, `Enter` opens one, `Esc` closes.

Everywhere

- `⌘K` or `/`: open the command line.
- `?`: show the shortcut sheet.
- `⌘Z`: undo the last change (when focus is not in a text field).
- `Esc`: close the top layer: menu, then inspector, then the selection.

Go to

- `g` then `h`: Home.
- `g` then `o`: Objects.
- `g` then `r`: Runs.
- `g` then `f`: Fix.
- `g` then `w`: Work.
- `g` then `u`: Rules.
- `g` then `m`: Master data.
- `g` then `i`: Migration.

- `g` then `p`: People.

Tables and lists

- `j` or ↓: next row. `k` or ↑: previous row.
- `Shift+J` or `Shift+↓`, `Shift+K` or `Shift+↑`: extend the selection.
- `Enter`: open the row in the inspector.
- `x`: select or clear the row.

Work ledger

- `a`: accept the suggested value on the selected rows.
- `n`: mark the selected rows as not an issue.
- `r`: accept the risk on the selected rows.
- `e`: export the fix sheet of what is on screen.
- `f`: move to the filter chips.
- `⌘D`: fill the New value of the first selected row down the selection.
- Paste several lines into a New value cell to fill the selected rows.

Inspector

- `]`: open the next row. `[`: open the previous row.
- `Tab` cycles inside the inspector.
- `Esc` closes it and returns you to where you were.

Where the numbers come from

Nothing typed in the command line is guessed: results come from the tenant's own records, rules, objects and runs, grouped by kind.

What you can do, by role

Every role can use every key. A key that acts (`a`, `n`, `r`) does nothing visible when your role does not hold the action; the ledger says "Read-only for your role".

When something is wrong

- **A key does nothing**: focus is in a text field. Press `Esc` or click the page, then try again. `⌘K` and `Esc` work everywhere.
- **"Select rows first: x selects the highlighted row, Shift+J extends the selection."** A ledger key was pressed with no row selected.
- **The command line finds nothing**: try a shorter word, or the record key as SAP stores it, with its leading zeros.
- **Keyboard hints are missing beside places**: turn them on under Preferences, Keyboard.

Walkthroughs by role

Walkthrough: programme lead

You run the data programme: you need to know how ready each system and object is for the next wave, what holds it back, who is working on it, and what to tell the steering committee. Your start screen is the **Landscape** with the next-step button for your nearest wave.

Your first five steps

1. **Landscape**: read the go-live banner on top (the verdict for your next wave, what blocks it, days to go-live and the cost of waiting), then press the next step, for example "Clear the 3 no-go objects in Wave 1: finance foundations".
2. **Waves**: switch the landscape to **Waves**, open **Edit** on the wave and set its gate: min score, zero criticals, contracts met. **Save the gate**.
3. **An object**: open a no-go object and read its receipt: measured, capped, by which critical rules.
4. **People**: see who holds the open work and **Balance work** if one person holds most of it.
5. **Reports**: schedule the board pack for the steering committee.

Your jobs and where they are done

- **Landscape readiness by system and wave**: Landscape, Coverage and Waves views.
- **Plan waves (objects, dates, stages)**: Migration waves panel, **Add a wave**, **Objects**.
- **Wave gate criteria**: Migration waves panel, **Gate**, **Save the gate**.
- **Gate status**: the Waves view, "Gate not passed, 2 of 3 not met".
- **Readiness certificate**: the waves panel shows the latest issued version and its state ("Version 2, waiting for signatures", "Version 2, signed" or "Not issued"); the certificate link opens it to issue a version, sign it or download its PDF.
- **Sign off a wave**: **Sign off** on the wave in the Migration waves panel (needs approve); only a wave whose verdict is go can be signed off, and any later change to the wave clears the sign-off.
- **Balance work between stewards**: People, **Balance work** (even counts, whole objects or by the assignment template), applied with undo.
- **Read data contracts and their breaches**: People, **Data contracts**.
- **Watch objects**: **Watch** on the object header for run, score and breach notices.
- **Export the ledger or any grid**: **Share and export**, **Export to Excel**.
- **Share a canvas with the steering committee**: **Share and export**, **Share this page as a link** (read-only, expires).
- **Start extraction and analysis, schedules**: System inspector; Administration, Start and schedule.
- **Reports and report schedules**: Reports.

The numbers you report

- **The shown score** is capped: 85 at one critical rule over tolerance, 70 at two or more. Report the measured score beside it when it matters; the receipt states both.
- **Value at risk** is from the cost model named in the footer; it is an estimate from your own cost settings, not a forecast.
- **Go, at risk, no-go** come from the score bands under Administration, Settings.

When something is wrong

- **The steering committee sees different numbers**: check which run each of you reads. The run is in the top bar and in every report.
- **A wave cannot pass its gate**: open each object named under **Zero criticals**; the critical rules over tolerance are listed in its receipt.
- **Contact**: SAP Basis for extraction failures, your administrator for access, the object owner for approvals that wait.

Walkthrough: data steward

You clean the data: you decide the right value for each failing field, accept what Meridian suggests where it is right, sign off what is not a problem, and hand your batch to a second person. Your start screen is the **Landscape** with "Accept N suggested values for" the top rule as your next step.

Your first five steps

1. **Landscape**: press **Accept N suggested values for** the top rule; Work opens on that rule's rows.
2. **Work**: read the top rows, ranked by value at risk. Select rows and press `a` to accept their suggested values; the notice has **Undo**.
3. **A record**: press `Enter` on a row without a suggestion. Read the possible values with their meanings, type the New value and press `Enter`.
4. **By rule**: switch the ledger to **By rule** and decide a whole rule at once, or **Export fix sheet** (`e`) to work offline and **Import filled sheet** when done.
5. **Waiting for approval**: check your draft batch under the chip; a data owner approves it.

Your jobs and where they are done

- **Work the ledger: fix, propose, not an issue, accept risk**: Work ledger and record inspector.
- **Accept suggested values in bulk**: **Accept all** on a rule, or the next-step button.
- **Build a fix batch for approval**: every accepted or entered value goes into your draft batch.
- **Undo a decision**: the notice's **Undo**, or `⌘Z`.
- **Comment and mention**: Inspector, **Comments**; type `@` and a name.
- **Watch a rule or record**: Inspector, **Watch**.
- **Fix the upstream object first**: the fix-order hint, for example "Fix Material master first: 3,640 upstream failures. Records there fail before this object can pass."
- **Out of office with a deputy**: Preferences or People, **Out of office**.
- **Saved and shared views of the ledger**: **Views**, **Save current view…**, shared with everyone if you choose.
- **Balance work in the team**: People, **Balance work** (needs assign).
- **Author rules**: Rules, **New rule draft** (needs manage rules).
- **Import rules in bulk**: Administration, Rules import and export; there is no link from the rules library yet.
- **See your own activity**: your row in People, **Activity**.

The numbers you work with

- **Value at risk** ranks the rows; it is the cost-model estimate of what each failing value costs.
- **Possible values** come from the rule; a value outside them still fails and is refused.

- **Suggested value** is the rule's proposal, "Blank when no single value is right."

When something is wrong

- **The ledger says "Read-only for your role"**: you need the apply permission. Ask your administrator.
- **A batch comes back**: the approver sent it back with a reason; it is in the batch inspector.
- **A fixed record returns to To fix**: the next run still found it failing after it was sent to SAP. Check with whoever loads the file.
- **Contact**: the object's owner for approvals, SAP Basis for extraction gaps.

Walkthrough: data owner

You are accountable for one or more objects and you are the second pair of eyes: you approve or send back the fix batches stewards prepare, sign off exceptions, approve contracts and glossary terms, and decide who covers for you when you are away. Your start screen is the **Landscape** with the line "2 batches wait for your approval, SLA met on 92% of your open work" and the next step "Approve batch B-241".

Your first five steps

1. **Landscape**: read the line under the sentence and press **Approve batch B-241**.
2. **Batch inspector**: read the facts, the changes ("`MARA-MEINS` STK to EA") and the approval trail. **Approve** or **Send back** with a reason.
3. **Waiting for approval**: in Work, the chip lists every batch waiting, grouped by batch.
4. **Your object**: open an object you own; **Watch** it and apply the steward template to it.
5. **Out of office**: before you are away, set who covers your approvals.

Your jobs and where they are done

- **Approve or send back a batch (four eyes)**: Batch inspector, **Approve**, **Send back**. You never approve a batch you built.
- **Out of office: approvals route to the deputy**: Preferences or People, **Out of office**; the batch and person inspectors show the cover.
- **Escalation chain after an SLA breach**: Work or People, **Triage and SLA**, **Escalation chain** (changing it needs manage triage).
- **Create and edit a data contract**: People, **Data contracts**, **New contract**; the contract inspector.
- **Approve a contract (never your own)**: Contract inspector.
- **Contract versions and revisions**: Contract inspector, **Versions**, **Revise as a new version**.
- **Contract breaches**: Contract inspector, **Breaches**.
- **Apply the steward template to your objects**: the object canvas, steward template.
- **Watch your objects; subscriptions**: Object canvas, **Watch**; Preferences, What you hear about.
- **Gate status of your objects' wave**: Landscape, Waves view.
- **Sign records off**: Work, **not an issue** or **accept risk** with a reason.
- **Merge pairs and approve glossary terms**: an object's Golden records and Glossary tabs.

The numbers you approve

- **Records** and **Value at risk** on a batch are the records in it and their cost-model estimate.
- **SLA met** is the share of your open work resolved within the SLA days set under Triage and SLA.

When something is wrong

- **Approve is missing**: you built the batch, or your role lacks approve.
- **Approved values do not appear in SAP**: approval does not touch SAP. The batch is exported and loaded by your team's process; the next run confirms each record.
- **Contact**: the steward named on the batch for its content, your administrator for access.

Walkthrough: SAP Basis

You keep the plumbing healthy: connections, extractions, schedules, configuration loads and the jobs that run them. Meridian reads from SAP with a read-only role; it writes back only where an administrator has turned on live write-back for material master, customer and business partner, through a separate SAP user. Your start screen is the **Landscape** with your own home above the coverage matrix: **Connections** (each system, its connection state and what the last test said), **Last extraction** (per system, with records and whether it was analysed), **Failed jobs to triage** (a job that failed and was not followed by a success of the same kind), **Authorisation errors** (refused credentials and missing authorisations among the connection tests and the jobs) and **Schedules** (system, object, when, next run and mode). The next step names the system that needs you first, for example "Extract SuccessFactors", and **Your next work** lists the system tasks.

Your first five steps

1. **The Basis home**: read Connections, Failed jobs to triage and Authorisation errors first, then the task list: "Extract ECC production again", "Run analysis on S/4HANA sandbox", "Extract SuccessFactors".
2. **System**: open the system and **Test connection**. A passing test reads "RFC connection succeeded (client 100, 41 ms)".
3. **Dry run**: in the system inspector choose the objects, run **Dry run** and read the rows and duration; then **Extract**, or **Extract changes only** where available.
4. **Schedules**: add a schedule with a rhythm and the mode **Changes only** for routine runs.
5. **Runs**: watch the job line and open a failed job for its stages and **Retry**.

Your jobs and where they are done

- **Test a connection, health**: System inspector, **Test connection**.
- **Extract objects**: System inspector, **Extract**.
- **Extraction dry run: row counts per table and a duration estimate**: System inspector, **Dry run**.
- **Delta extraction (changes only)**: System inspector, **Extract changes only**; schedules in **Changes only** mode.
- **Schedules, retry, cancel**: the system's Schedules view; Administration, Start and schedule.
- **Field mappings**: Administration, Field mappings.
- **The read-only SAP role and the table reader**: `docs/sap/authorisation-role-template.md` lists the S_RFC function modules and the S_TABU_NAM tables; `MERIDIAN_RFC_READ_FUNCTION` in `.env` chooses `RFC_READ_TABLE` (default), `/BODS/RFC_READ_TABLE2`, `BBP_RFC_READ_TABLE` or `auto`.
- **Load configuration and compare with the target**: the system's Configuration view.
- **Custom objects**: the system's Custom objects view.
- **Source design**: the system's data dictionary against the SAP standard, configuration deviation and drift between discoveries.

- **Watch a system**: possible through the API only; the System inspector has no Watch button yet.

The numbers you watch

- **Records** per extraction and per table, from the job registry; the Basis home's Last extraction shows the latest per system.
- **Failed jobs to triage** and **Authorisation errors** count what still needs you.
- **Areas loaded**, for example "11 of 12", and per area what could not be read and why.
- **Health checks** under Administration, Platform and updates.

When something is wrong

- **HTTP 401 or "Sign-in refused"**: rotate the credentials in Connection settings and test again.
- **"No authorisation to read this table. The user needs read access to it."** Grant the RFC user read access from the role template, then **Load again**. SU53 shows the failed check.
- **The table reader is refused**: set `MERIDIAN_RFC_READ_FUNCTION` to the reader your role grants, or `auto`.
- **A scheduled extraction did not start**: due runs start within 5 minutes; check Runs, Jobs and the background workers health check.
- **Contact**: your administrator for permissions on a system; the programme lead for which objects to extract.

Walkthrough: process analyst

You read the data through the business process: which steps of Procure to Pay, Plan to Produce or Order to Cash the data blocks, what fixing a rule would release, and whom to hand a blocking rule to. Your start screen is the **Landscape** with "Open PTP step 3" as your next step.

Your first five steps

1. **Landscape**: press **Open PTP step 3**.
2. **Processes tab**: read "2 of 4 process steps are blocked by Material master: PTP step 3, PP step 1." and the fields behind each blocked step.
3. **What if**: tick the rules to fix under **What if these rules were fixed** and see which steps would be ready.
4. **Step inspector**: **Hand off** the blocking rule to a steward with a note.
5. **Process designer**: map the fields each activity reads, then publish the model.

Your jobs and where they are done

- **Process readiness by step**: Object, **Processes** tab.
- **What if these rules were fixed (simulation)**: Processes tab; computed in the browser, nothing changes and no data is read.
- **Hand a blocking rule to a steward**: Step inspector, **Hand off** (needs handoff rule).
- **Export the as-is process**: Processes and the designer (needs export process).
- **Comment on a rule or a step's rule**: Rule inspector, **Comments**.
- **Related objects and fix order**: Object header, related objects; the fix-order hint in Work.

The numbers you read

- "A step is blocked when a field it writes fails a critical or high rule, or passes under 70%."
- Failing counts are records, and a record can fail several rules.

When something is wrong

- **A step you expect blocked is Ready**: its fields are not mapped to the activity. Map them in the designer; it shows on the next run.
- **Hand off is missing**: your role lacks handoff rule.
- **Contact**: the object's owner for priorities, the steward the rule was handed to for progress.

Walkthrough: auditor and viewer

You read: every canvas, every number and its receipt, and, as an auditor, the audit log and each person's activity. Nothing you do changes data. Your start screen is the **Landscape**, the same one the programme lead sees, read-only. A viewer may also arrive through a share link, which opens one page read-only with the bar "Shared view. Read-only, as a viewer."

Your first five steps

1. **Landscape**: read the go-live banner, the sentence, the landscape map and the coverage matrix; hover a tile or a cell for its reading.
2. **An object**: open the receipt and every **i**: the arithmetic behind each figure.
3. **Approvals**: in Work, **Waiting for approval** and **Exceptions** show the four-eyes evidence: who drafted, who approved, who signed off and why.
4. **Audit log** (auditor): Administration, **Audit log**; filter by person and dates and **Export Excel**.
5. **Person activity** (auditor): People, open a person, **Activity**.

Your jobs and where they are done

- **Read every canvas**: all canvases, read-only.
- **Audit log and export (auditor)**: Administration, Audit log.
- **A person's audit trail (auditor)**: Person inspector, **Activity**.
- **Open a read-only share link**: `/share/{token}`; the bar says "Shared view", and **Leave the shared view** returns you.
- **Export a grid (auditor)**: **Share and export**, **Export to Excel**.
- **Read comments**: Inspector, **Comments**.
- **Rule change trail joining versions and runs**: not joined yet; read the rule's **Versions** and **History** side by side.

The numbers you check

- Every object score has a receipt: dimension points, measured, cap and shown. The formula line states the weights, the severity weights, the tolerance and the cap rule.
- Comparisons state their method: the common set of rules and the change threshold.
- The audit log's **AI calls** list every call to the language model, which service made it and whether a deterministic rule answered first; prompts and records are never stored there.

When something is wrong

- **A share link says it expired or was revoked**: ask the person who shared it for a new one.
- **The audit log is missing**: it needs the view audit permission.
- **Contact**: your administrator for access; the programme lead for the run a report was built on.

Walkthrough: administrator

You set the tenant up and keep it right: who signs in and with which roles, how scoring and alerts behave, the licence, the platform version and its health, and the configuration you move between tenants. Your start screen is the **Landscape**; **Administration** is in your account menu.

Your first five steps

1. **Users**: invite people and give each a role; scope a role to one system or object from their grants.
2. **Single sign-on and sign-in page**: connect your identity provider, test it, then decide whether to require it.
3. **Settings**: check the dimension weights, score bands, cap rule and tolerance, cost model and retention.
4. **Stewardship**: name an owner per object and stewards per slice, **Check what applying does**, then **Apply**; close the **Gaps**.
5. **Platform and updates**: read the health checks and update the platform when a new version is available.

Your jobs and where they are done

- **Invite users, set roles**: Administration, Users.
- **Custom roles and scoped grants**: Administration, Roles and permissions.
- **Single sign-on**: Administration, Single sign-on and Sign-in page.
- **API keys**: Administration, API keys.
- **Audit log and export**: Administration, Audit log.
- **Export and import the tenant configuration with a dry-run diff**: Administration, Configuration bundle.
- **View the product as a role (support)**: Administration, View as a role; the bar shows "Viewing as".
- **Import and export rules in bulk (YAML and Excel)**: Administration, Rules import and export.
- **Govern share links**: Administration, Share links, **Revoke**.
- **See everyone's out of office**: Administration, Out of office, **End it**.
- **Steward assignment template, apply with a dry-run diff, audit the gaps**: Administration, Stewardship.
- **Licence, platform update, health**: Administration, Licence; Platform and updates.
- **Tenant settings (scoring, cost, retention)**: Administration, Settings.
- **SCIM provisioning and group to role mapping**: Administration, Users, Provisioning (SCIM).
- **Theme and compliance record**: Administration, Settings, Appearance and Compliance.
- **Webhooks, digests, the default digest**: Administration, Webhooks, Digests, Default digest.
- **Diagnostics bundle for support**: Administration, Platform and updates, Support.

- **Your own training**: the administrator training guide, `docs/admin-training/meridian-administrator-training.pdf`.

The numbers you set

- Dimension weights must add up to 100.
- The caps are fixed at 85 and 70; you set the tolerance: the failing share and the failing records that count as beyond.
- Score bands: target, warning and critical.

When something is wrong

- **Someone cannot see a section**: compare their role with the matrix under Roles and permissions, or use View as a role.
- **Email is not delivered**: Health checks, Email delivery.
- **Your first sign-in asks for a new password**: the seeded account and any reset account must change it first: 12 characters or more, not the email name, not a common password.
- **Sign-in events**: the audit log records sign-in, failed sign-in, lockout, sign-out and password change; filter by action `auth.login` or `auth.login_failed`.
- **The licence shows fewer modules than expected**: modules come from the licence in Meridian HQ; **Check with Meridian HQ now**, then contact Meridian support.
- **Support**: open help, copy the diagnostics under **Contact support** and send it with your request. It holds no record values and no secrets.